

Learned Cost Models for Query Optimization: From Batch to Streaming Systems

VLDL 2025



Roman Heinrich

TU Darmstadt & DFKI



Xiao Li

ITU Copenhagen



Manisha Luthra Agnihotri

TU Darmstadt & DFKI



Zoi Kaoudi

ITU Copenhagen

Presenters

Dr. Manisha Luthra Agnihotri

Athene Young Investigator (TU Darmstadt)
& *Deputy Head* (DFKI GmbH)

Research focus:

- AI-enhanced streaming systems
- Multimodal streaming
- Benchmarking AI-enhanced streaming



Prof. Zoi Kaoudi

Associate Professor
(IT University of Copenhagen)

Research focus:

- ML-based query optimization
- Cross-engine data systems



Roman Heinrich

PhD Student
(TU Darmstadt and DFKI GmbH)

Research focus

- Learned cost models &
- Query optimization of data systems



Dr. Xiao Li

Postdoc
(IT University of Copenhagen)

Research focus

- Query optimization with ML
- Data cleaning with ML

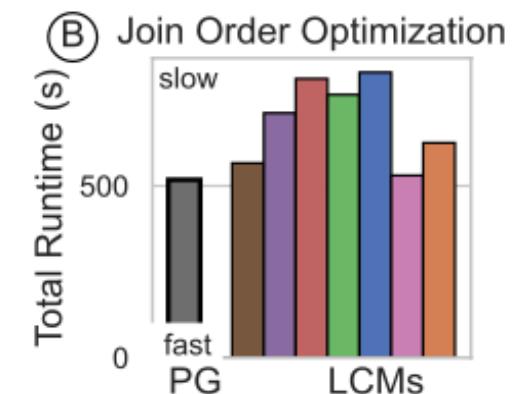
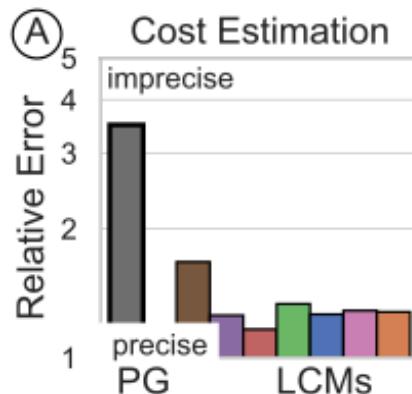


| Why this Tutorial?

*No unified (batch & stream systems)
overview of Learned Cost Models for
query optimizers yet!*

Spoiler: Batch systems

- Trad. approach (PG)
- Flat Vector
- MSCN
- E2E
- Zero-Shot
- QPP-Net
- QueryFormer
- DACE

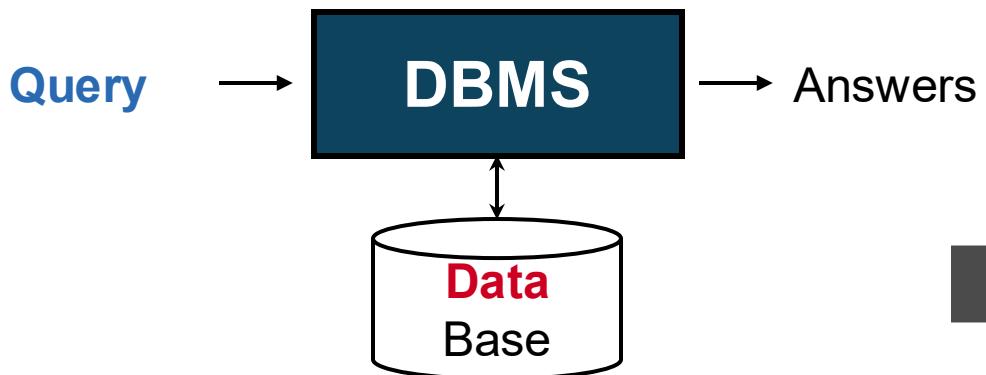


Heinrich R, Luthra M, Wehrstein J, Kornmayer H, Binnig C: How Good are Learned Cost Models, Really?
Insights from Query Optimization Tasks, SIGMOD 2025

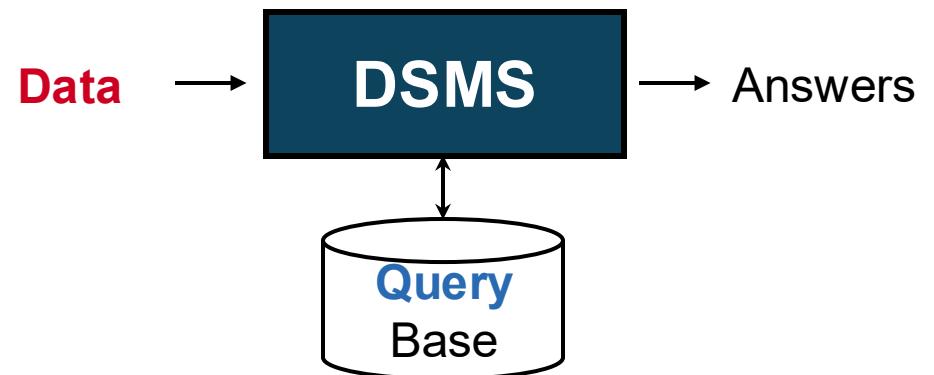
Tutorial: Learned Cost Models for Query Optimization: From Batch to Streaming Systems

What are Batch and Stream Systems?

Database Management System



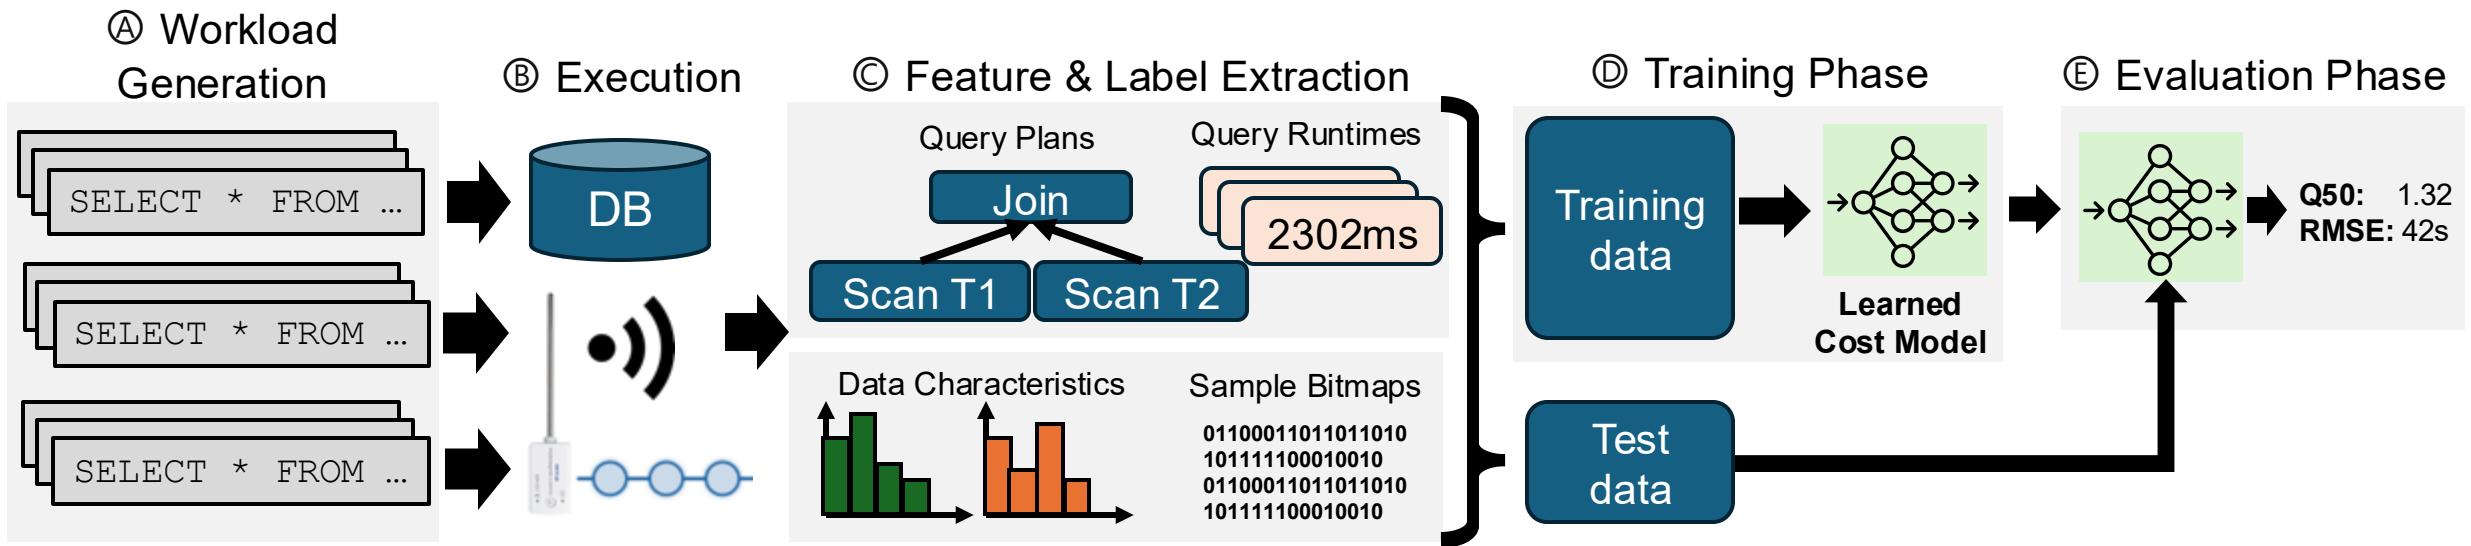
Data Stream Management System



Differences in workloads mean very different requirements for learned cost models

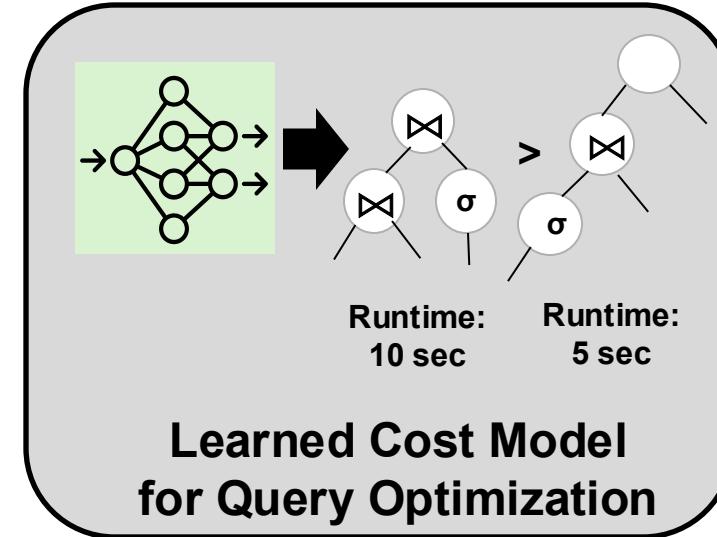
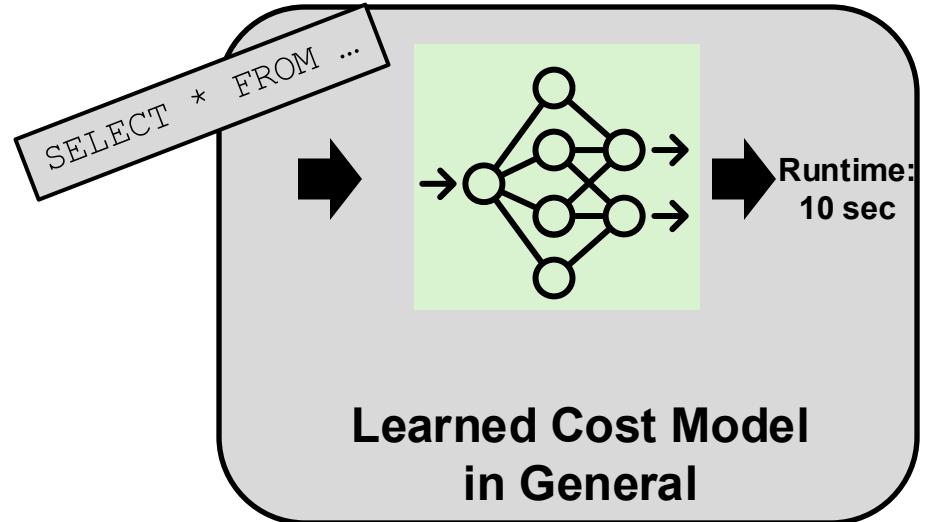
Emergence of Learned Cost Models

- Learned cost models: powerful tool overcoming limitations of traditional cost models
- Key Idea: Instead of relying on hand-crafted analytical models, let data and ML guide the estimation



Heinrich R, Luthra M, Wehrstein J, Kornmayer H, Binnig C: How Good are Learned Cost Models, Really?
Insights from Query Optimization Tasks, SIGMOD 2025

Our Tutorial: Cost Models (in) Query Optimizers



For both Batch and Streaming Systems

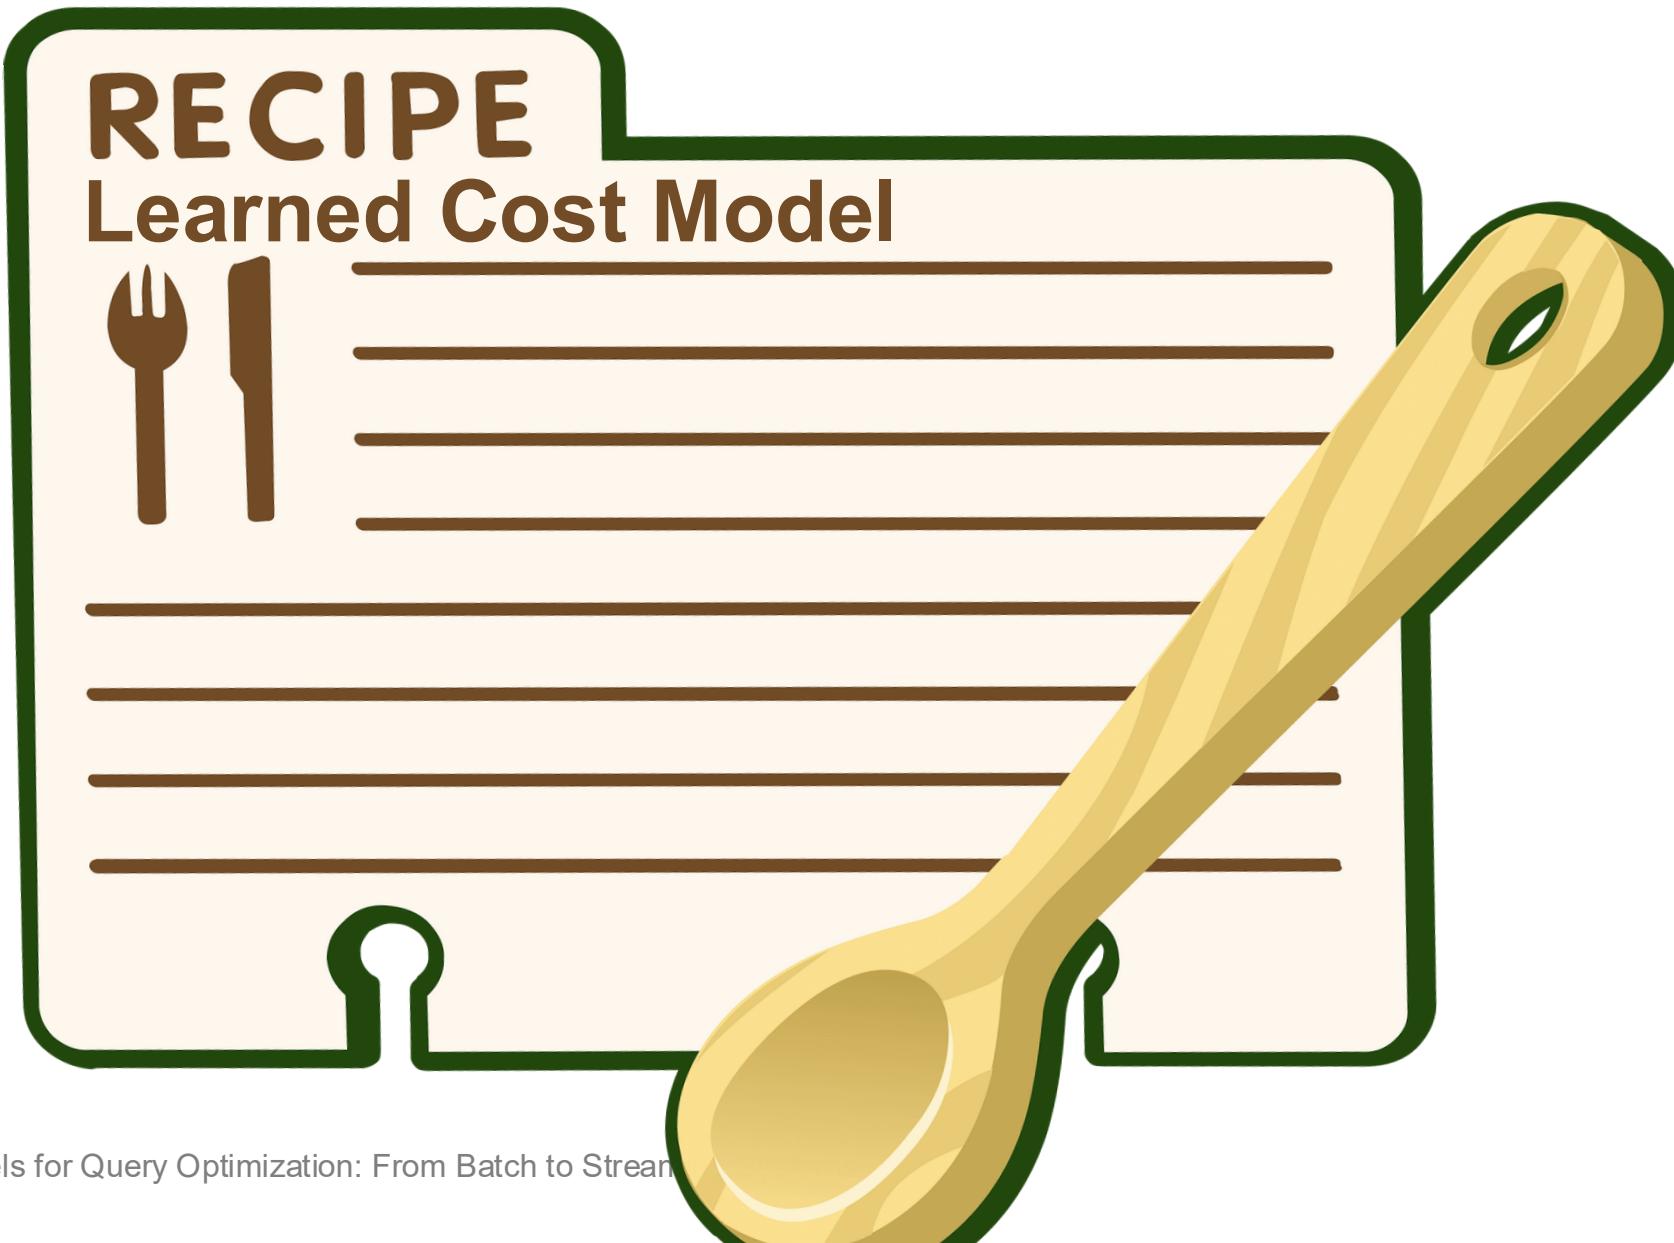
Agenda

- LCMs in Batch Systems
- LCMs in Streaming Systems
- Road Ahead

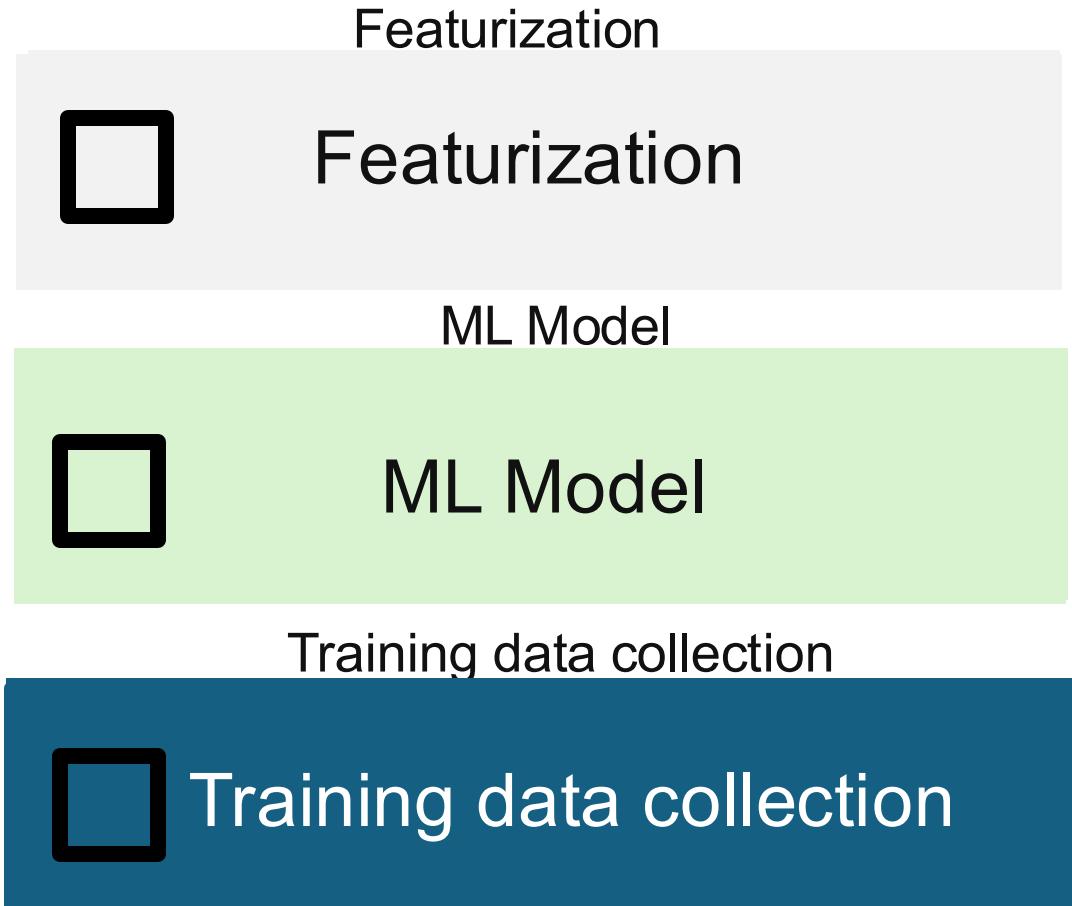
Agenda

- **LCMs in Batch Systems**
- LCMs in Streaming Systems
- Road Ahead

| What's the ingredients of Learned Cost Models?



Learned Cost Models Ingredients



What features can we build?

Featurization

- Query Encoding
- Plan Encoding
- Cardinality/Cost Estimates
- DB Statistics

Query encoding

Featurization

- Given a query SQL statement, convert it into a feature vector, e.g.,

```
SELECT * FROM A, B, C  
WHERE A.a = B.a AND B.b = C.b AND A.a < 51  
GroupBy C.c OrderBy A.a;
```



```
[0, 0, 1, ..., 1, 0, 0.25, 0.73, 1, 0, 0, 0, .....]
```

Elements in a query

Featurization

```
SELECT * FROM A, B, C  
WHERE A.a = B.a AND B.b = C.b AND A.a < 51  
GroupBy C.c OrderBy A.a;
```

- Major elements in a query to consider
 - tables, columns, predicates, joins, aggregator (group by) /sorter (order by)
- Most of these elements are categorical variables
 - one-hot
 - multi-hot
 - learnable embedding

Encode elements of a query

Featurization

```
SELECT * FROM A, B, C  
WHERE A.a = B.a AND B.b = C.b AND A.a < 51  
GroupBy C.c OrderBy A.a;
```

Tables/columns: one-hot/multi-hot encoding

- table set: [A, B, C, D] => table “A”: [1, 0, 0, 0]
- column set: [A.a, A.b, A.c, B.a, ...] => column “A.a”: [1, 0, 0, 0, ...]
- if multiple tables/columns are involved, multi-hot is used
 - [A, B] => [1, 1, 0, 0]
 - [A.a, A.c] => [1, 0, 1, 0, 0, 0, ...]

Encode elements of a query

Featurization

```
SELECT * FROM A, B, C  
WHERE A.a = B.a AND B.b = C.b AND A.a < 51  
GroupBy C.c OrderBy A.a;
```

Predicates (*<column, predicate operator, value> triplet*): e.g., “A.a < 51”

- concatenate: one-hot column + one-hot operator + value, or
- directly one-hot encode the existence of a column predicate

A.a	A.b	A.c	B.a	...
1	0	0	0	...

column predicates

- obtain estimated selectivity of column predicates with histogram or bitmap
 - e.g., “[1, 0, 0, 0, ..., 1, ...]” replaced by “[0.55, 0, 0, 0, ..., 0.76, ...]”
- semantic embedding: map a predicate to an embedding vector

Encode elements of a query

Featurization

```
SELECT * FROM A, B, C  
WHERE A.a = B.a AND B.b = C.b AND A.a < 51  
GroupBy C.c OrderBy A.a;
```

- Joins: e.g., “A.a = B.a”
 - directly one-hot encode the joins: [0, 1]
 - concatenate “a”, “A”, and “B”’s representation
 - embedding
- Groupby or Orderby operators
 - boolean indicator: 0/1 => the sql includes the operator or not

```
SELECT * FROM A, B, C  
WHERE A.a = B.a AND B.b = C.b AND A.a < 51  
GroupBy C.c OrderBy A.a;
```

Encode a query

Featurization

Encode a query:

- use element features individually

SELECT COUNT(*) FROM title t, movie_companies mc WHERE t.id = mc.movie_id AND t.production_year > 2010 AND mc.company_id = 5

Table set { [0 1 0 1 ... 0], [0 0 1 0 ... 1] } Join set { [0 0 1 0] } Predicate set { [1 0 0 0 0 1 0 0 0.72], [0 0 0 1 0 0 1 0 0.14] }
table id samples join id column id value operator id

- or concatenate the representation of considered elements

There is no consideration of the query/join graph structure!

Andreas Kipf et al. Learned Cardinalities: Estimating Correlated Joins with Deep Learning. CIDR 2019.

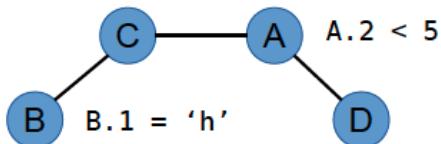
Encode a query

Featurization

Join graph: a table as a node and a join (e.g., $A.a = B.a$) as an edge

Adjacency matrix

```
SELECT * FROM A, B, C, D WHERE  
A.3=C.3 AND A.4=D.4 AND C.5=B.5  
AND A.2<5 AND B.1='h';
```



$\begin{matrix} A.1 & A.2 & \dots & B.1 & B.2 & \dots & E.1 & E.2 \\ [0 & 1 & \dots & 1 & 0 & \dots & 0 & 0] \end{matrix}$

Column Predicates

$[0 \ 1 \ 1 \ 0 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 1 \ \dots \ 1 \ 0 \ \dots \ 0 \ 0]$

Query-level Vector

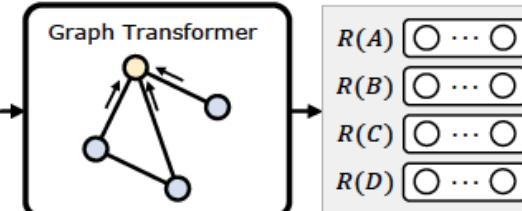
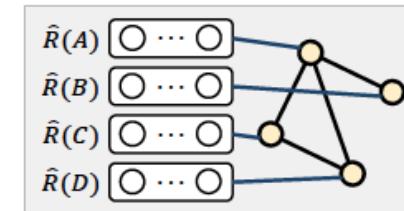
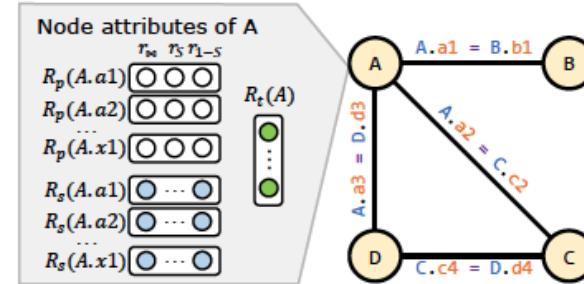
$$\begin{array}{c} \text{A} \ B \ C \ D \ E \\ \left(\begin{array}{ccccc} 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{array} \right) \end{array}$$

Join Graph

Graph embedding

```
select count(*)  
from A, B, C, D  
where A.a1 = B.b1  
and A.a2 = C.c2  
and A.a3 = D.d3  
and C.c4 = D.d4  
and A.x1 >= 100  
and B.x2 <= 3.0  
and C.x3 like "%c%" ;
```

(a) An example of a select-project-join query



Embed the structure info of join graph into the query encoding!

Ryan Marcus et al. NEO: A Learned Query Optimizer. VLDB 2019.

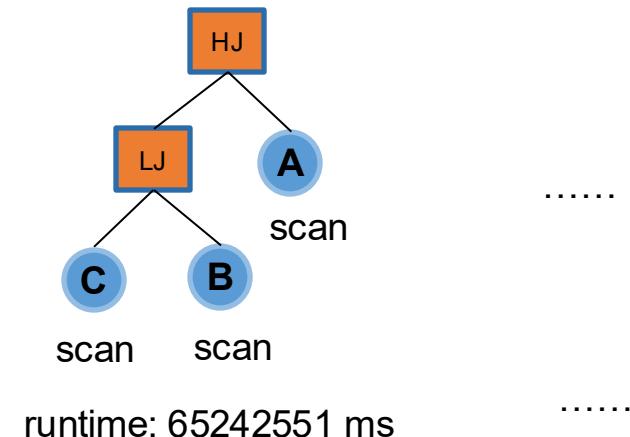
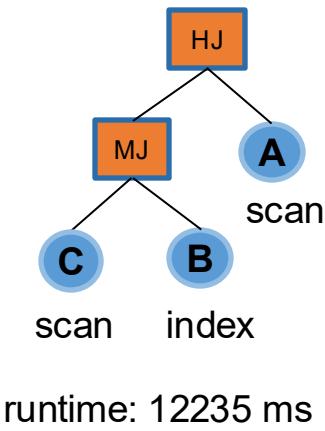
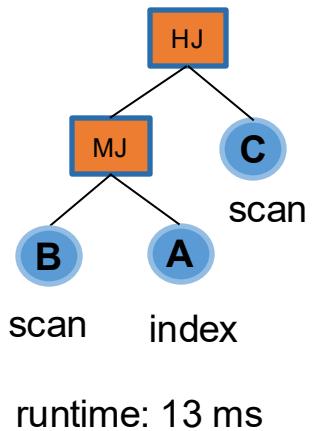
Tianyi Chen et al. LOGER: A Learned Optimizer Towards Generating Efficient and Robust Query Execution Plans. VLDB2023.

Plan encoding

Featurization

```
SELECT * FROM A, B, C  
WHERE A.a = B.a AND B.b = C.b AND A.a < 51  
GroupBy C.c OrderBy A.a;
```

DBMS execute the query as per a tree-structured query plan.



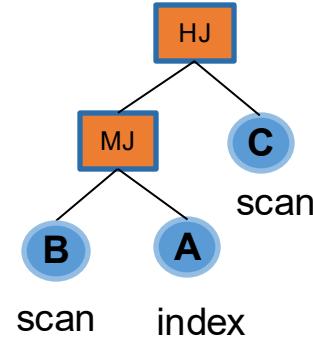
It is the execution plan that determines the cost of executing a query!

How to encode a query plan

Featurization

Major elements in a query plan:

- operators, e.g.,
 - join operator: e.g., hash join(HJ), merge join (MJ), loop join (LJ)
 - data access operator: e.g., index scan (index), seq scan (scan)
 - aggregate (hash/stream)
- tables and columns to be joined
- join order: embodied by the tree structure in a bottom-up way



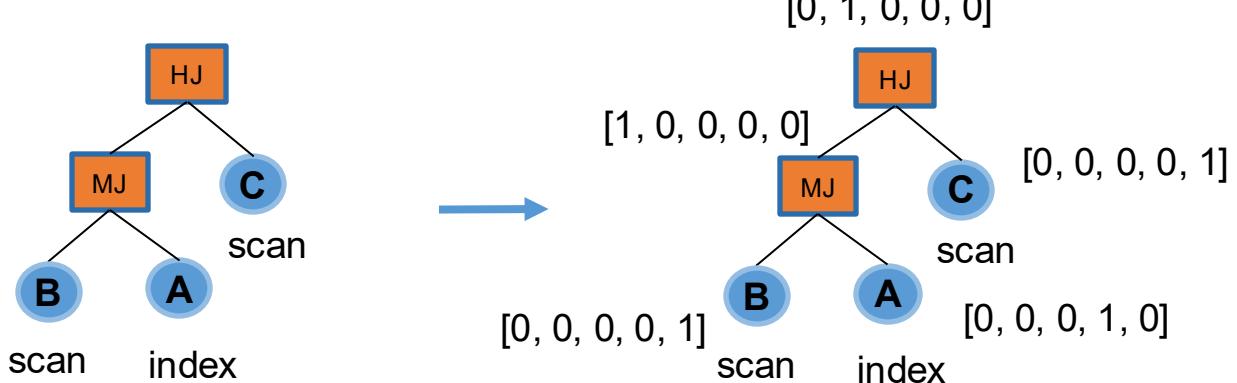
How to encode a query plan

Featurization

Procedures to encode a query plan:

1) encode the node:

- one-hot encode the operators
- aggregate other node features such as tables/columns (discussed before)



How to encode a query plan

Featurization

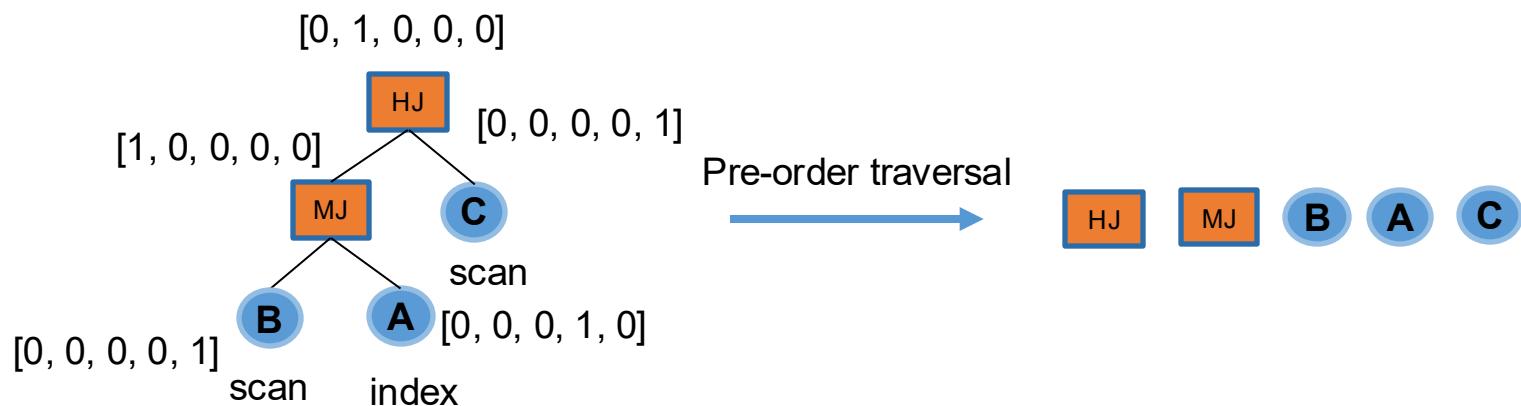
2) exploit the structure info (multiple choices).

- flatten it into a **vector** with tree traversing algorithm, e.g., depth first searching
- preserve the **tree** shape for the subsequent tree-based NN
- extend it to a more informative structure such as a **graph**

Examples of plan encoding

Featurization

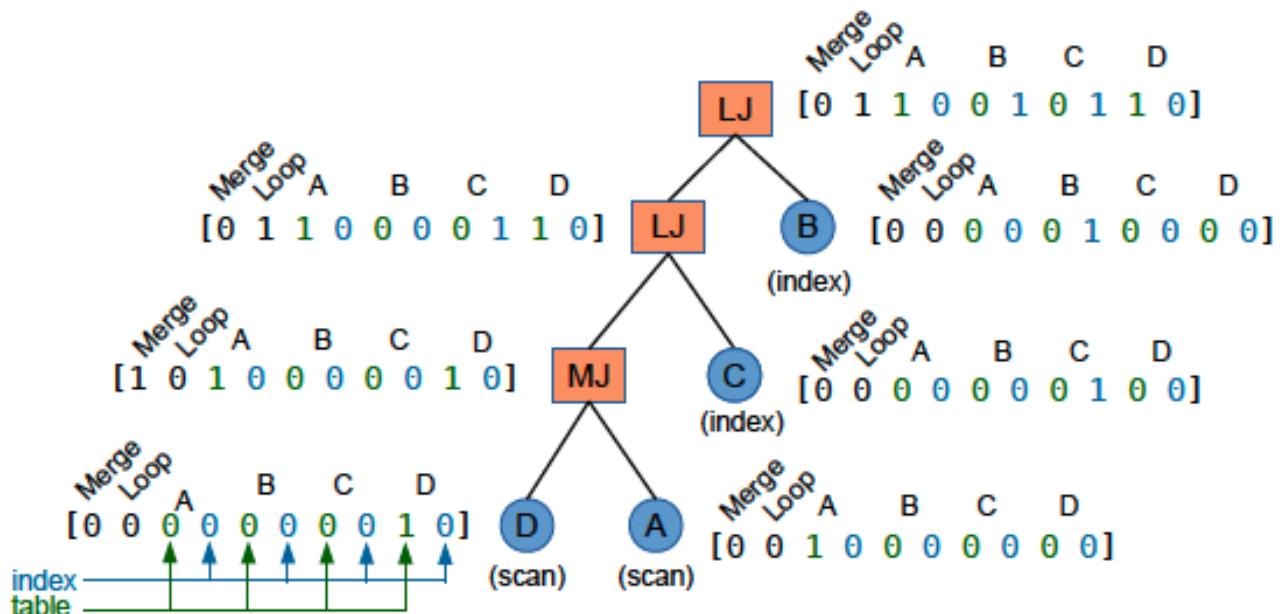
Flat vector



Examples of plan encoding

Featurization

Vector tree



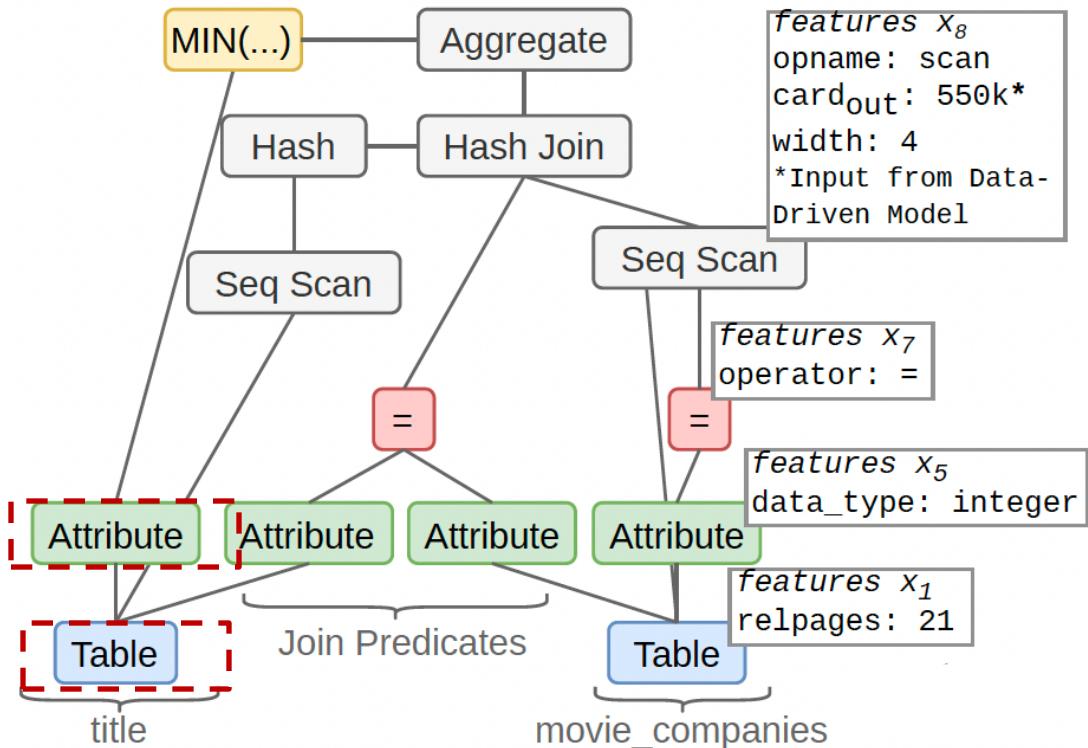
Keep the tree structure and take it into the subsequent tree neural network

Ryan Marcus et al. NEO: A Learned Query Optimizer. VLDB 2019.

Examples of plan encoding

Featurization

Graph

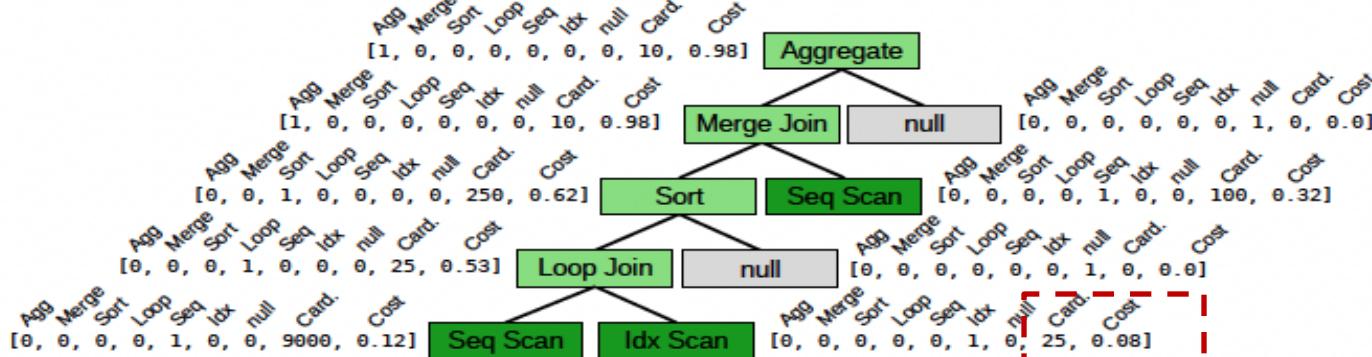


Besides the tree nodes, there are more nodes like attributes and tables.

Cardinality & cost estimates

Featurization

- Obtain these estimates from traditional cost-based optimizer
- Incorporate these estimates into the plan encoding



Bao takes cardinality and cost estimates into plan encoding

Ryan Marcus et al. Bao: Learning to Steer Query Optimizers. SIGMOD 2021.

DB statistics

Featurization

- Database statistics: histograms and sample bitmaps
 - combine their usage with table/column encoding or predicate encoding
- Other statistical features from DB
 - the number of rows in a table
 - the number of unique values in a column
 - the number of null values in a column
 -

Notes on featurization

Featurization

Transferability in the features:

- DB-specific features such as table/column name/identifier which cannot transfer across DBs
- transferable features such as card./cost estimates and statistics

Feature selection:

- some features may be redundant such as predicate encoding and cardinality estimates
- query encoding only / plan encoding only / both encoding

Plan encoding is often there!

Learned Cost Models Ingredients



Featurization



ML Model



Training data collection

What kinds of models can we build?

ML Model

- Regression task: the target is runtime (more common)
- The architecture of cost models are related to the shape of input features
- Cost model for flat features:
 - FlatVector [1] leverage a regression tree model
 - An multi-layer perceptron (MLP) even can do this task or more competent design, e.g., MSCN [2]

[1] Archana Ganapathi et al. *Predicting Multiple Metrics for Queries: Better Decisions Enabled by Machine Learning*. ICDE 2009.

[2] Andreas Kipf et al. *Learned Cardinalities: Estimating Correlated Joins with Deep Learning*. CIDR 2019.

Cost models for tree-shape features

ML Model

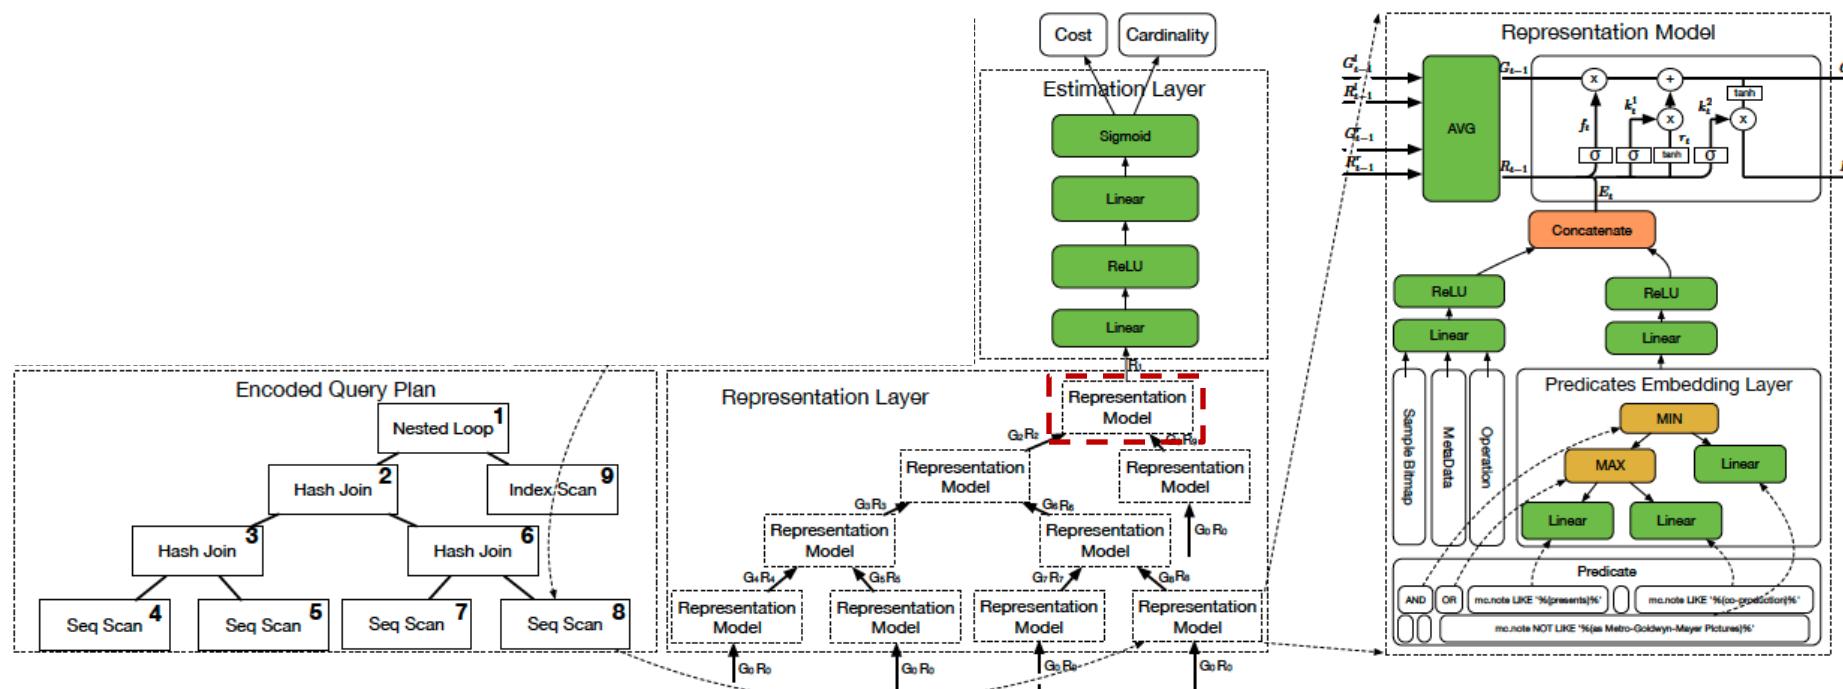
Cost model for tree-shape features:

- the model aims to capture the useful relations in the tree structure of the input
- the **popular** model architectures in LCMs: e.g.,
 - treeLSTM
 - treeCNN
 - tree-based transformer

Cost models for tree-shape features

ML Model

TreeLSTM

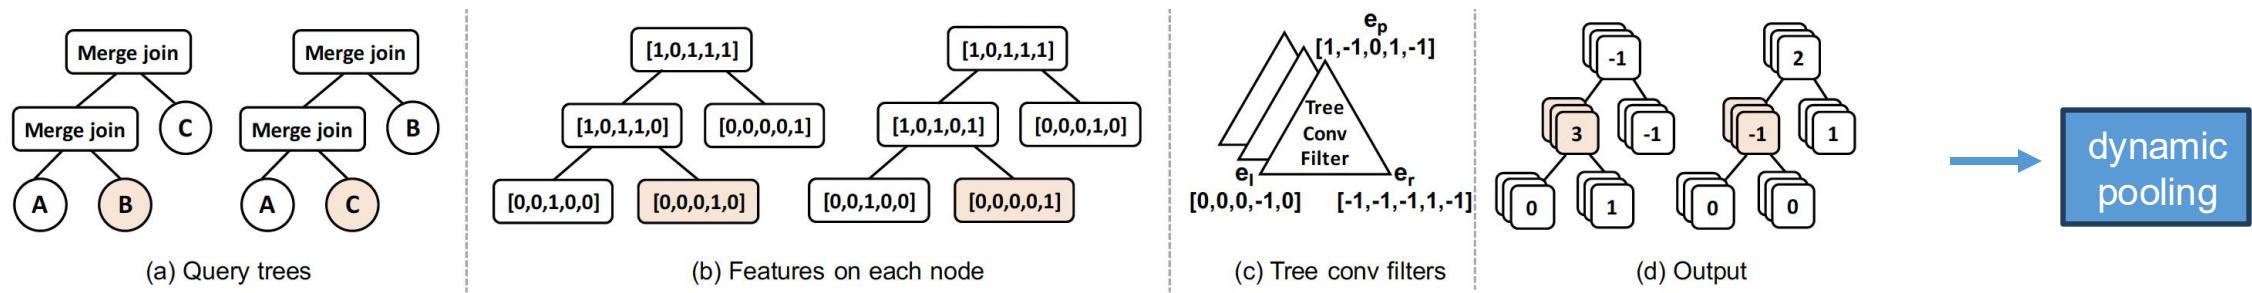


Ji Sun et al. An End-to-End Learning-based Cost Estimator. VLDB 2019.

Cost models for tree-shape features

ML Model

TreeCNN



treeConv filters slide in the plan tree to get a convolved representation of vector tree

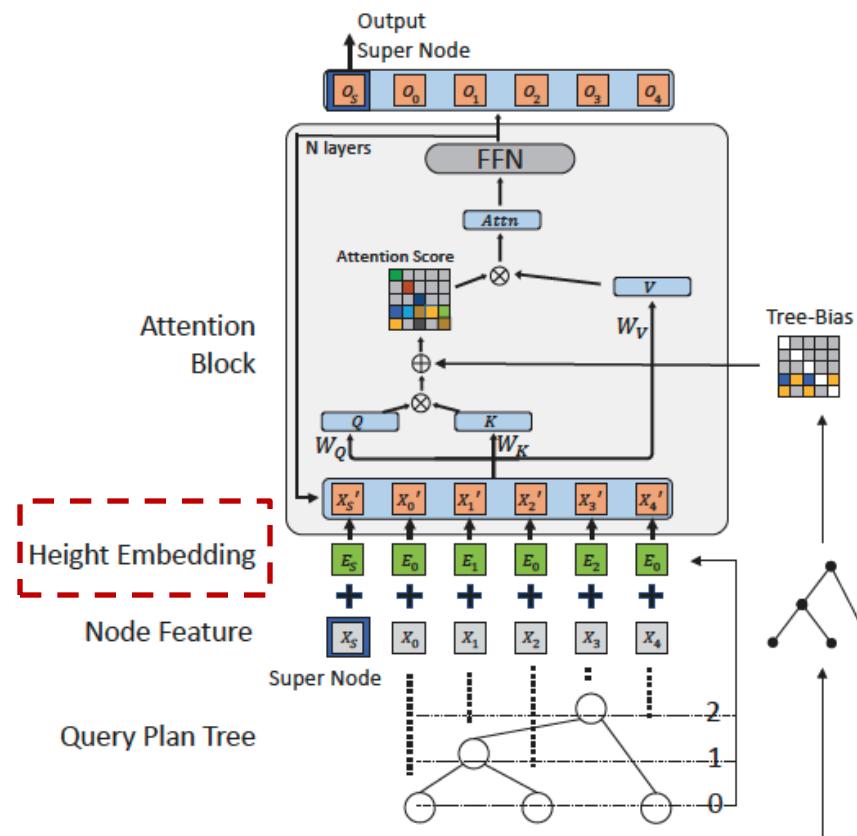
Ryan Marcus et al. NEO: A Learned Query Optimizer. VLDB 2019.

Cost models for tree-shape features

ML Model

Tree-transformer

Height embedding similar to position embedding in transformer records a node's position in a plan tree.

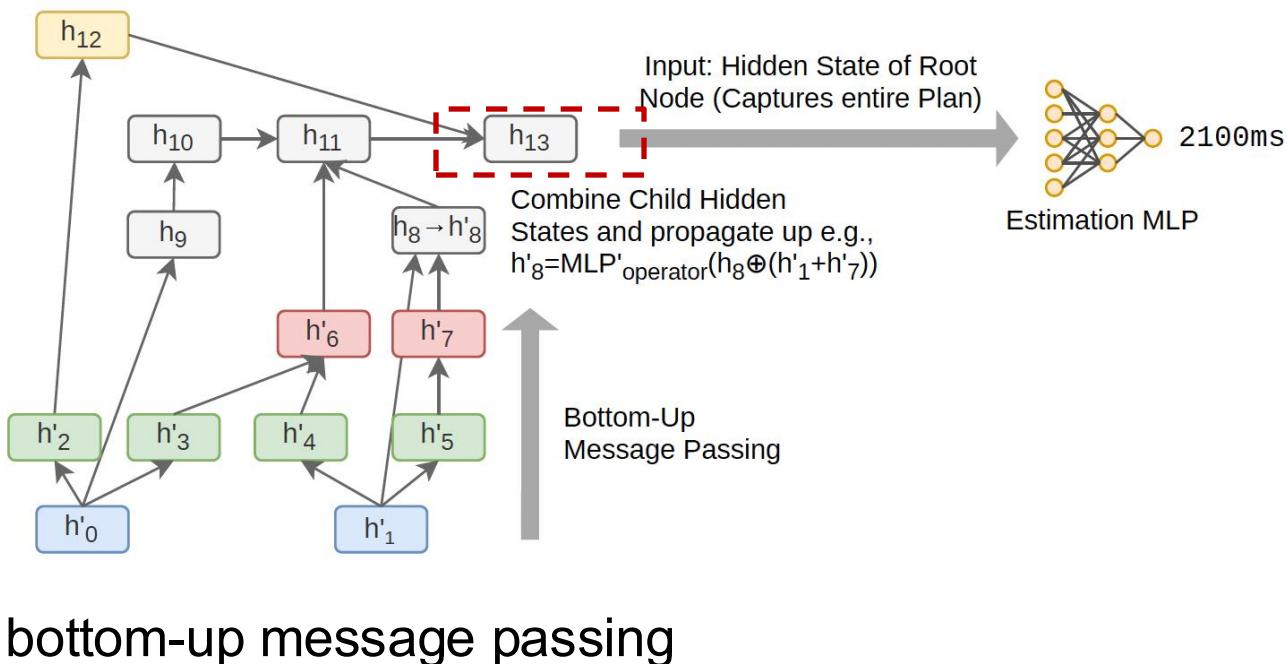


Yu Zhao et al. QueryFormer: A Tree Transformer Model for Query Plan Representation. VLDB 2022.

Cost models for graph-shape features

ML Model

Cost model for graph-structured features: graph neural network (GNN)



Benjamin Hilprecht et al. Zero-Shot Cost Models for Out-of-the-box Learned Cost Prediction. VLDB 2022.

General LCMs vs. LCMs for query optimizer

ML Model

Architecture differences

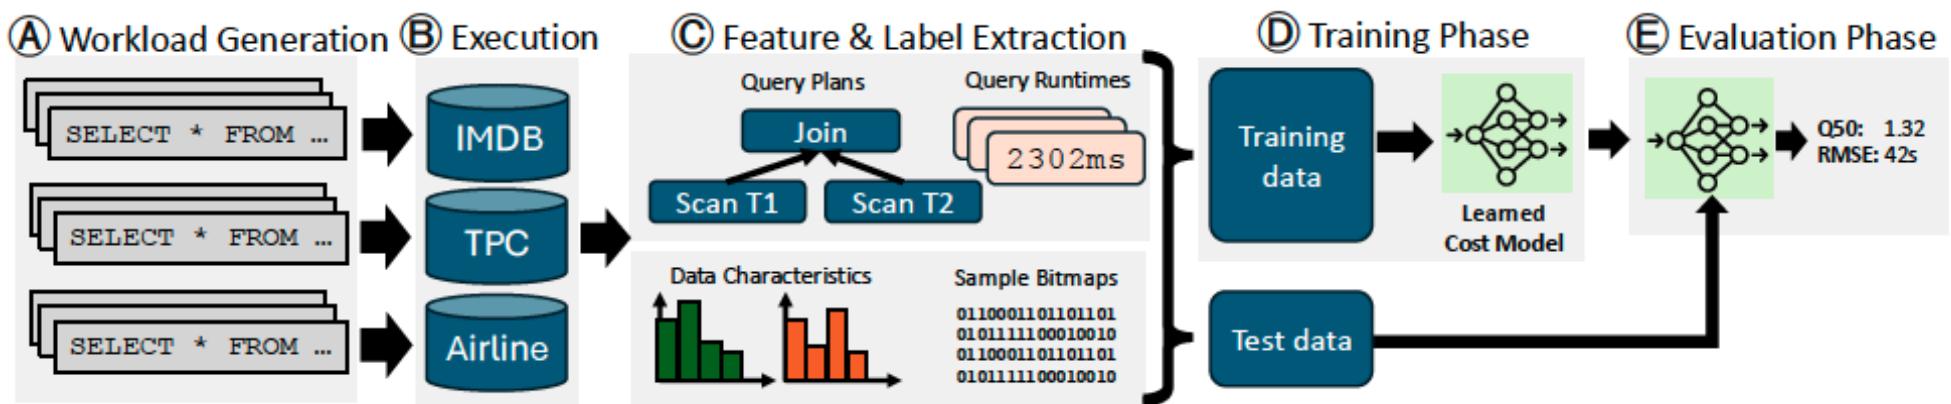
Research Topic	Existing Work	Architecture of Its Learned Cost Model
General LCMs	FlatVector MSCN End-to-End QPP-Net QueryFormer Zero-shot DACE	RegressionTree Deep multisets TreeLSTM TreeNN Transformer GNN Transformer
Learned Query Optimizer	NEO RTOS Bao Balsa HybridQO LEON LOGER	Tree-CNN Tree-LSTM Tree-CNN Tree-CNN Tree-LSTM Tree-CNN Tree-LSTM

General LCMs vary in the model architectures, while LCMs for query optimizers mostly use tree-CNN or tree-LSTM!

General LCMs vs. LCMs for query optimizer

ML Model

- Workflow differences
 - General LCMs:

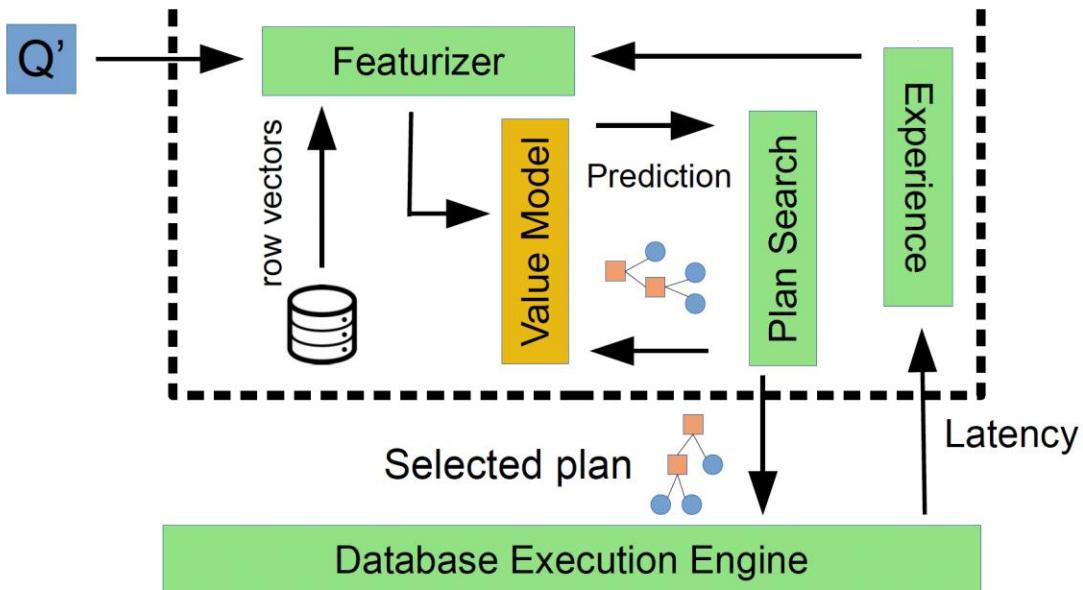


- follow typical “**two-stage**” working flow: i.e., training + testing
- training data and testing data are both **pre-optimized plans** obtained by DBMS
- the input of the model is a **complete plan** which corresponds to an input query

General LCMs vs. LCMs for query optimizer

ML Model

- LCMs for query optimizers

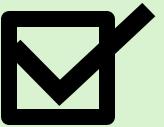


- works as a **value model**, embedded in RL framework
- the input plan to the model may be **not pre-optimized**
- the input plan to the model can be a **subplan**

Learned Cost Models Ingredients



Featurization



ML Model



Training data collection

Training data vital for ML models

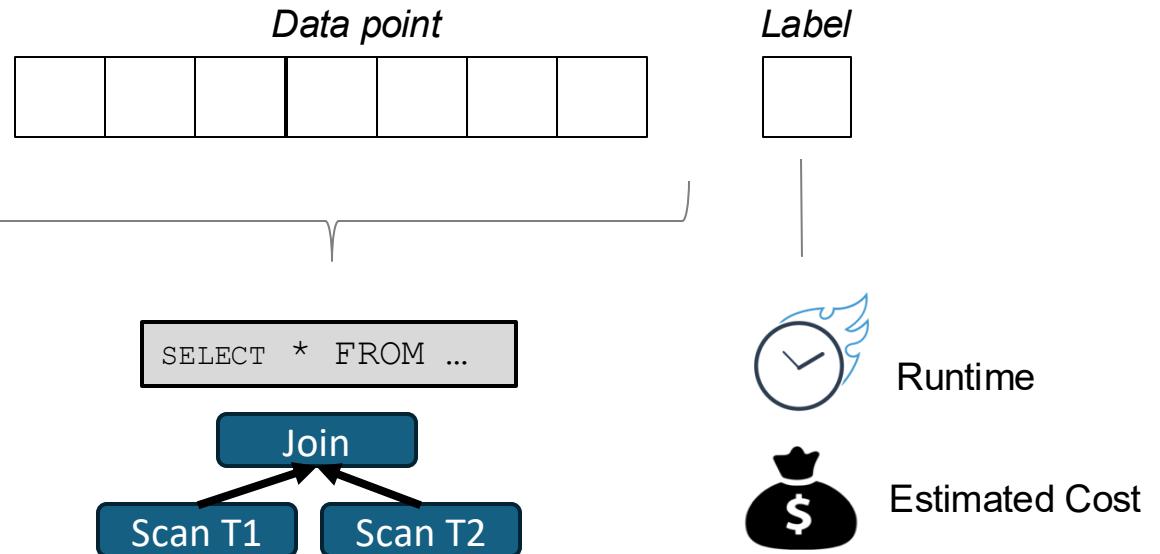
Training data collection

**ML models are
data-hungry**



What is training data in LCMs?

Training data collection

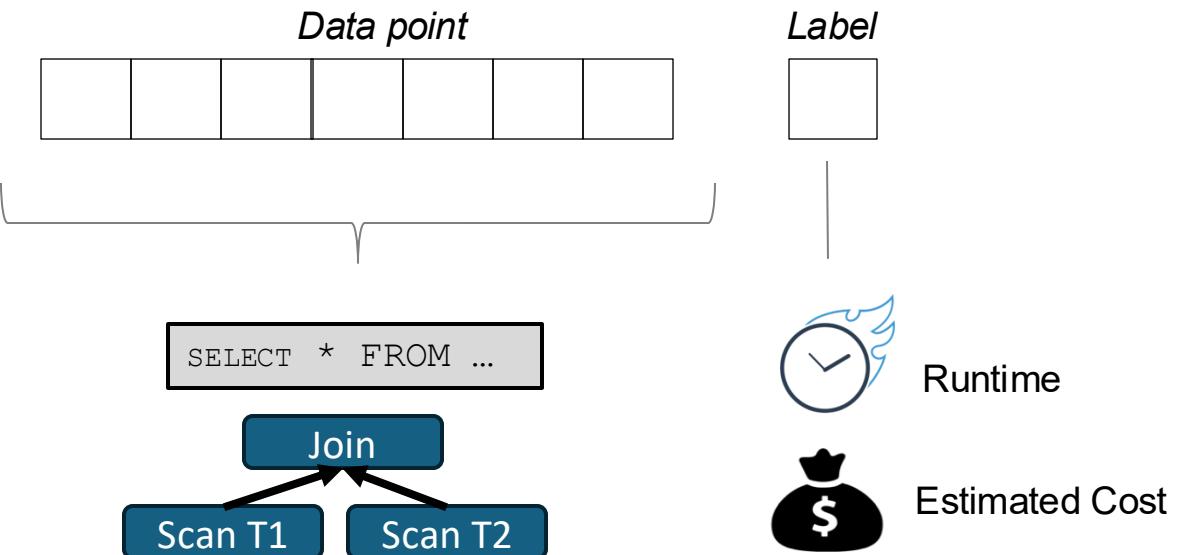


➤ How to find thousands of SQL queries and plans?

➤ How to obtain their label?

What is training data in LCMs?

Training data collection



➤ How to find thousands of SQL queries and plans?

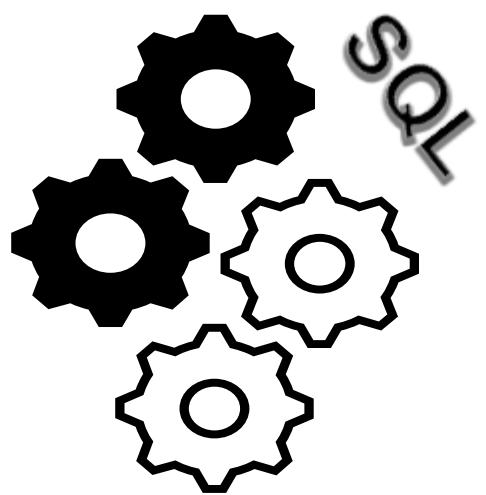
➤ How to obtain their label?

SQL queries

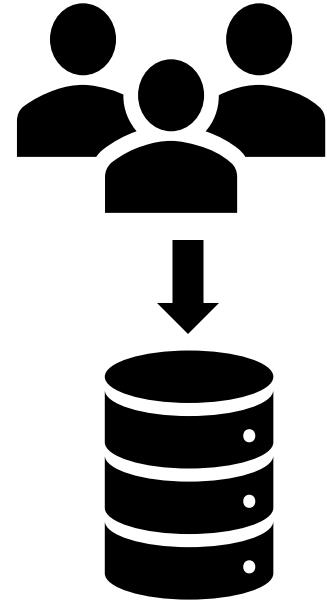
Training data collection



Benchmarks



Synthetic
query generators

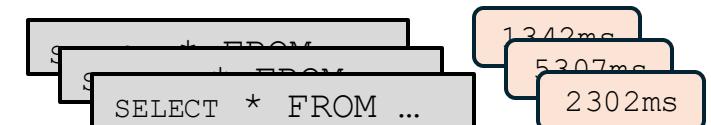
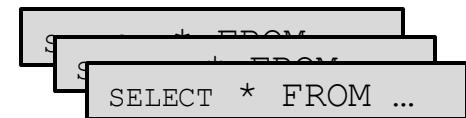
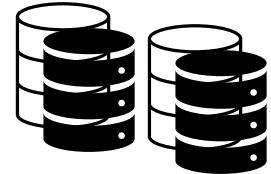


Real user queries

SQL query generation for zero-shot model

Training data collection

- **Data:** 20 databases from real-world datasets & benchmarks
- **Queries:**
 - Benchmark queries
 - Workload generator
 - Standard mode → SPJA queries with conjunctive predicates
 - Complex mode → SPJA with disjunctive complex predicates (e.g., IN)
 - Index mode → random indices in foreign keys and predicate columns
- **Bonus:** Workload traces (queries with runtimes)



SQL query collection in Amazon Redshift

Training data collection

- Training data collected as queries run in **production**
- **Sliding window** (one query in, one out)

Problem

Mostly short-running queries

Catastrophic predictions for long-running queries



Solution

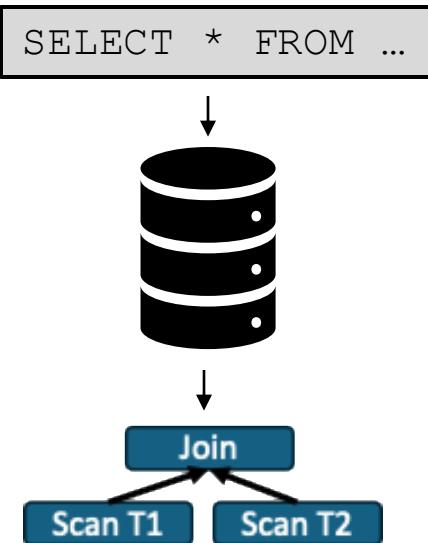
Partition training set into buckets, e.g.:

- o Bucket 1: 0-10 sec
- o Bucket 2: 10-30 sec etc.

Training happens in the production cluster!

Query plans from SQL workloads

Training data collection



SQL query execution

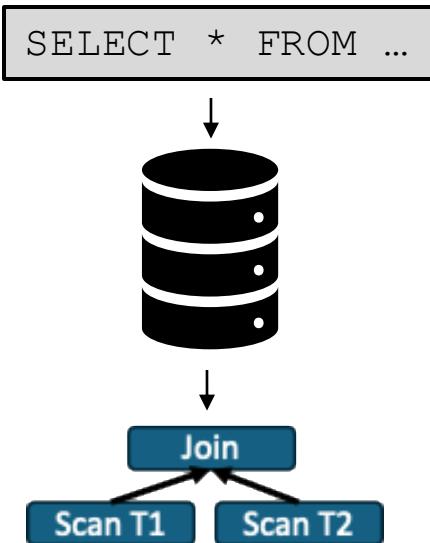
[Balsa] Z. Yang et al. Balsa: Learning a Query Optimizer Without Expert Demonstrations. SIGMOD 2022.

[LTR] H. Behr et al. Learn What Really Matters: A Learning-to-Rank Approach for ML-based Query Optimization. BTW 2023.

[Neo] R. Marcus et al. Neo: a learned query optimizer. PVLDB 2019

[Bao] R. Marcus et al. Bao: Making Learned Query Optimization Practical. SIGMOD 2021.

[LOGER] T. Chen et al. LOGER: A Learned Optimizer Towards Generating Efficient and Robust Query Execution Plans. PVLDB 2023

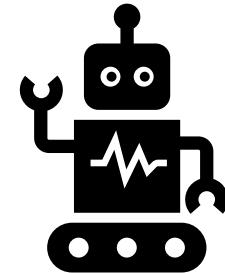
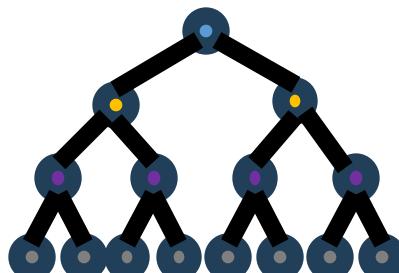


SQL queries with hints

For LCMs in QO only



Dynamic programming Dynamic programming

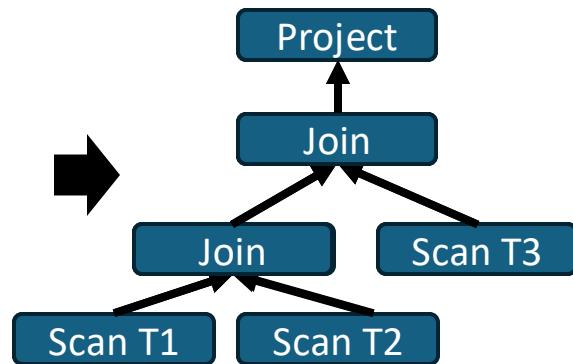


Plan enumeration

Query plan augmentation

Training data collection

Use subplans



Cost=12346

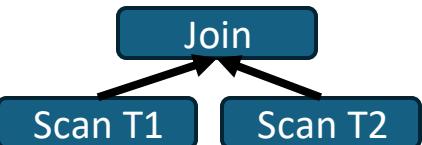


Cost=12346

[Balsa] Z. Yang et al. *Balsa: Learning a Query Optimizer Without Expert Demonstrations*. SIGMOD 2022.

Change cardinalities

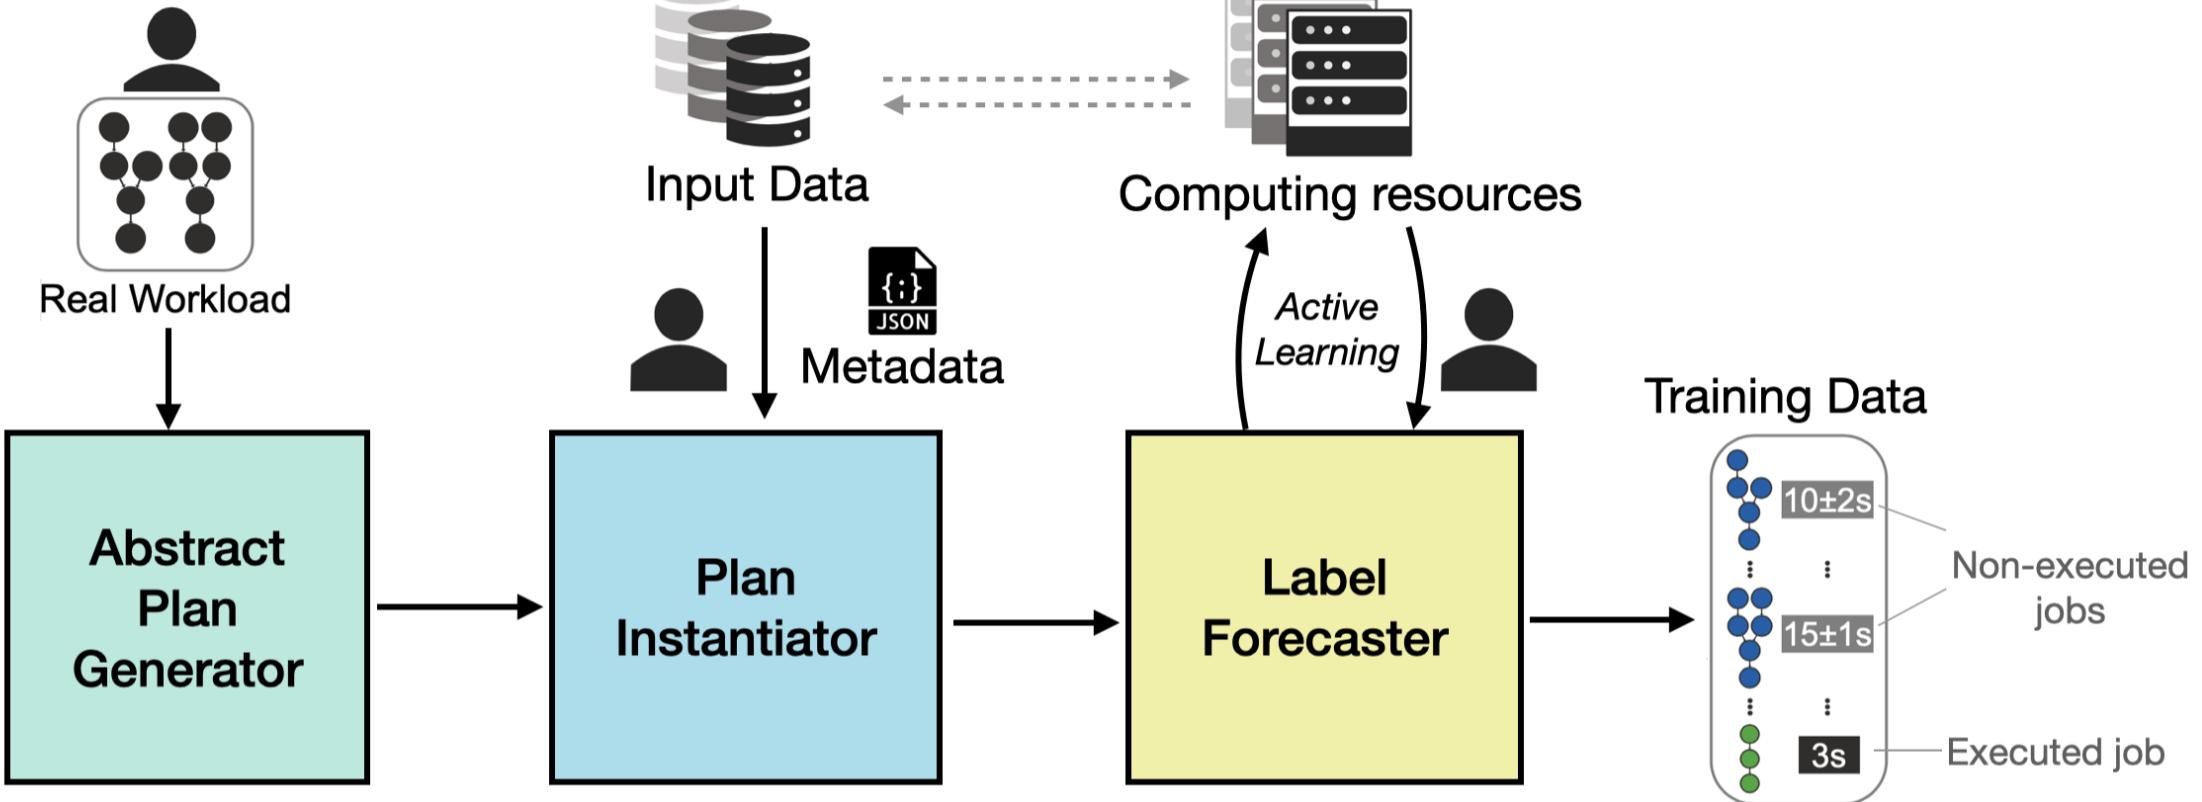
SELECT * FROM ...



R. Zhu et al. *Balsa: Lero: A Learning-to-Rank Query Optimizer*. PVLDB 2023.

Synthetic plan generation with DataFarm

Training data collection



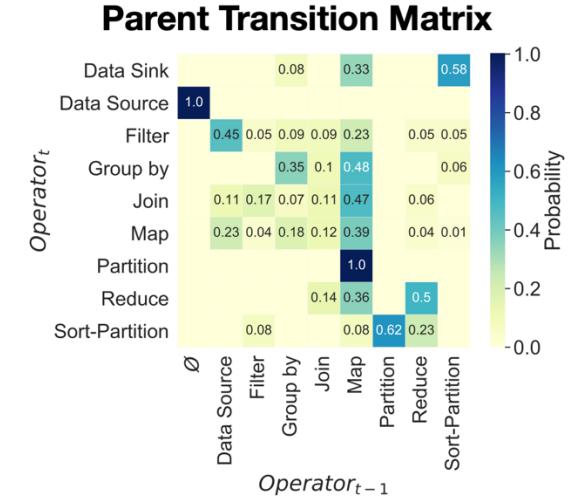
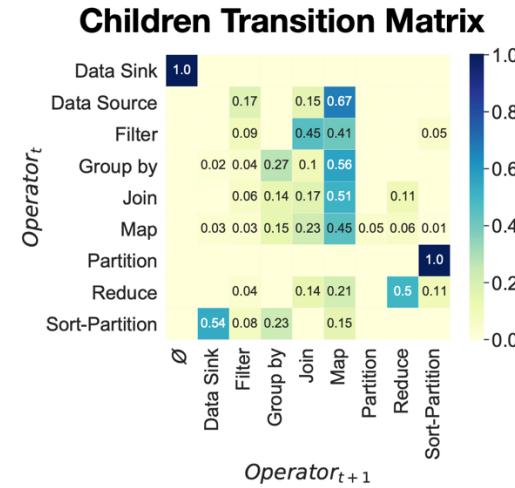
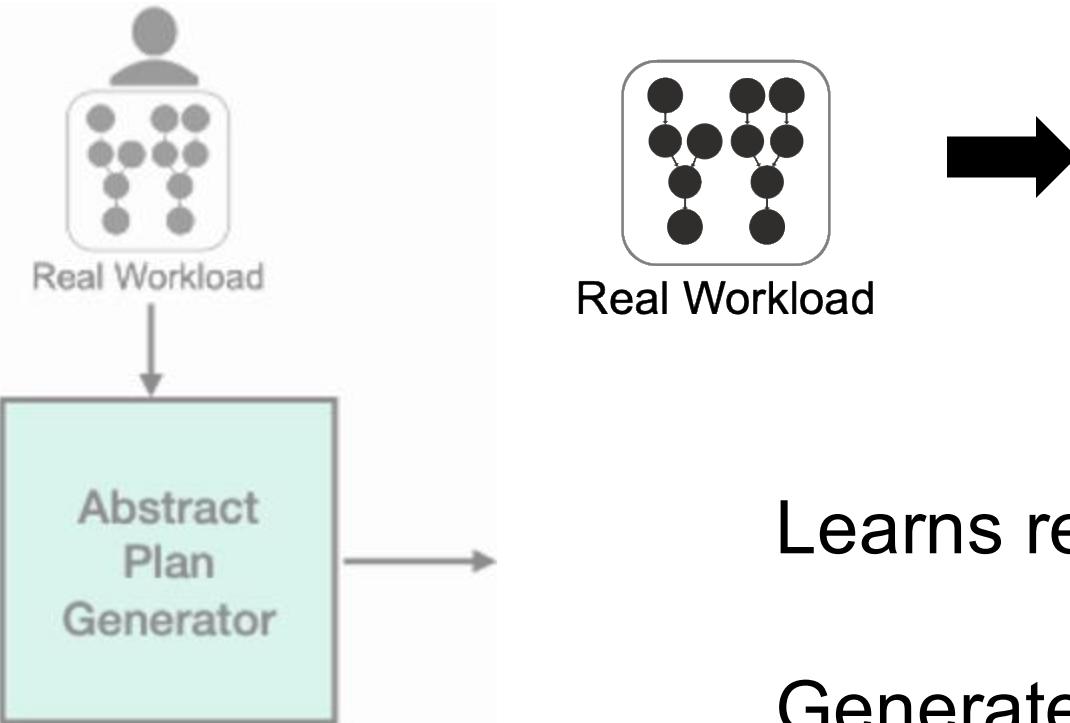
F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.

R. van de Water. DataFarm: Farm Your ML-based Query Optimizer's Food! – Human-Guided Training Data Generation. CIDR 2022

R. van de Water. Farming Your ML-based Query Optimizer's Food. ICDE 2022 (best demo award)

Synthetic plan generation with DataFarm

Training data collection



Learns real execution patterns

Generates new representative plans

F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.

R. van de Water. DataFarm: Farm Your ML-based Query Optimizer's Food! – Human-Guided Training Data Generation. CIDR 2022

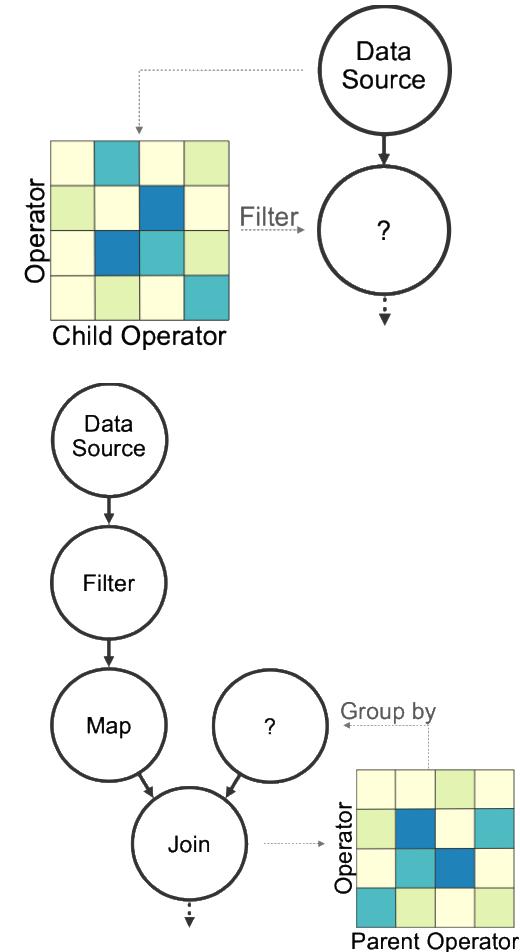
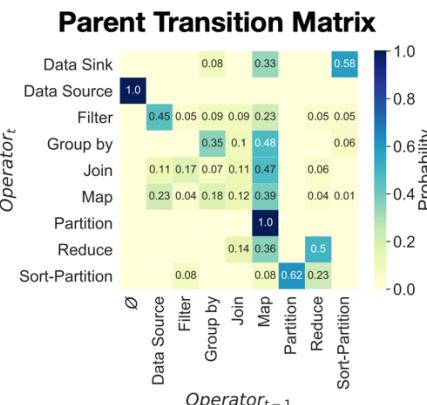
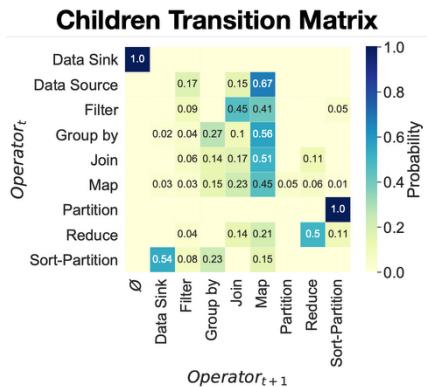
R. van de Water. Farming Your ML-based Query Optimizer's Food. ICDE 2022 (best demo award)

Synthetic plan generation with DataFarm

Training data collection



Real Workload



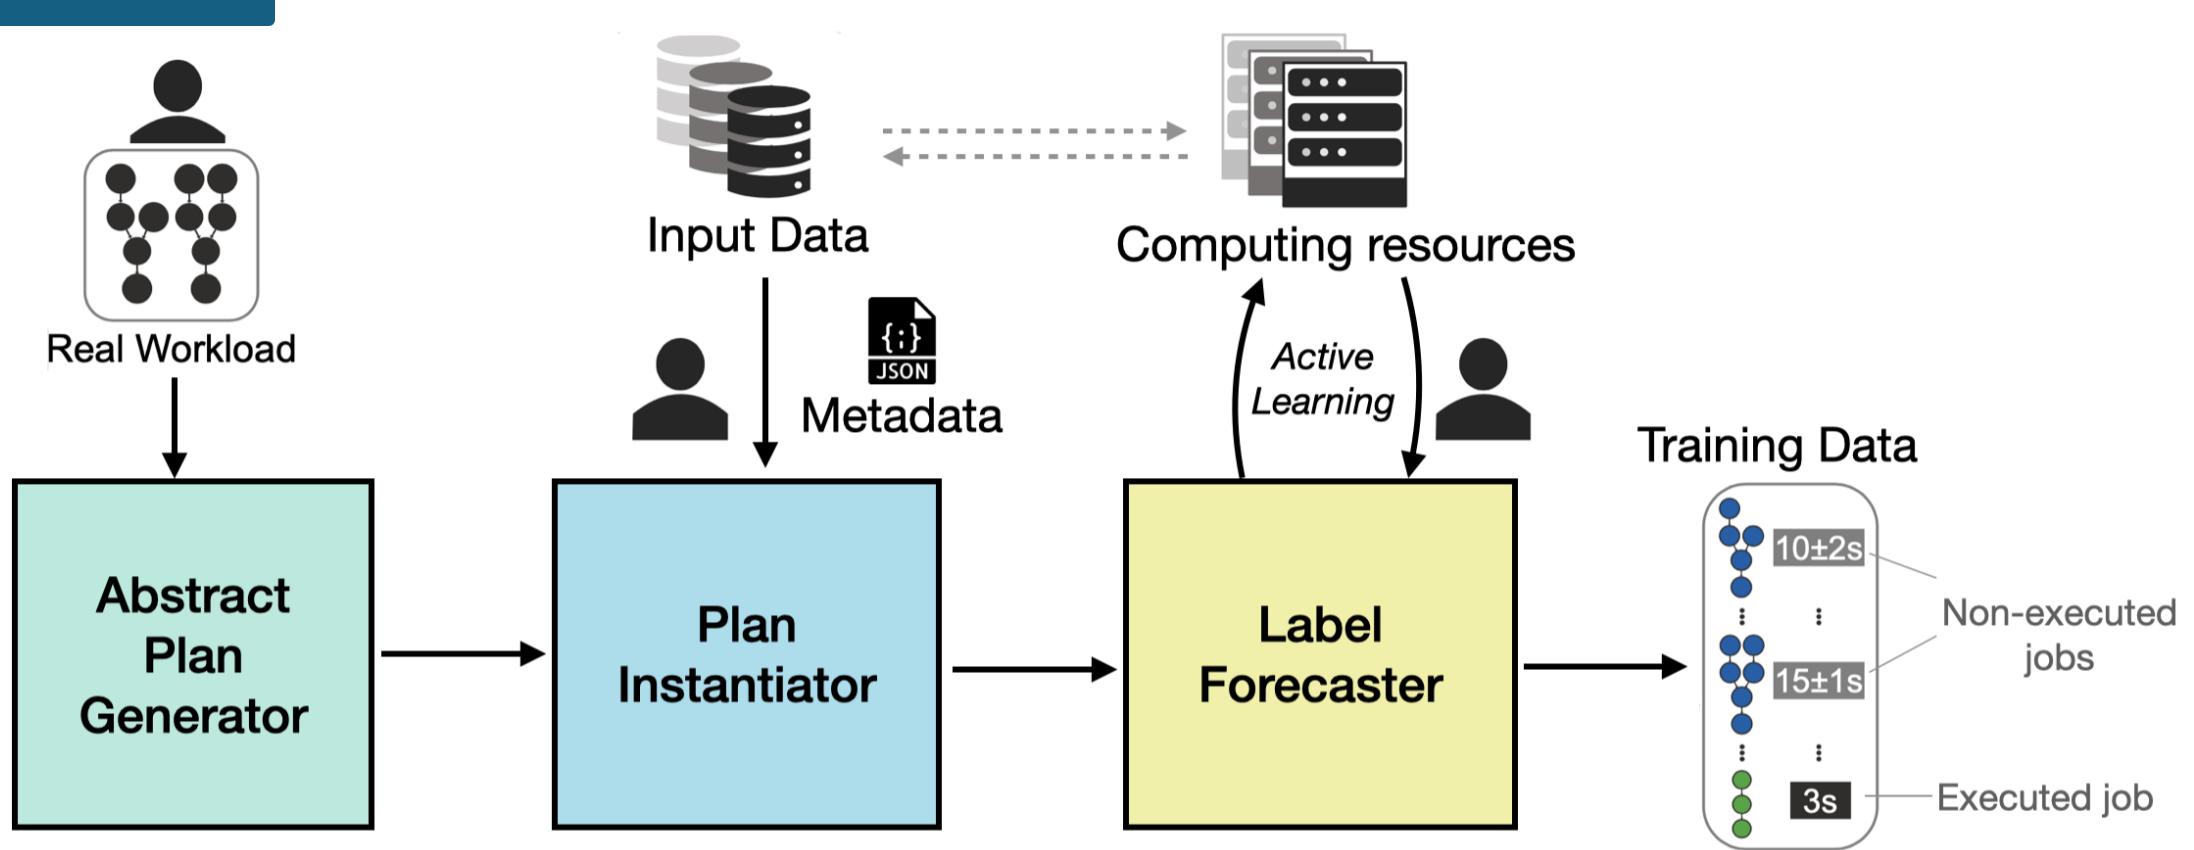
F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.

R. van de Water. DataFarm: Farm Your ML-based Query Optimizer's Food! – Human-Guided Training Data Generation. CIDR 2022

R. van de Water. Farming Your ML-based Query Optimizer's Food. ICDE 2022 (best demo award)

Synthetic plan generation with DataFarm

Training data collection

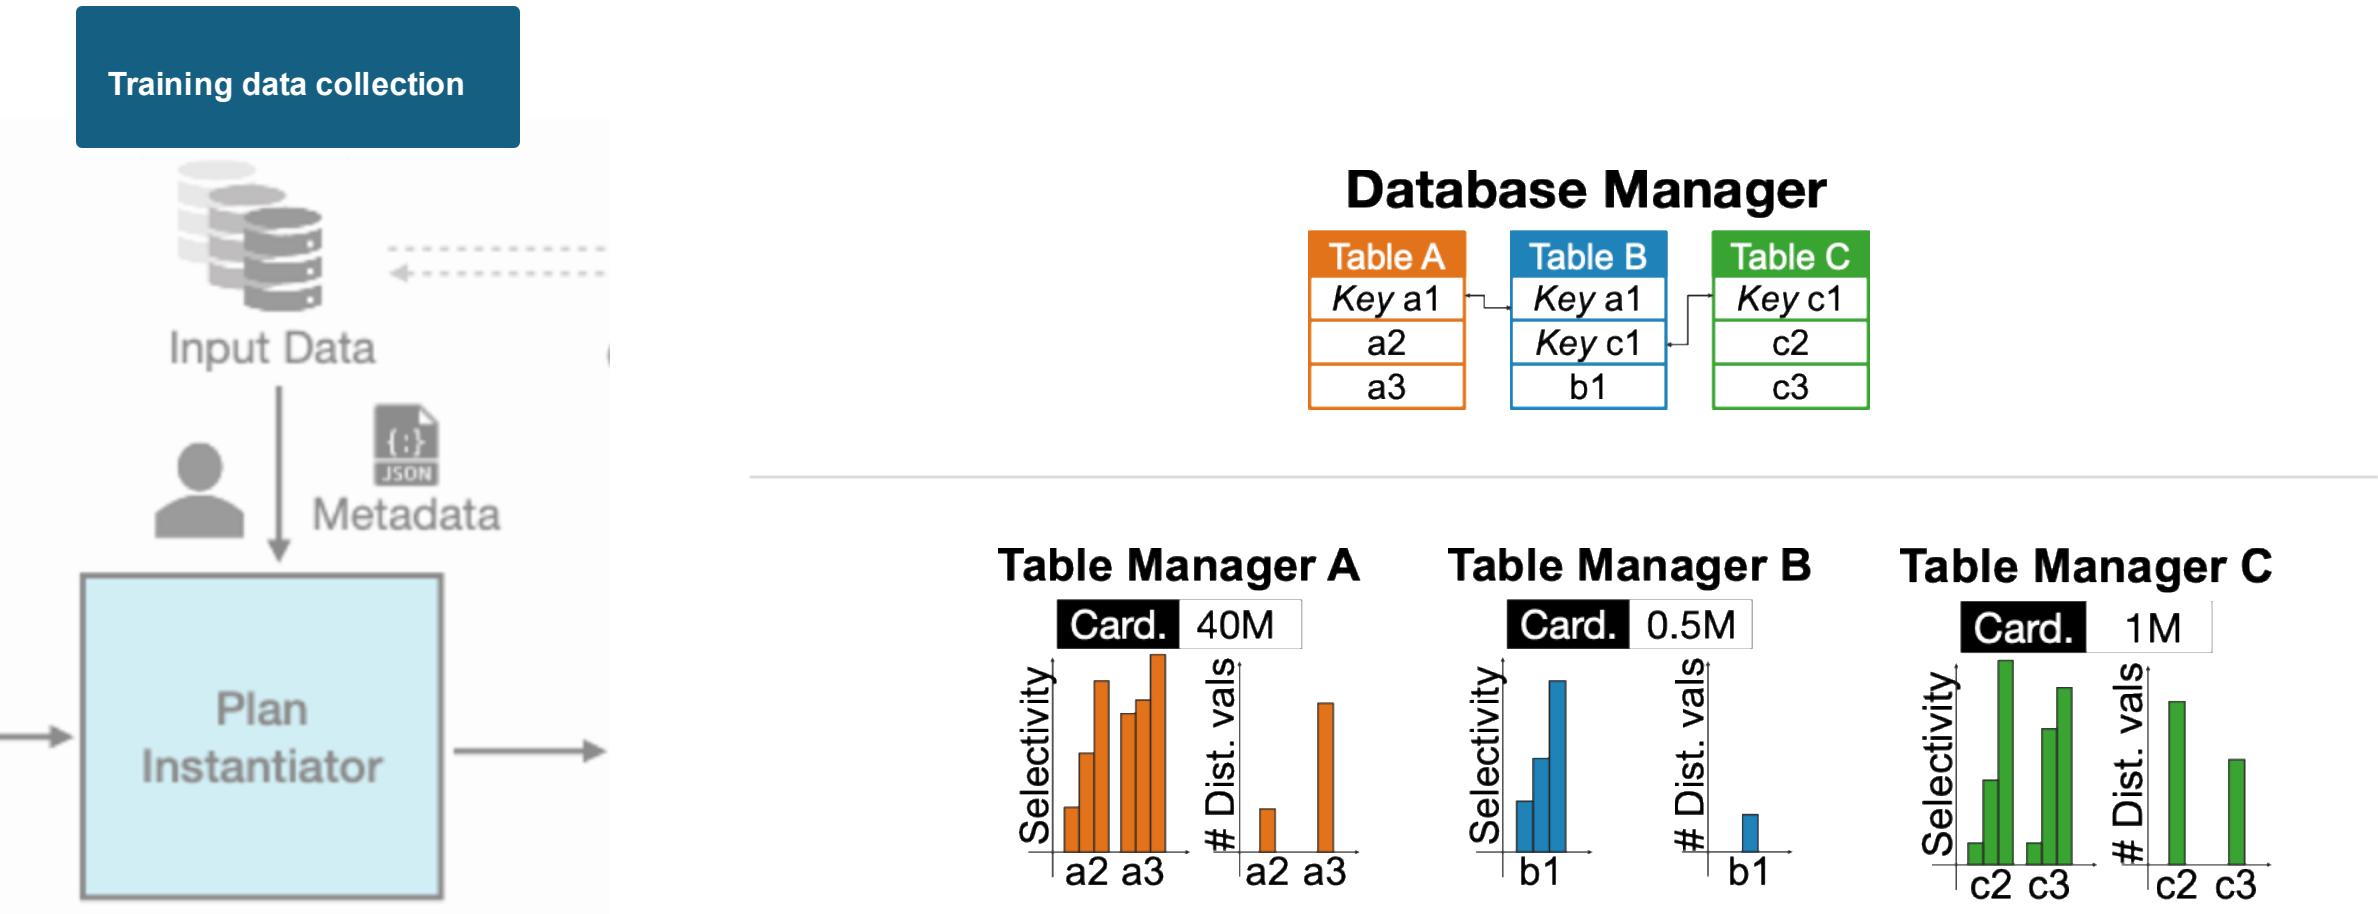


F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.

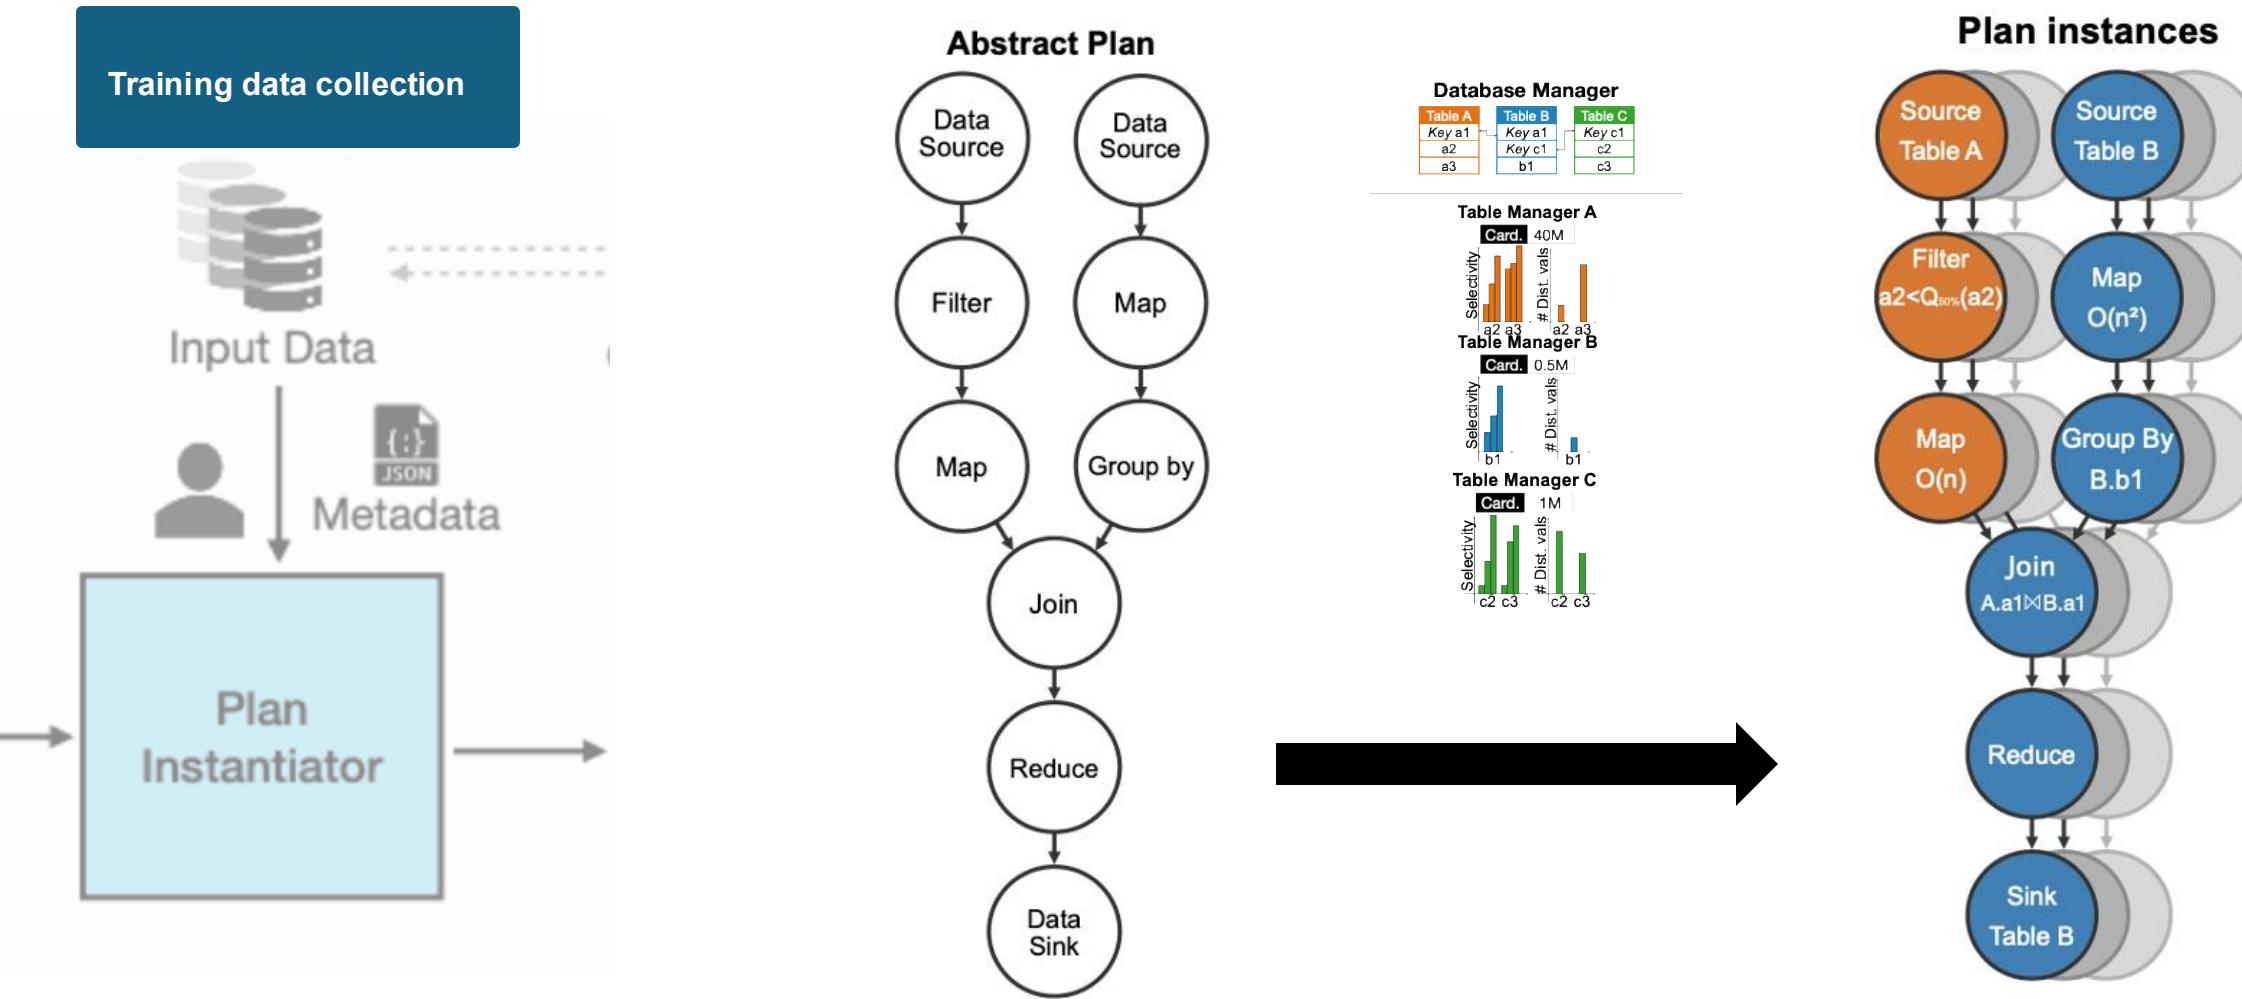
R. van de Water. DataFarm: Farm Your ML-based Query Optimizer's Food! – Human-Guided Training Data Generation. CIDR 2022

R. van de Water. Farming Your ML-based Query Optimizer's Food. ICDE 2022 (best demo award)

Synthetic plan generation with DataFarm



Synthetic plan generation with DataFarm



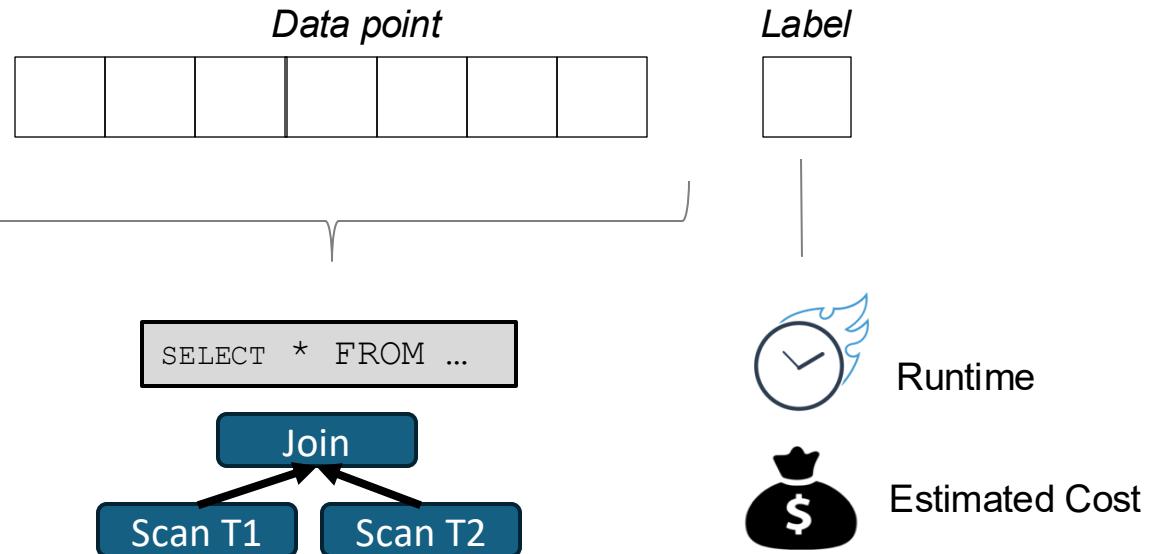
F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.

R. van de Water. DataFarm: Farm Your ML-based Query Optimizer's Food! – Human-Guided Training Data Generation. CIDR 2022

R. van de Water. Farming Your ML-based Query Optimizer's Food. ICDE 2022 (best demo award)

What is training data in LCMs?

Training data collection



➤ Where to find thousands of queries and plans?

➤ How to obtain their label?

Runtime collection

Training data collection

Executing queries can be very time-consuming!



Runtime collection - examples

Training data collection

Zero-shot

3.9GB GB data
300k queries
PostgreSQL

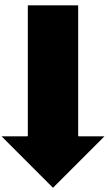


10 days

B. Hilprecht et al. Zero-Shot Cost Models for Out-of-the-box Learned Cost Prediction. VLDB 2022.

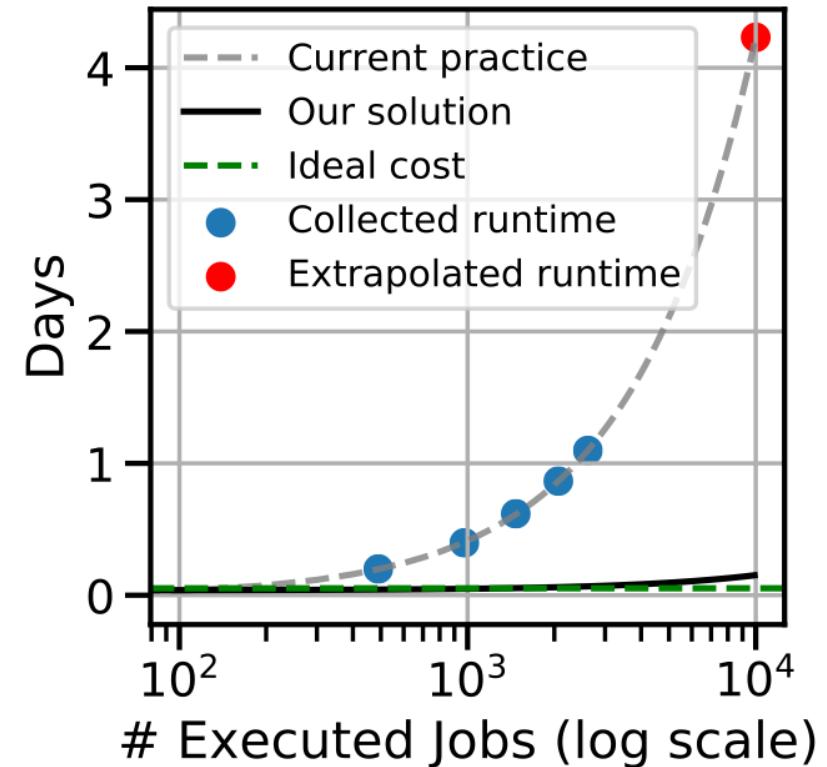
DataFarm

1GB data
10k jobs
Flink



5 days

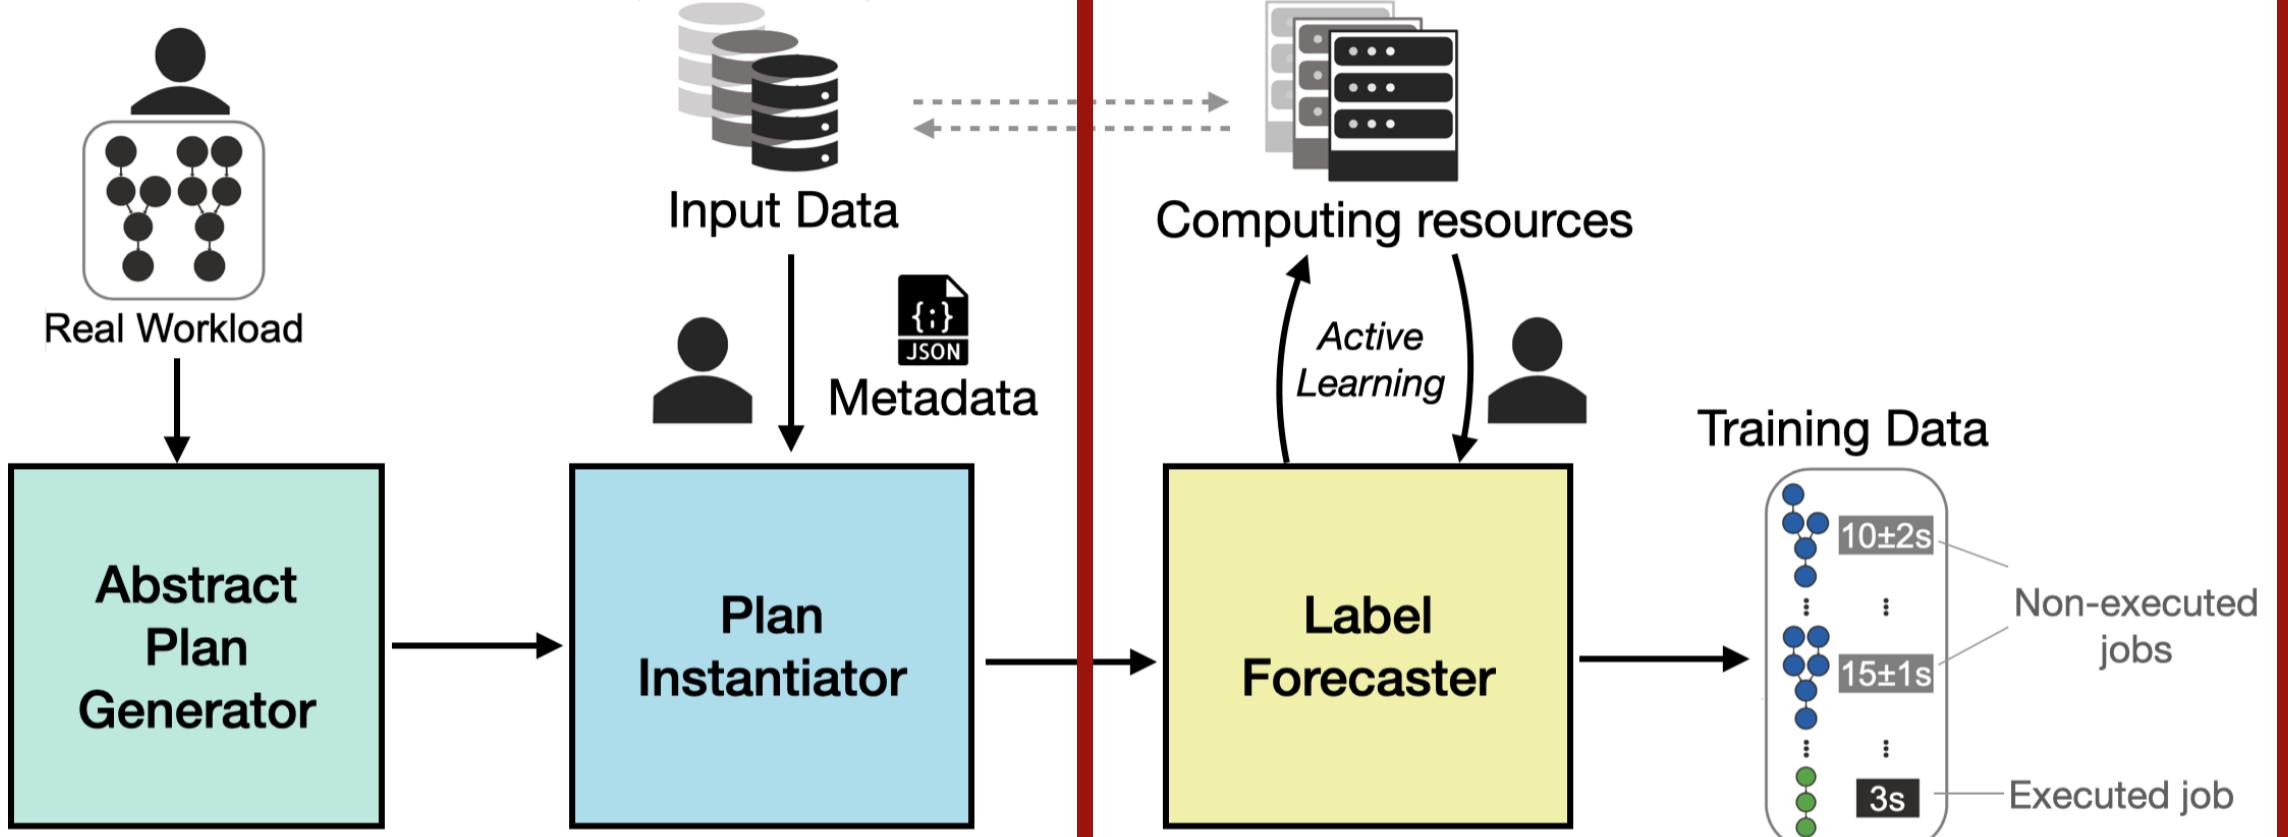
F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.



Extrapolated cost of 10,000 plans with 1TB input data > 6 months*

Label collection in DataFarm

Training data collection

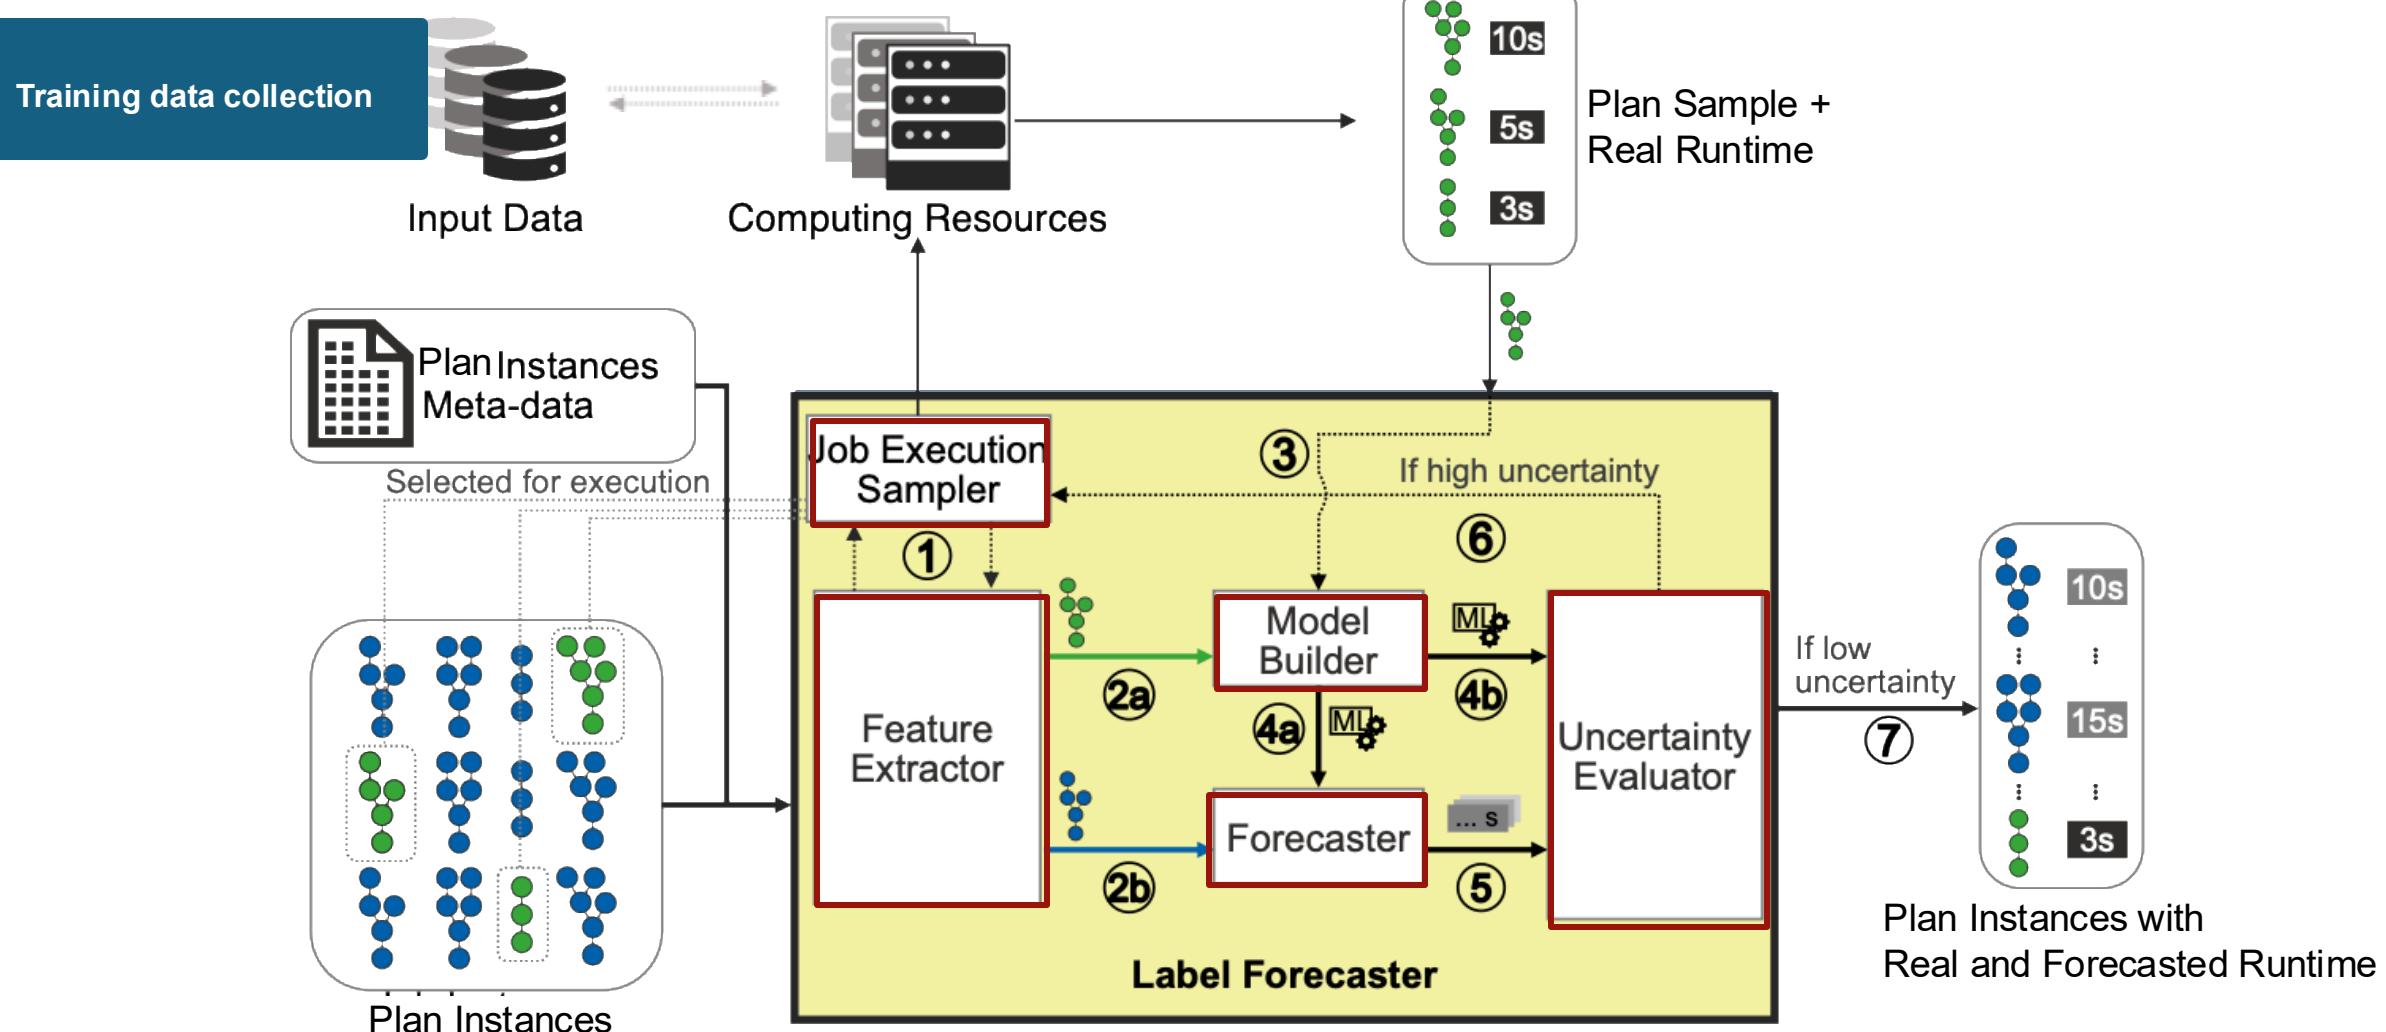


F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.

R. van de Water. DataFarm: Farm Your ML-based Query Optimizer's Food! – Human-Guided Training Data Generation. CIDR 2022

R. van de Water. Farming Your ML-based Query Optimizer's Food. ICDE 2022 (best demo award)

Label collection in DataFarm



F. Ventura et al. Expand your Training Limits! Generating Training Data for ML-based Data Management. SIGMOD 2021.

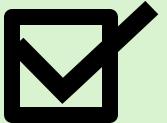
R. van de Water. DataFarm: Farm Your ML-based Query Optimizer's Food! – Human-Guided Training Data Generation. CIDR 2022

R. van de Water. Farming Your ML-based Query Optimizer's Food. ICDE 2022 (best demo award)

Learned Cost Models Ingredients



Featurization



ML Model

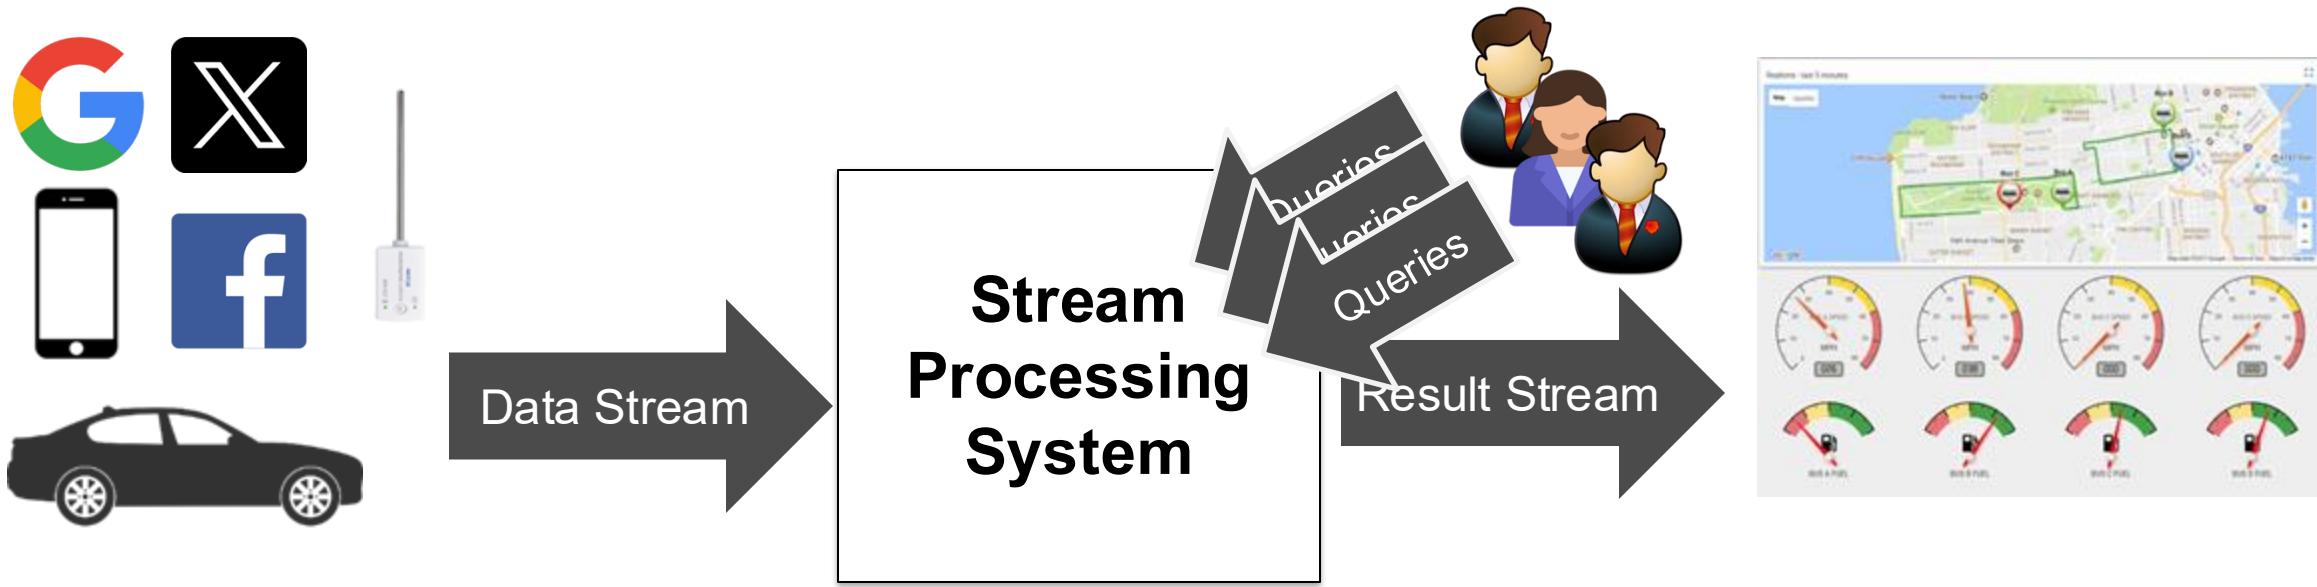


Training data collection

Agenda

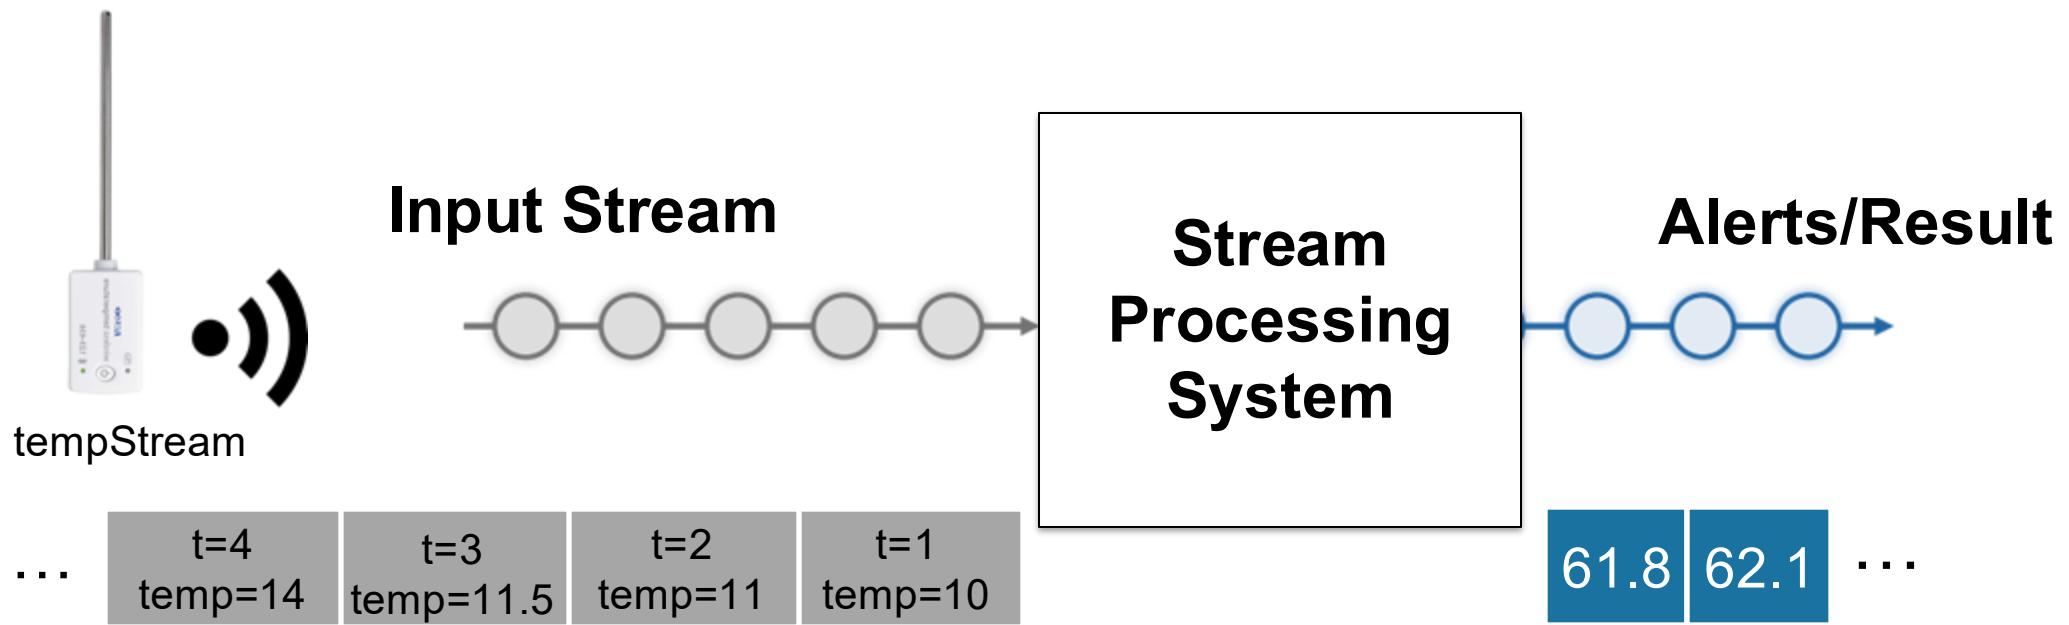
- LCMs in Batch Systems
- **LCMs in Streaming Systems**
- Road Ahead

What is Stream Processing in a Nutshell?



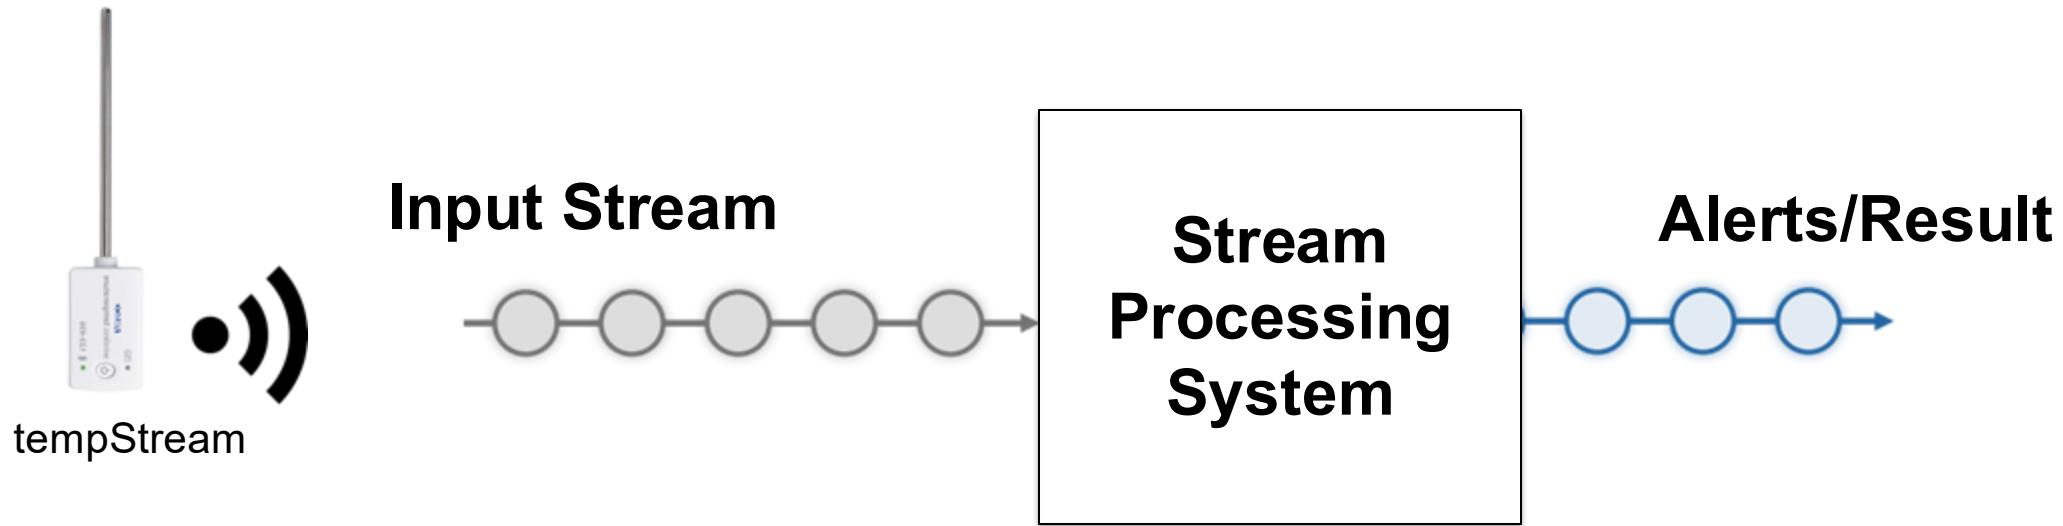
- **Take inputs:** **Continuous data** from devices (cars/buses, health devices, card transactions, social networks, sensors)
- **AND Standing queries** for monitoring (e.g., positions/speed/# of cars)
- **Output:** **Continuous results on standing queries** (time-series)
- **Objectives:** (often) **low latency** and **high throughput**

Stream Processing 101



Query: Notify when average values of temperature is higher than 60°C
(in the last minute, for the last three sensor values, ...)?

Stream Processing 101



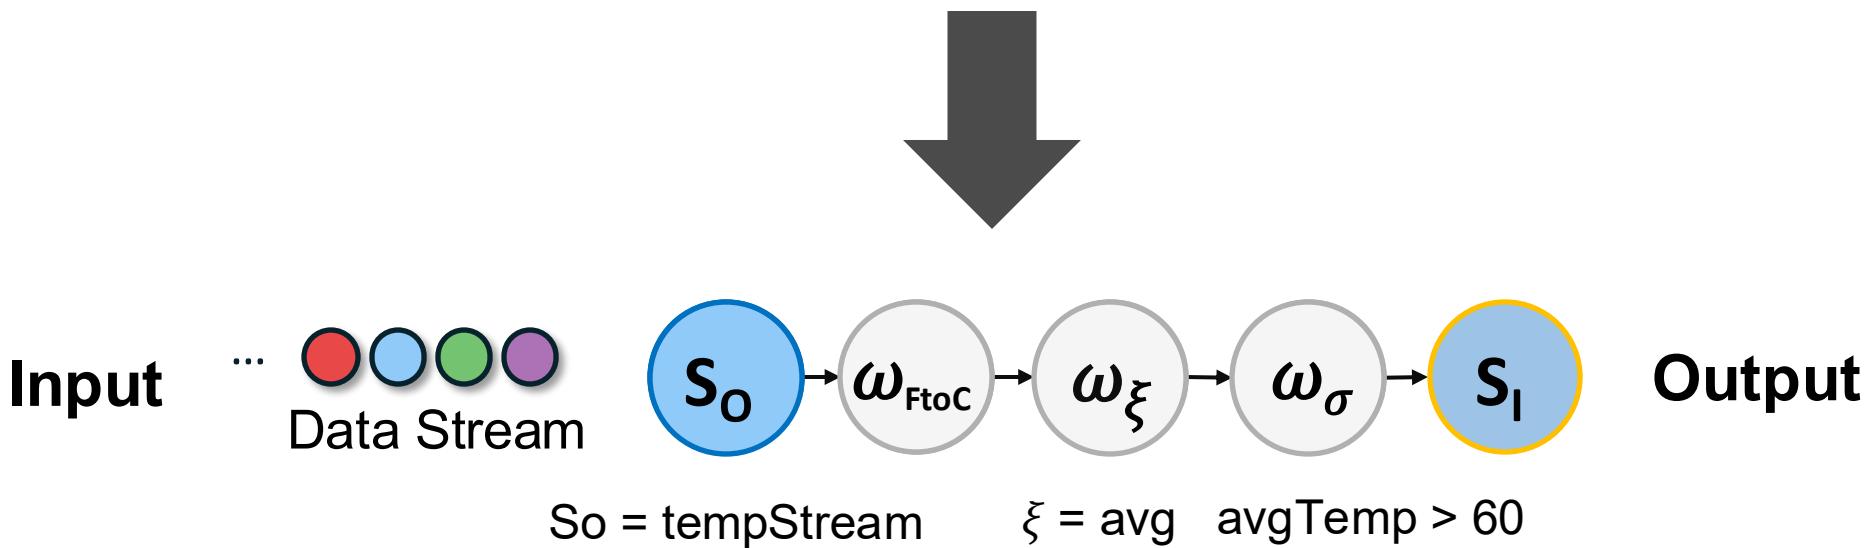
```
1SELECT AVG(FtoC(temp)) as avgTempStream  
FROM tempStream [ROWS 3, ADVANCE BY 1 MIN]  
HAVING avgTemp > 60
```

¹ Query expressed in CQL (continuous query language), a SQL-like query language for streaming.

Stream Processing 101

Queries are compiled **into data flow graph (DFG) of stream operators**

```
SELECT AVG(FtoC(temp)) as avgTempStream  
FROM tempStream [ROWS 3, ADVANCE BY 1 MIN]  
HAVING avgTemp > 60
```



| Poll

*Can traditional cost models of databases
be adapted to estimate costs of data flow
graphs of streaming systems?*

NO FREE LUNCH



No Traditional Cost Models in Streaming!

A Catalog of Stream Processing Optimizations

MARTIN HIRZEL, IBM Watson Research Center

ROBERT SOUË, University of Lugano

SCOTT SCHNEIDER, IBM Watson Research Center

BUĞRA GEDIK, Bilkent University

ROBERT GRIMM, New York University

Avenues for future work. Finding the right sequence in which to apply optimizations is an interesting problem when there are variants of optimizations with complex interactions. Furthermore, while there is literature with cost models for individual optimizations, extending those to work on multiple optimizations is challenging; in part, that is because the existing cost models are usually sophisticated and custom-made for their optimization. Furthermore, models for optimizations must capture characteristics not just of the application, but also of the system and the input data. These characteristics accurately and with moderate cost is another avenue for future work.

Apache Flink™: Stream and Batch Processing in a Single Engine

Paris Carbone[†]
Asterios Katsifodimos*

Stephan Ewen[‡]
Volker Markl*

Seif Haridi[†]
Kostas Tzoumas[‡]

[†]KTH & SICS Sweden
parisc.katsifodimos@kth.se

[‡]data Artisans
first@data-artisans.com

^{*}TU Berlin & DFKI
first.last@tu-berlin.de

and interesting-property propagation. However, the arbitrary UDF-heavy DAGs that make up Flink's dataflow programs do not allow a traditional optimizer to employ database techniques out of the box [17], since the operators hide their semantics from the optimizer. For the same reason, cardinality and cost-estimation methods are equally difficult to employ. Flink's runtime supports various execution strategies including repartition and

[2]

[1] Hirzel, M., Soulé, R., Schneider, S., Gedik, B., & Grimm, R. (2014). A Catalog of Stream Processing Optimizations. ACM Computing Surveys (CSUR), 46(4).

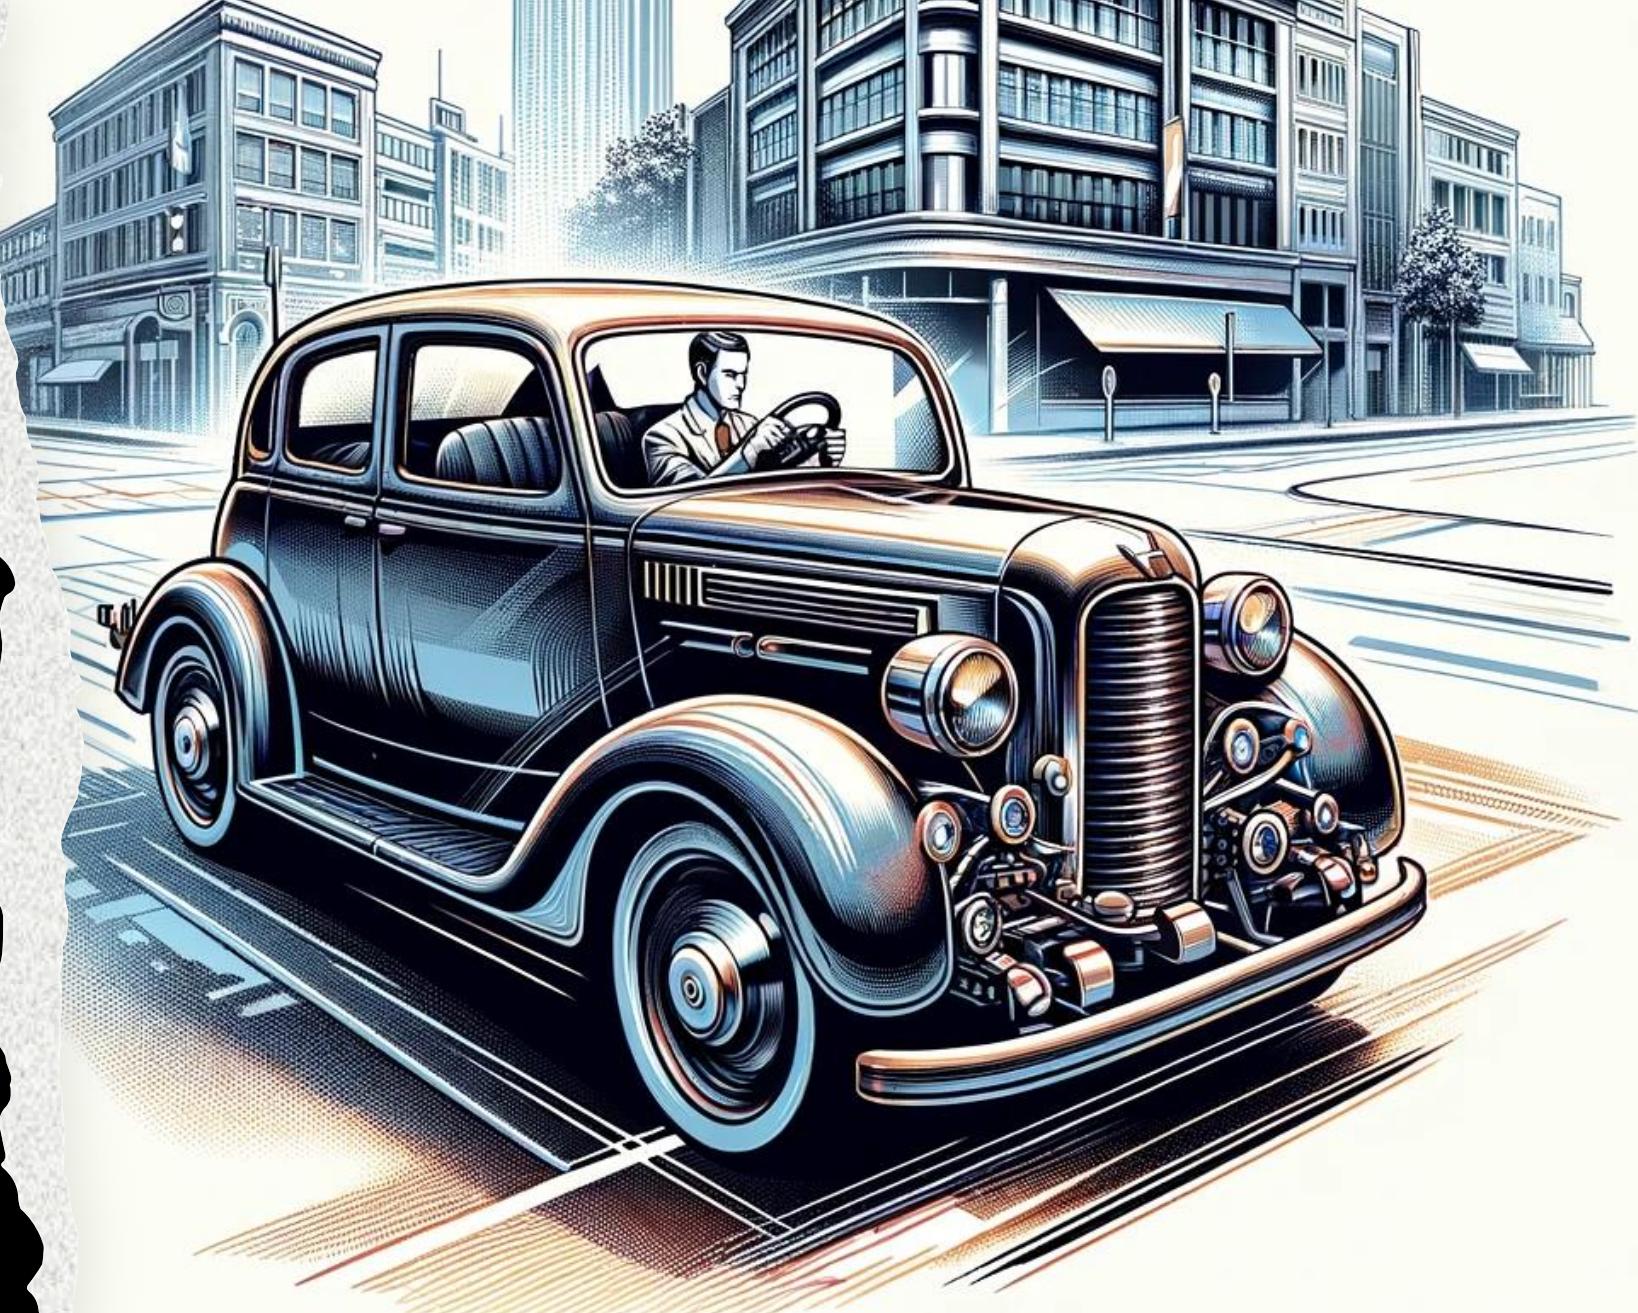
[2] Carbone, P., Katsifodimos, A., Ewen, S., Markl, V., Haridi, S., & Tzoumas, K. (2015). Apache Flink™: Stream and Batch Processing in a Single Engine. IEEE Data Engineering Bulletin.

Using Current Streaming Systems Feels Like

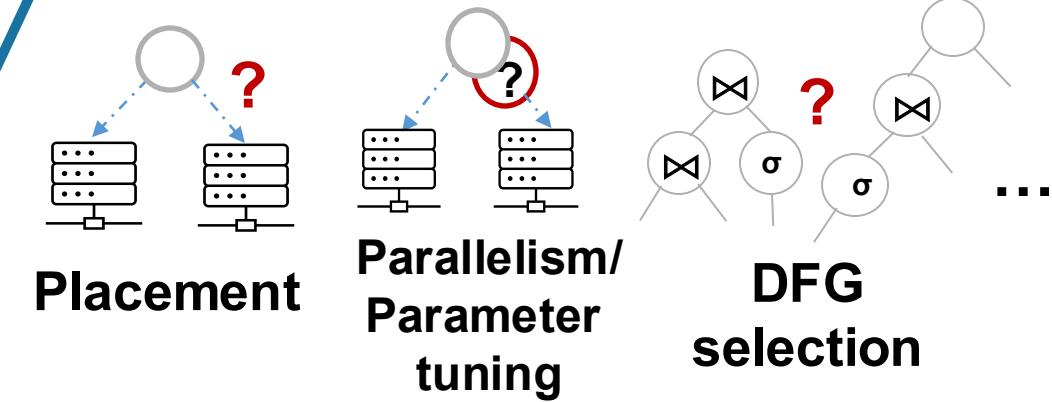
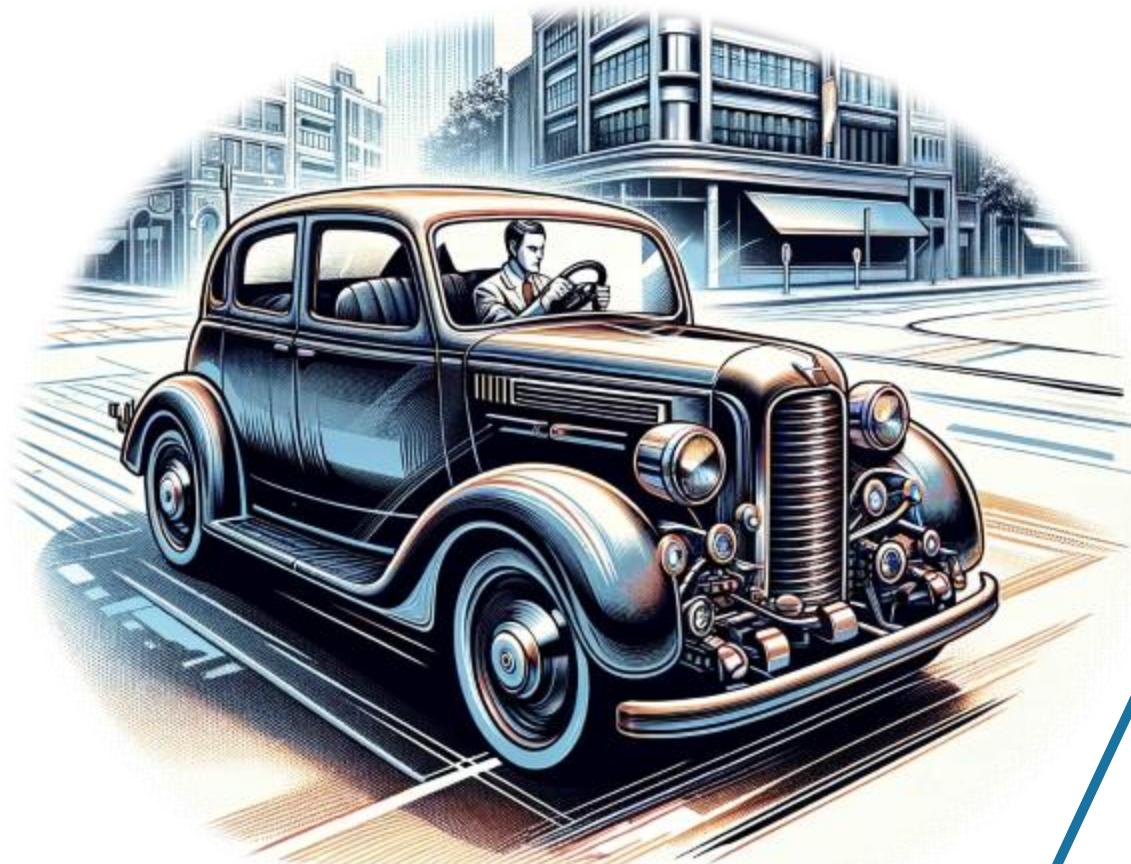


Tutorial: Learned Cost Models for Query Optimization: From Batch to Streaming Systems

Image Source: ChatGPT - DALLE



Optimization Parameters in Stream Processing



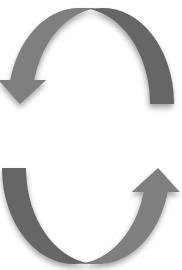
**Stream
Processing
System**

Expert Tuning in Streaming Systems

Notify when average number of cars on the street is greater than 60



Domain Expert



Data Engineer

Extensive Tuning Needed!
→ Expert tuning to meet performance constraints



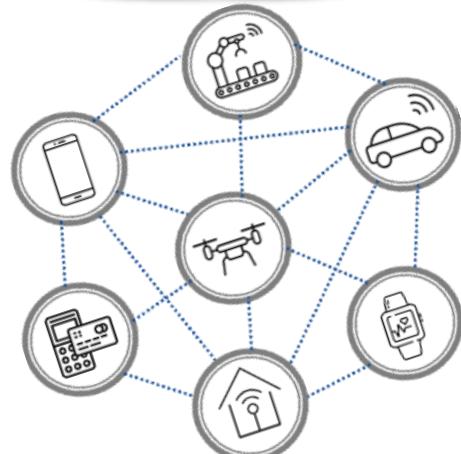
APACHE
STORM™



Spark
Streaming



Apache Flink

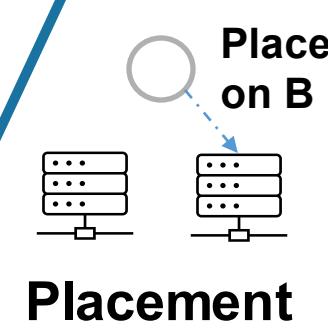


IoT infrastructure

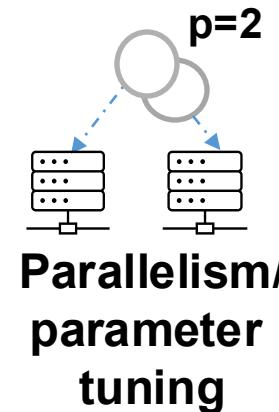
Learned Cost Models To the Rescue



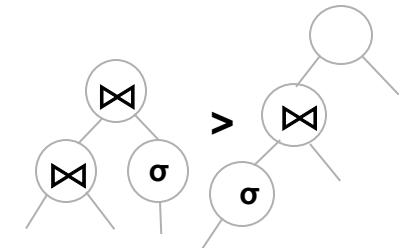
Opportunity: LCMs Enabled Optimizations



Placement



Parallelism/
parameter
tuning

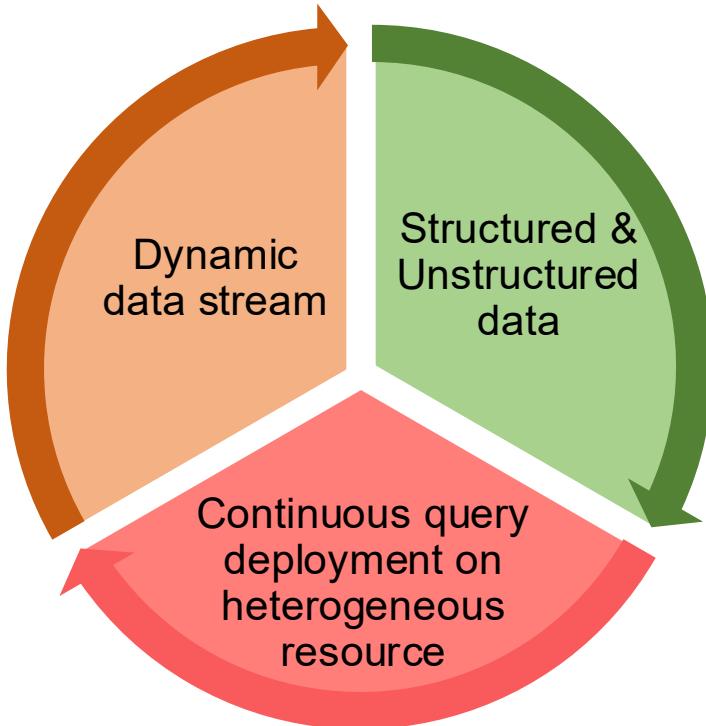


DFG
selection

Enables
Optimizations

Learned
Cost Model

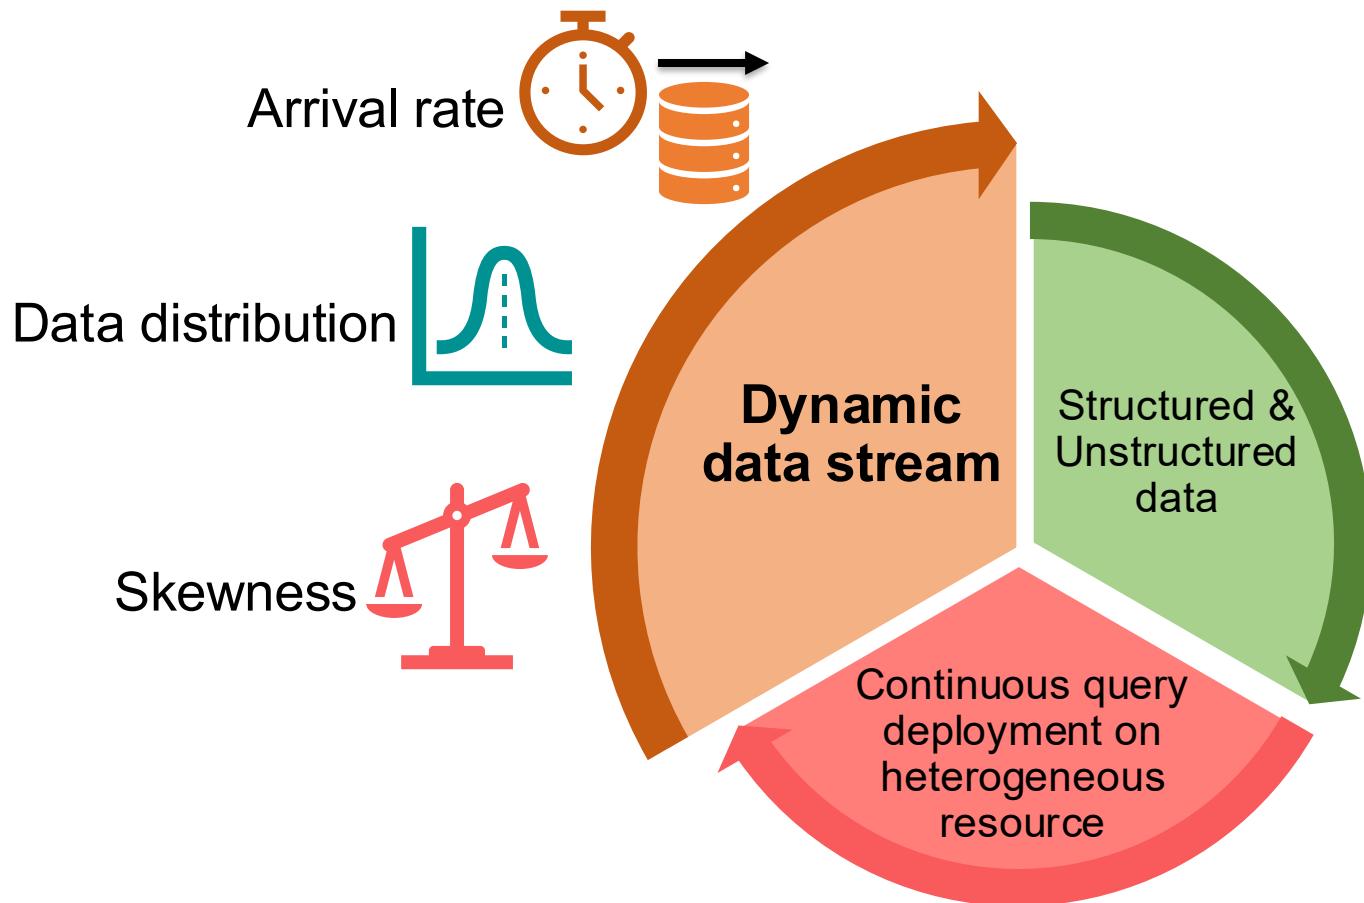
Challenges for LCMs for Streaming



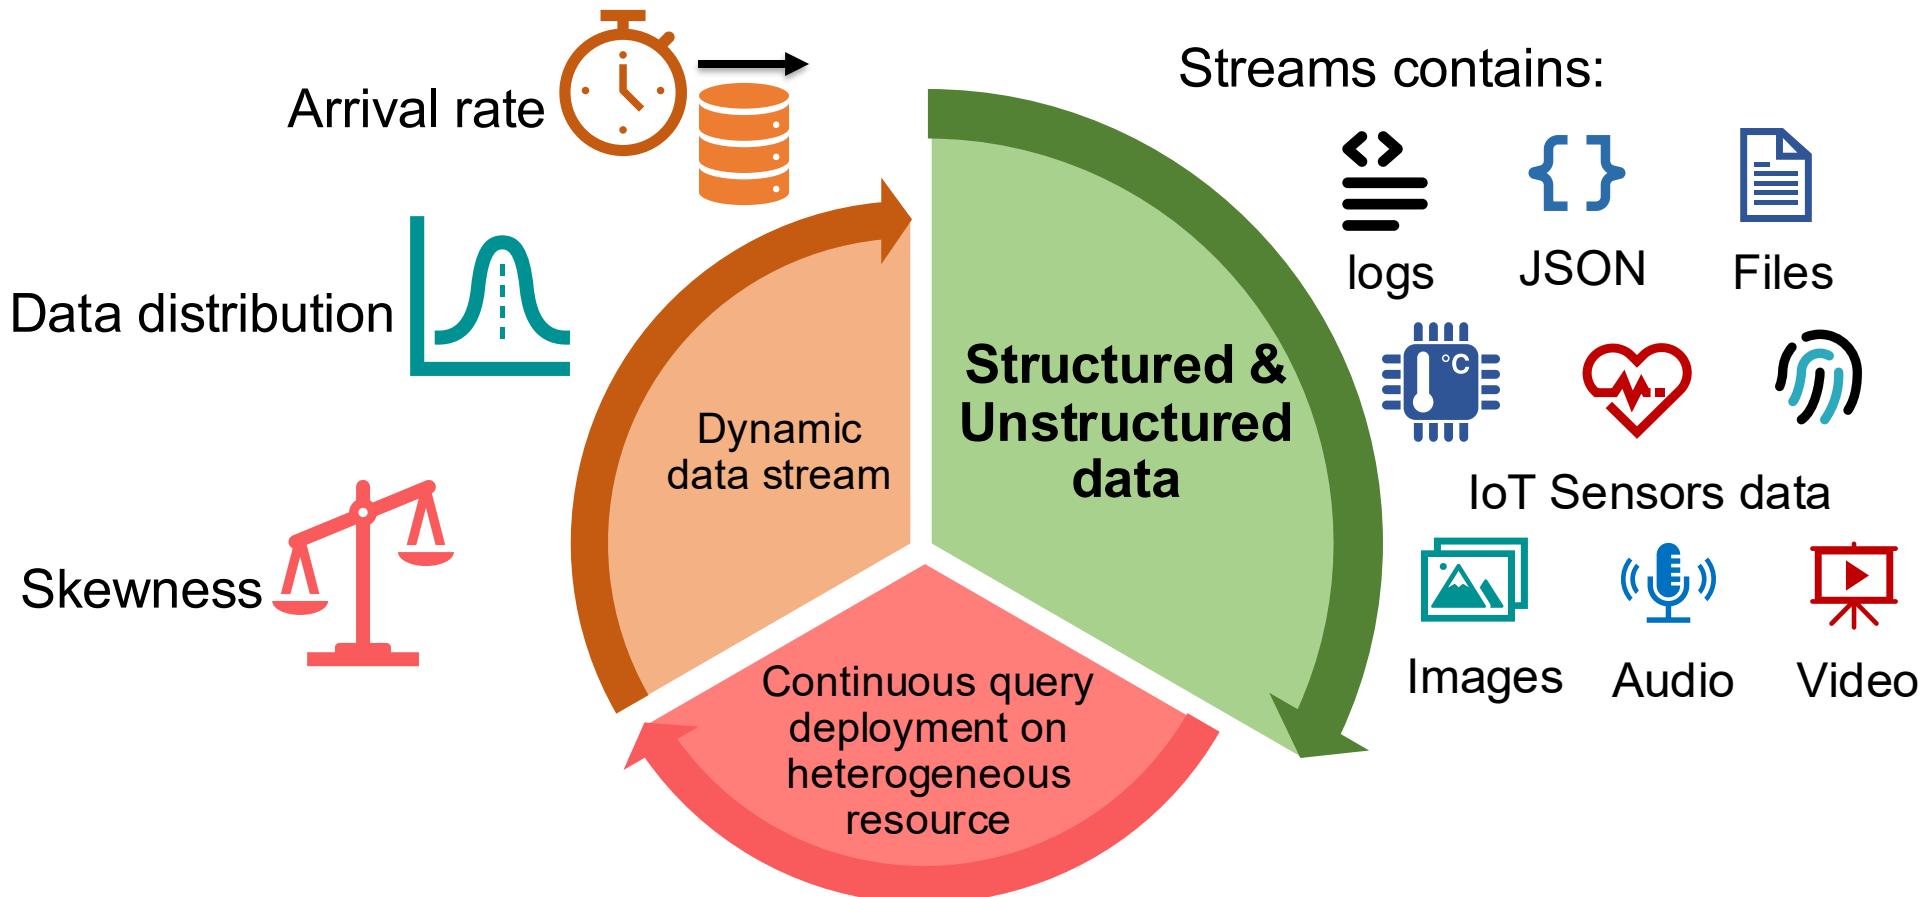
G Rosinosky, D Schmitz, and E Rivière. 2024. StreamBed: Capacity Planning for Stream Processing. DEBS '24

Tutorial: Learned Cost Models for Query Optimization: From Batch to Streaming Systems

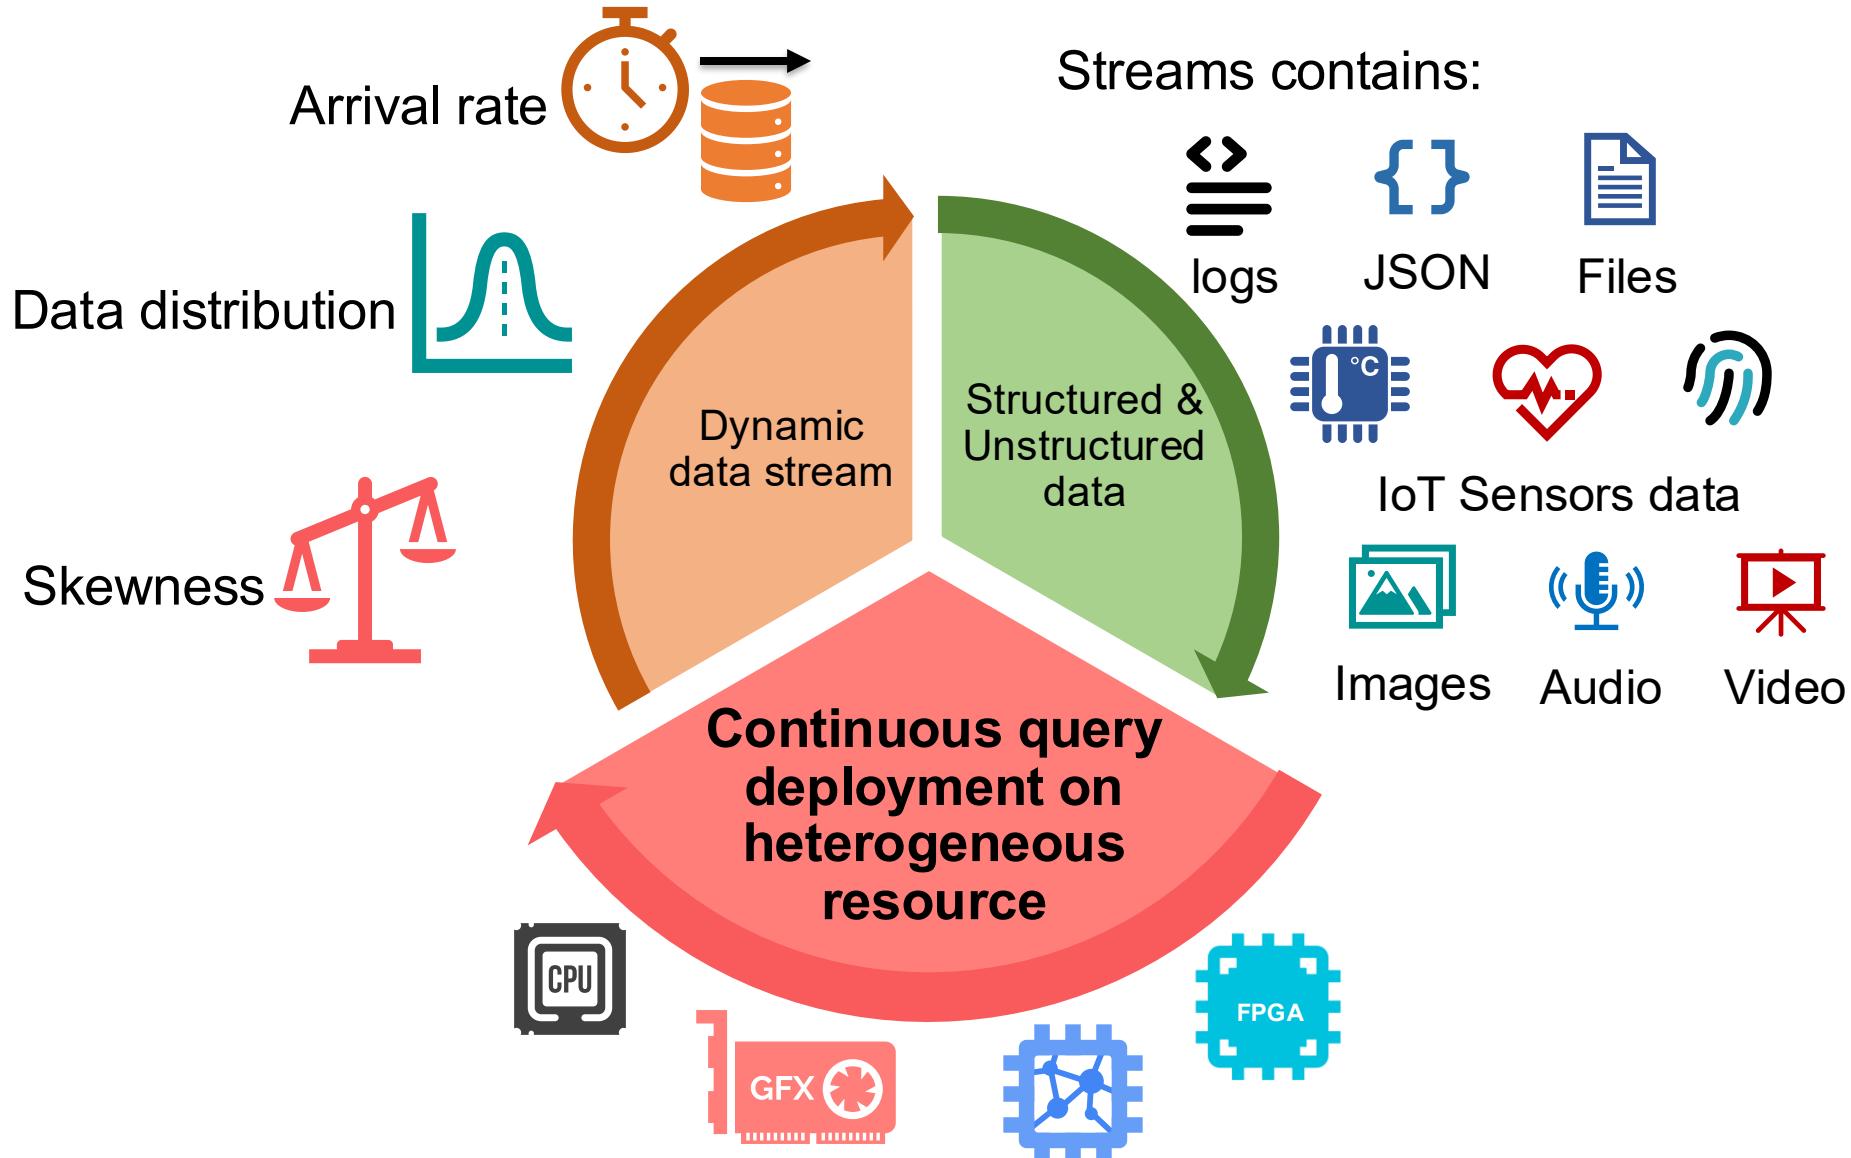
I Challenges for LCMs for Streaming



Challenges for LCMs for Streaming

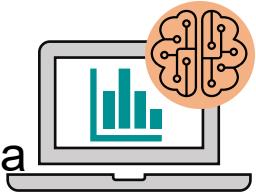


Challenges for LCMs for Streaming



| How existing LCMs deals with them?

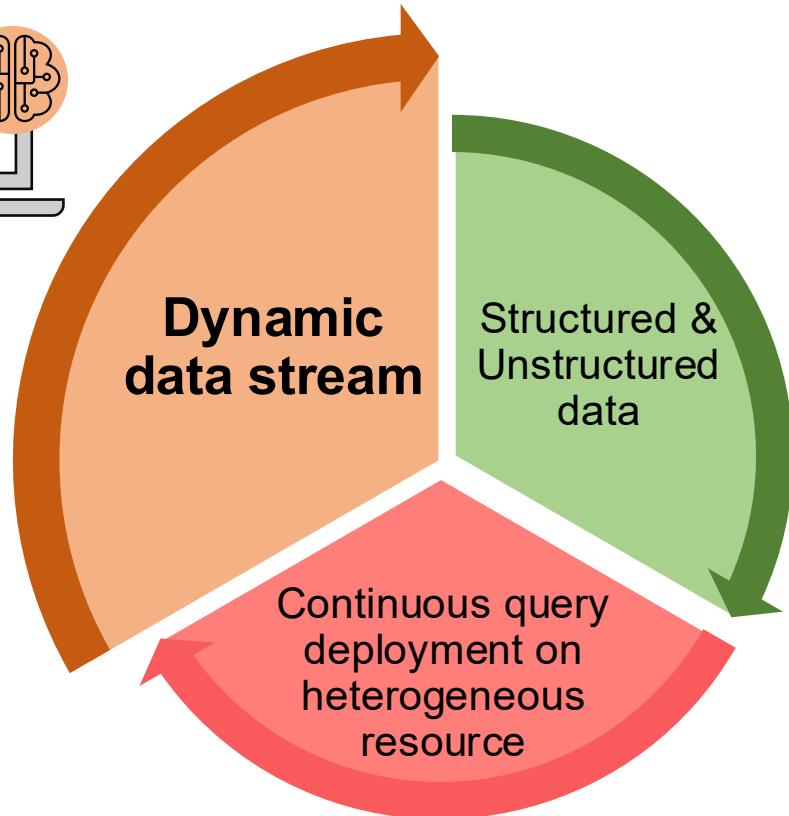
LCMs are (re)trained or fine-tuned on dynamic data



Learn from **feedback loops** (monitoring)

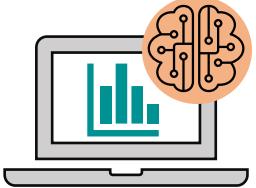


E.g., RL approaches like Decima

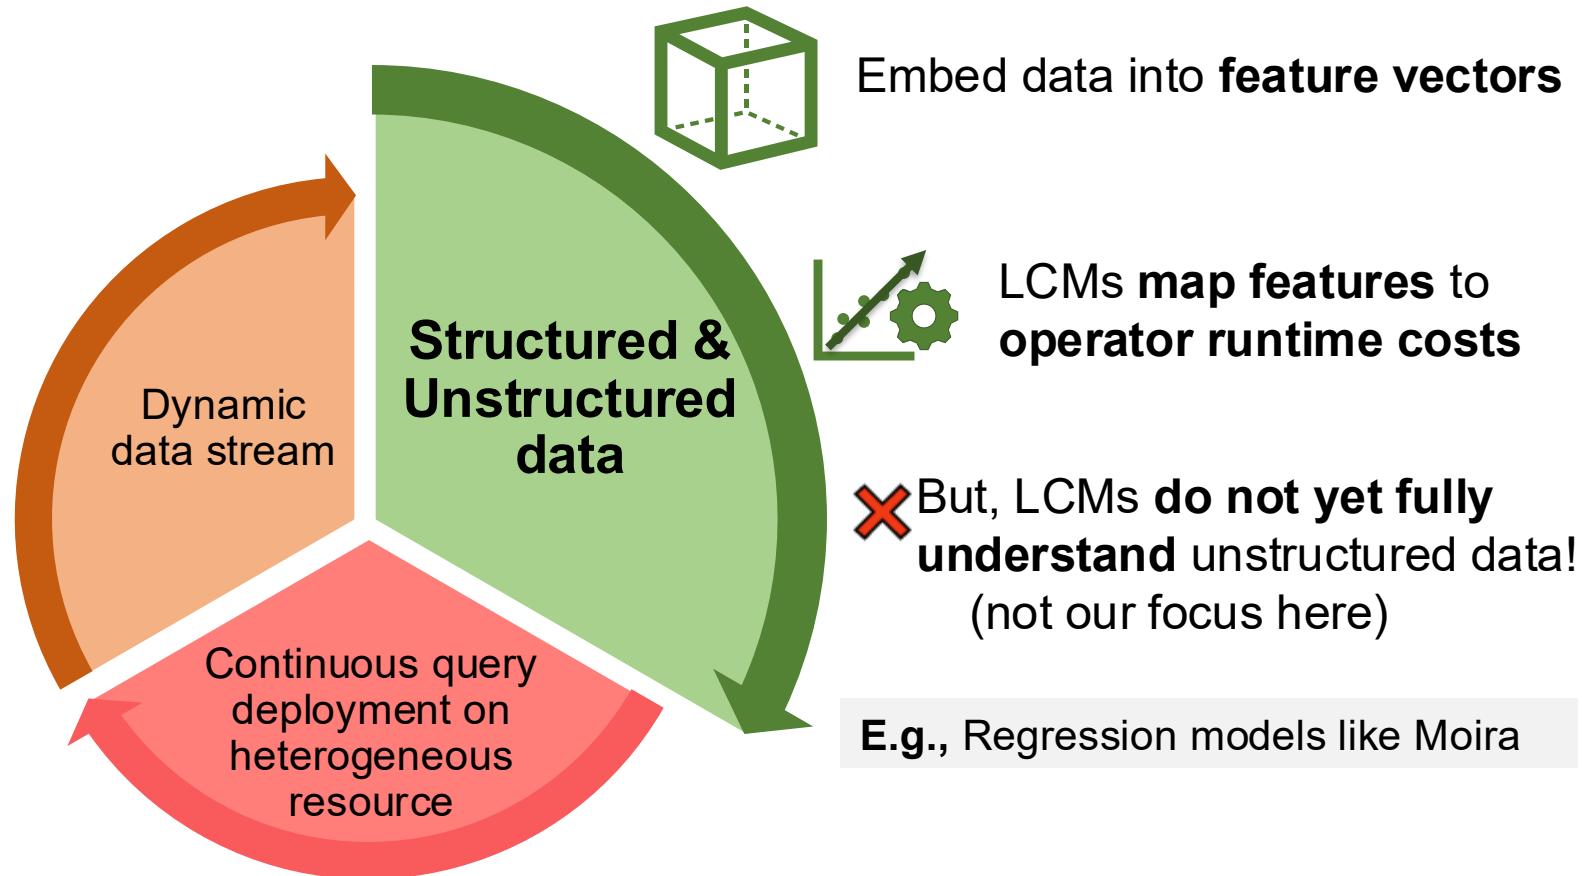


| How existing LCMs deals with them?

LCMs are **trained** or **fine-tuned** on dynamic data

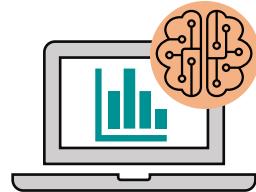


Learn from **feedback loops** (monitoring)



| How existing LCMs deals with them?

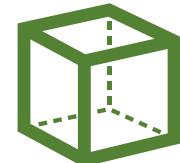
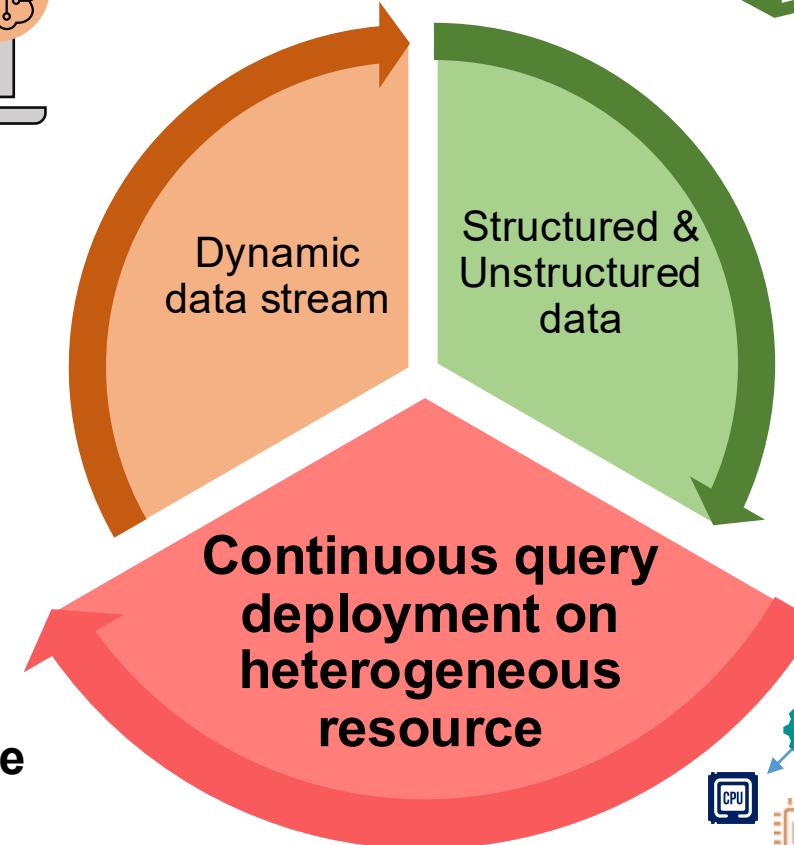
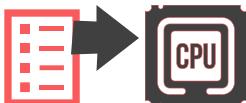
LCMs are **trained** or **fine-tuned** on dynamic data



Learn from **feedback loops** (monitoring)



LCMs include hardware descriptors



Embed data into **feature vectors**



LCMs map **features** to operator runtime costs



LCMs can guide placement decisions

E.g., Optimization oriented LCMs like COSTREAM

Taxonomy of LCMs in Streaming

	Model	Intended Task	Model Architecture	Input Features			
				Stream Statistics	Hardware Charact.	Hardware Monitoring	Query Plan
General LCMs	Moira	Cost Estimation (latency, throughput)	SVM		✓		✓
	Imai et al.	Cost Estimation (throughput)	Linear Reg.			✓	✓
	Li et al. 2014	Cost Estimation (latency)	SVR			✓	✓
	Zero-shot	Cost Estimation (latency, throughput)	GNNs		✓		✓
Optimization-oriented LCMs	ZeroTune	Operator Parallelism	GNNs		✓	✓	✓
	COSTREAM	Operator Placement	GNNs		✓	✓	✓
	Li et al. 2016	Operator Placement	RL		✓		✓
	Decima	Operator Placement	RL		✓		✓
	Ni et al.	Operator Placement	RL		✓		✓

Generalization vs. Specialization

Generalizable LCM
(e.g., Zero-shot)

VS

Specialized LCM
(e.g. RL tuned for workload)



- Transfer across workloads and hardware
- Use **transferable features**
 - Shows high **accuracy on unseen workloads**



- Optimize for a given workload/task
- Use **runtime-driven features**
 - Shows high **accuracy on known workloads**



Advantages

can better deal with workload drifts,
adaptable



Disadvantages

high (one-time) training effort



Advantages

better accuracy (overfit to workload),
adapts online



Disadvantages

retraining required to deal with
workload drift

I Era of Generalizable Cost Models for Streaming

Zero-Shot Cost Models for Distributed Stream Processing

Roman Heinrich
DHBW Mannheim

Manisha Luthra
Technical University
of Darmstadt

Harald Kornmayer
DHBW Mannheim

Carsten Binnig
Technical University of
Darmstadt & DFKI

ABSTRACT

This paper proposes a learned cost estimation model for Distributed Stream Processing Systems (DSPS) with an aim to provide accurate cost predictions of executing queries. A major premise of this work is that the proposed learned model can generalize to the dynamics of streaming workloads *out-of-the-box*. This means a model once trained can accurately predict performance metrics such as *latency* and *throughput* even if the characteristics of the data and workload or the deployment of operators to hardware changes at runtime. That way, the model can be used to solve tasks such as optimizing the placement of operators to minimize the end-to-end latency of a streaming query or maximize its throughput even under varying conditions. Our evaluation on a well-known DSPS, Apache Storm, shows that the model can predict accurately for unseen workloads and queries while generalizing across real-world benchmarks.

CCS CONCEPTS

- Information systems → Stream management.

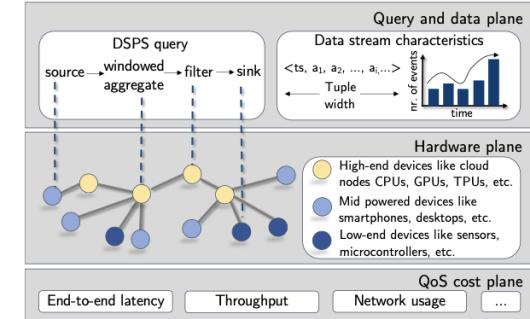
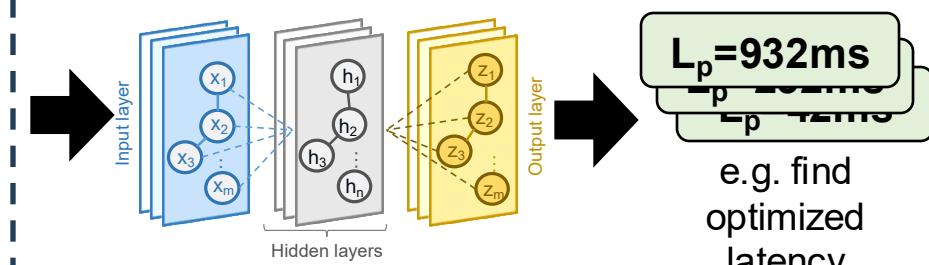
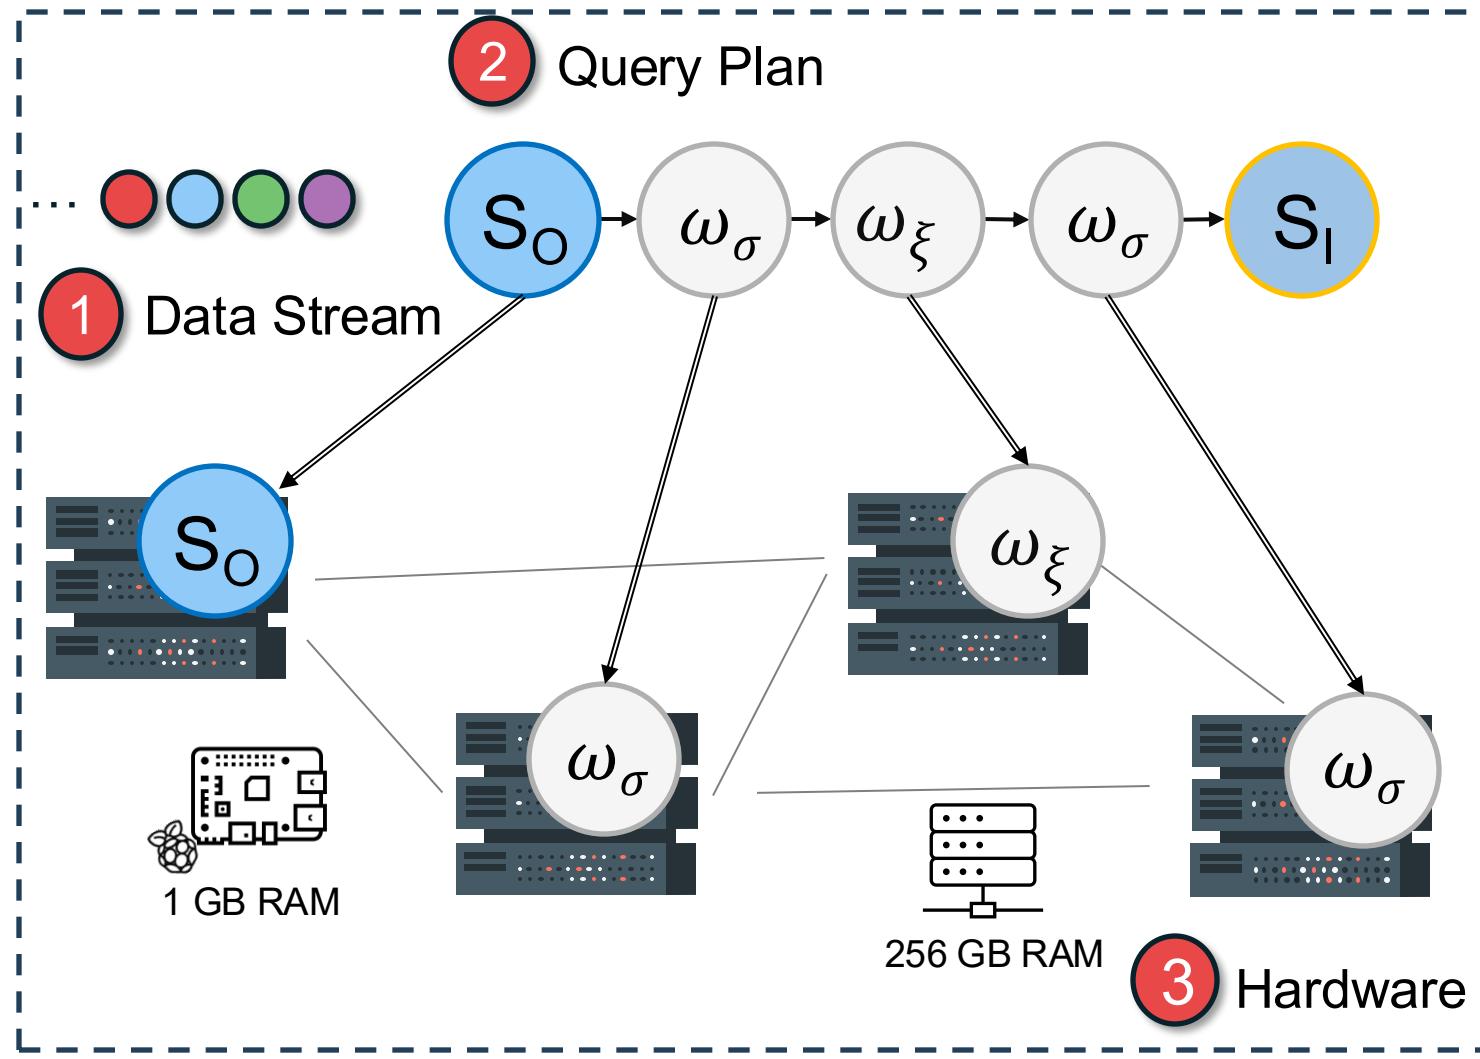


Figure 1: A DSPS has to provide guarantees in terms of one or more quality-of-service (QoS) cost metrics such as latency and throughput. The challenge is that DSPS serve a wide range of workloads on potentially diverse hardware, which makes the cost estimation harder.

Typically, a DSPS provides QoS guarantees using optimization mechanisms such as *operator placement* that usually monitors the costs to decide on the mapping of operators to hardware as shown in Figure 1 [2]. Moreover, frequent reconfigurations of the operator placement are required based on the observed changes of the work-

Zero-shot Cost Model in a Nutshell



GNN model

Cost Predictions
Latency, Throughput, Backpressure, ...

Generalizable Models for Optimizations

Generalizable Resource Allocation in
Stream Processing via Deep Reinforcement Learning

Xiang Ni,^{*}
^{*}Citadel
xiang.ni@citadel.com

Abstract

This paper considers the problem of resource allocation in stream processing, where continuous data is processed in real time in a large distributed system. To maximize system throughput, the resource allocations must consider the computation tasks of a stream processing system and the communication between computing devices. This problem is known to be NP-hard. Graph partitioning is a key component of stream processing systems, and its performance is crucial to practical streaming systems. Many algorithms have been developed to find good partitions. In this paper, we present a general framework to learn a *generalized* graph partitioning strategy that can properly distribute the computation tasks of a stream processing system across multiple computing devices.

COSTREAM: Learned Cost Model for Operator Placement in Edge-Cloud Environments

Roman Heinrich
DHBW Mannheim

Abstract—In this work, we propose COSTREAM, a learned cost model for Distributed Stream Processing (DSP) that provides accurate predictions for operator placement. The learned cost model can be used to find an optimal placement of operators across heterogeneous hardware in edge-cloud environments. In our experiments, COSTREAM can produce highly accurate initial operator placement recommendations, queries, and hardware configurations. We show that COSTREAM can optimize the placements of streaming queries up to 21x faster than existing approaches.

I. INTRO

Pratyush Agnihotri*, Boris Koldehofe†, Paul Stiegele*,
^{*}Technische Universität Darmstadt, [†]Technische Universität Berlin

Abstract—This paper introduces ZEROTUNE, a novel cost model for parallel and distributed stream processing that can be used to effectively set initial parallelism degrees of streaming queries. Unlike existing models, which rely majorly on online learning statistics that are non-transferable, context-specific, and require extensive training, ZEROTUNE proposes *data-efficient zero-shot learning techniques* that enable very accurate cost predictions without having observed any query deployment. To overcome these challenges, we propose ZEROTUNE, a graph neural network architecture that can learn from the structural complexity of parallel distributed stream processing systems, enabling them to adapt to unseen workloads and hardware configurations. In our experiments, we show when integrating ZEROTUNE in a distributed streaming system such as Apache Flink, we can accurately set the degree of parallelism, showing an average speed-up of around 5x in comparison to existing approaches.

Index Terms—Zero-shot cost models, Parallelism tuning

ZEROTUNE: Learned Zero-Shot Cost Models for Parallelism Tuning in Stream Processing

Learning from the Past: Adaptive Parallelism Tuning for Stream Processing Systems

Yuxing Han¹, Lixiang Chen^{1,2}, Hui Wang¹, Chengcheng Yang²,
¹ByteDance Inc, ²East China Normal University
1{hanyuxing, chenlixiang.3608, wanghaoyu}@bytedance.com
2ccyang@dase.ecnu.edu.cn



v2 [cs.DC] 7 Jul 2025

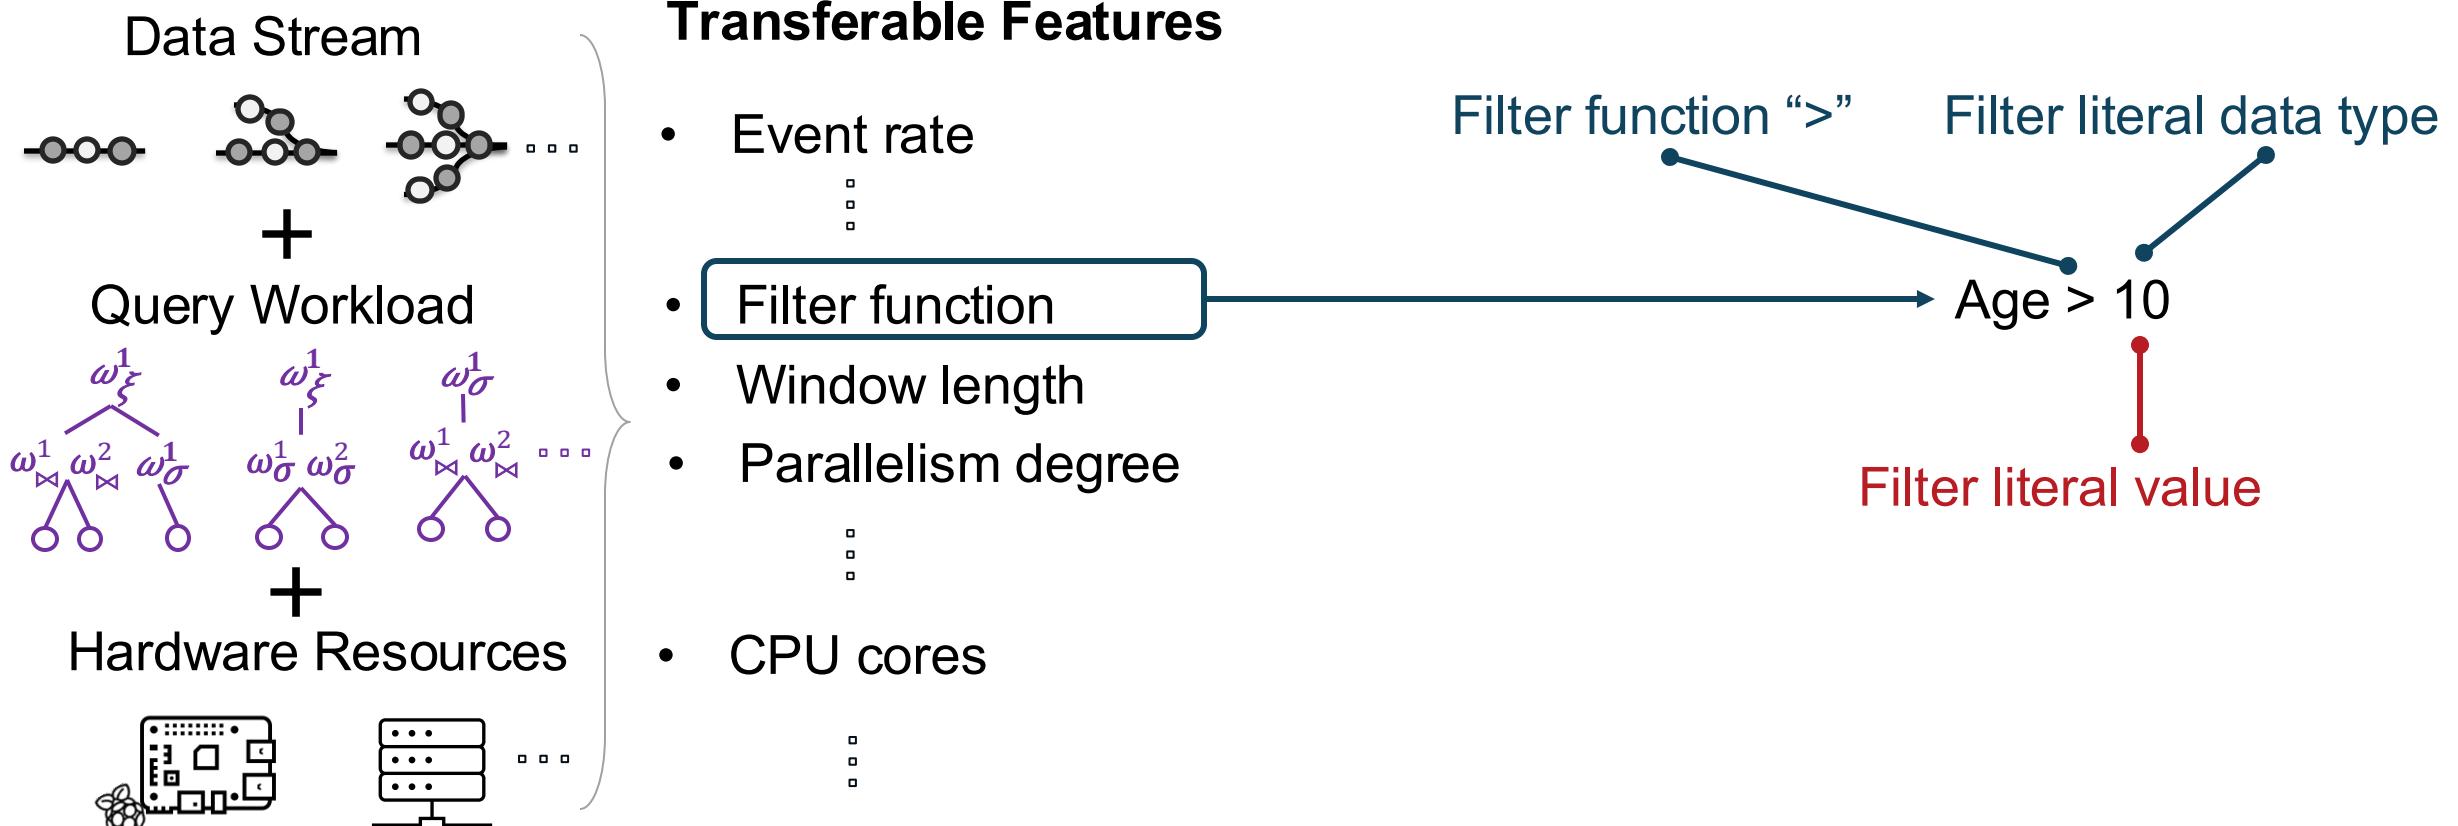
Abstract—Distributed stream processing systems rely on the dataflow model to define and execute streaming jobs, organizing computations as Directed Acyclic Graphs (DAGs) of operators. Adjusting the parallelism of these operators is crucial to handling fluctuating workloads efficiently while balancing resource usage and processing performance. However, existing methods often fail to effectively utilize execution histories or fully exploit DAG structures, limiting their ability to identify bottlenecks and determine the optimal parallelism. In this paper, we propose StreamTune, a novel approach for adaptive parallelism tuning in stream processing systems. StreamTune incorporates a pre-training and fine-tuning framework that leverages global knowledge from historical execution data for job-specific parallelism tuning. In the pre-training phase, StreamTune clusters the historical data with Graph Edit Distance and pre-trains a Graph Neural Network-based encoder per cluster to capture the correlation between the operator parallelism, DAG structures, and the identified operator-level bottlenecks. In the online tuning phase, StreamTune iteratively refines operator parallelism recommendations using an operator-level bottleneck prediction model enforced with a monotonic constraint, which aligns with the observed

capture data dependencies between these operators. Dataflow execution relies on asynchronous message passing, allowing each operator to process data independently, achieving both high throughput and fault tolerance. In real-world applications, dataflow execution should accommodate fluctuating workload characteristics, such as varying data arrival rates.

Traditionally, system engineers manage these fluctuations by manually adjusting the parallelism of dataflow operators to match different workload demands. This process involves increasing the parallelism (i.e., scaling out) during peak periods to maintain performance and decreasing the parallelism (i.e., scaling in) during off-peak times to conserve resources. However, manual tuning is labor-intensive and error-prone, which often results in suboptimal resource allocation. Ineffective adjustments might lead to over-provisioning during periods of low demand, resulting in resource wastage; or under-provisioning during sudden workload peaks, potentially causing violations of Service Level Objectives (SLOs) [12]. As

Zero-Shot Cost Model: Training

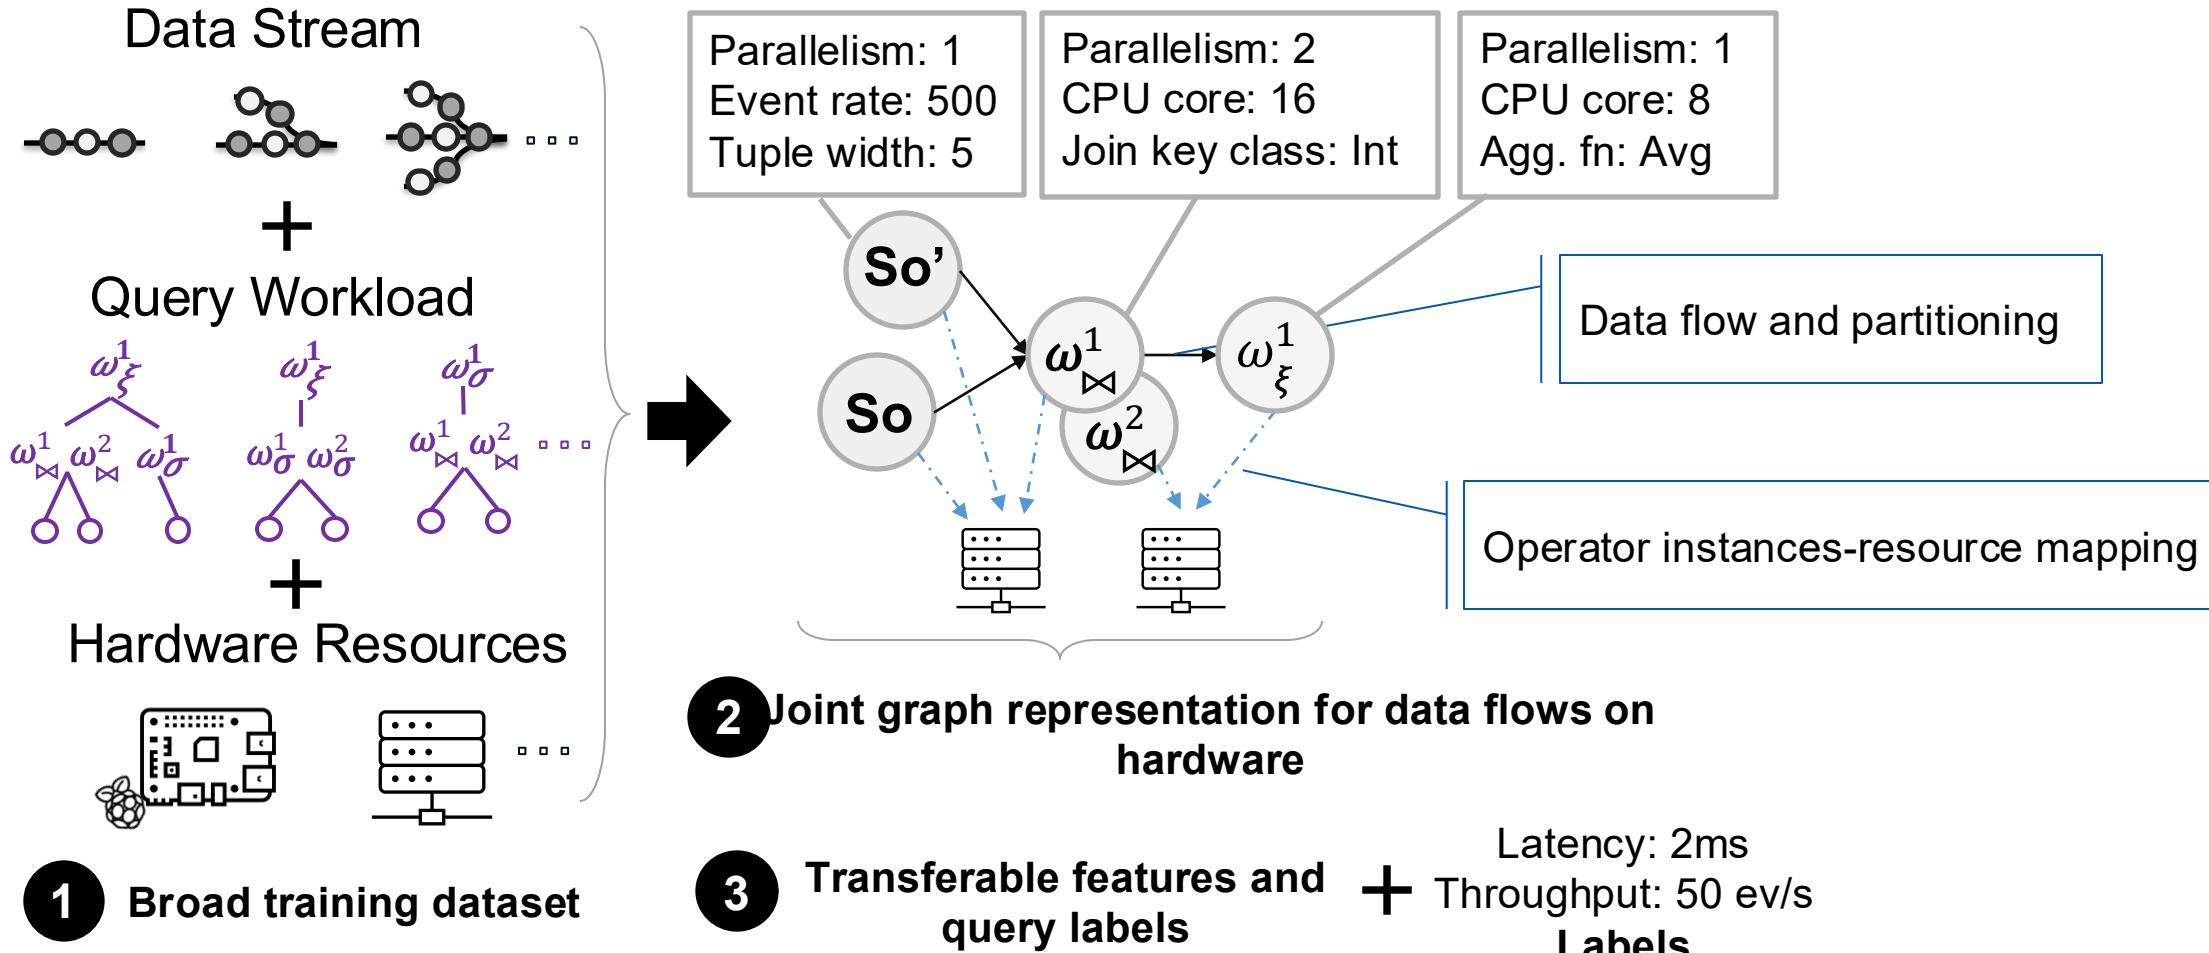
Training Zero-Shot Cost Model



1 Broad training dataset

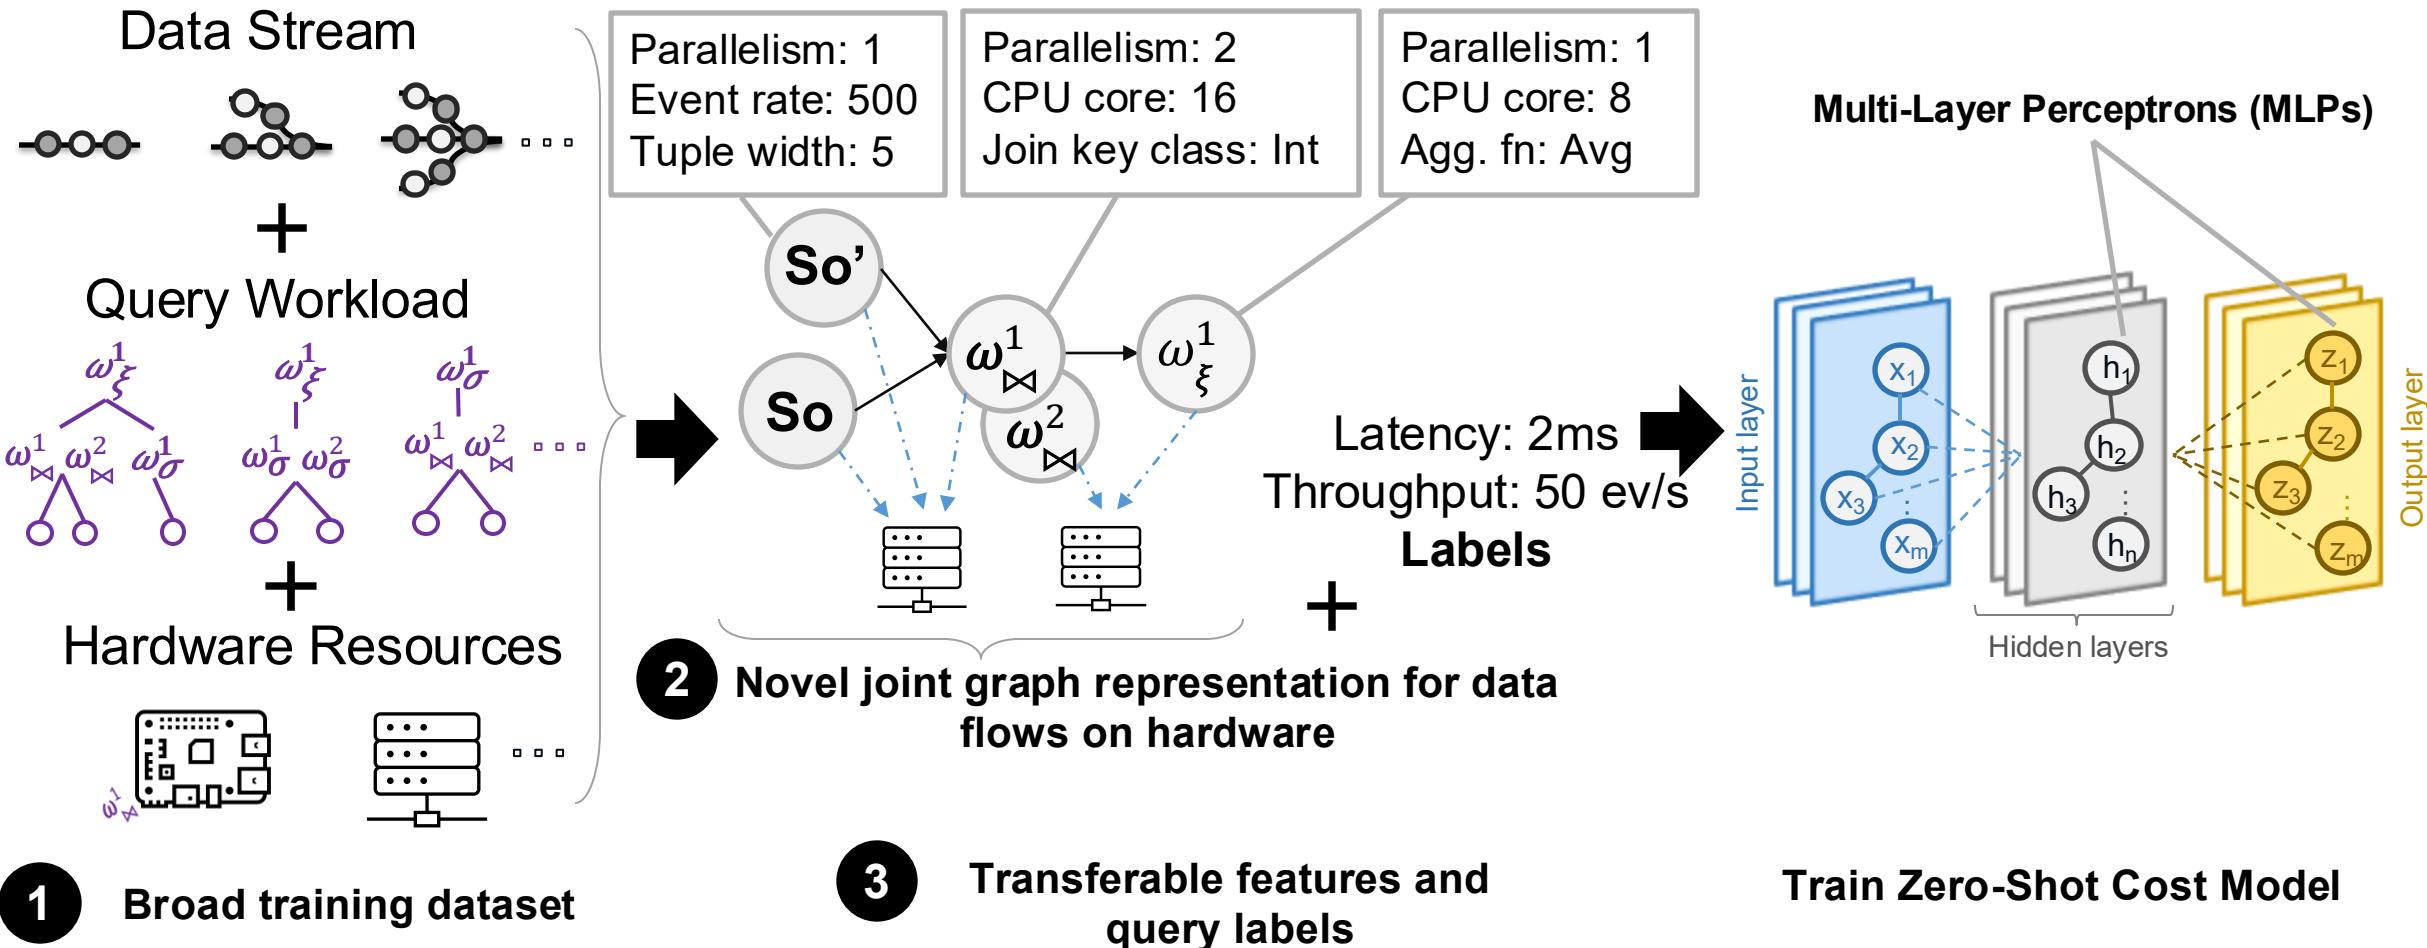
Zero-Shot Cost Model: Training

Training Zero-Shot Cost Model



Zero-Shot Cost Model: Training

Training Zero-Shot Cost Model

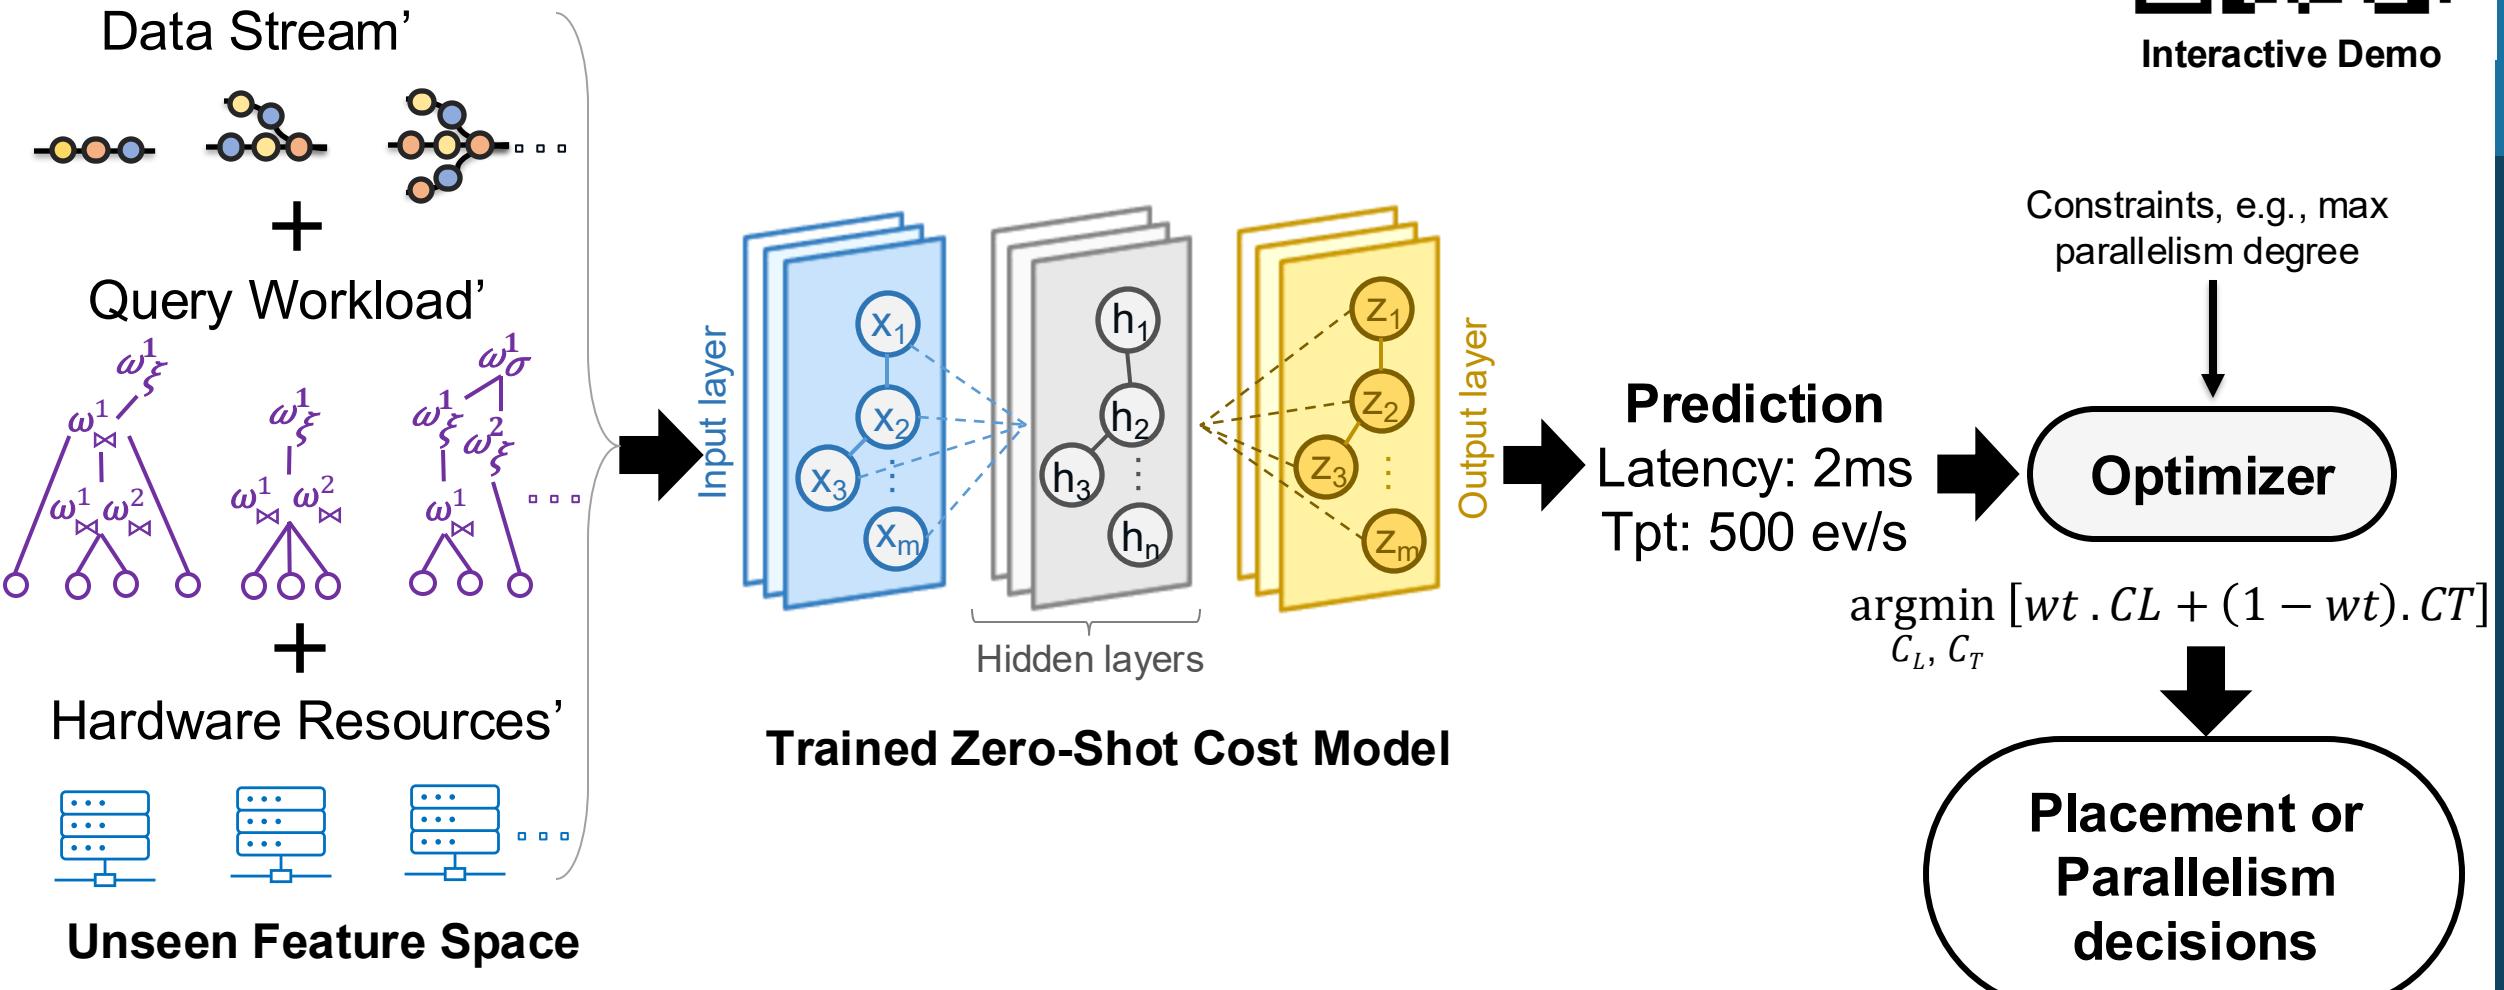


Zero-Shot Cost Model: Inference



Interactive Demo

Inference and Optimization using Zero-Shot Cost Model



Learning Query Placement Costs with GNN

Neural Encoding

Transferable features

CPU: 4 cores
RAM: 1024 MB
Bandwidth: 20 MBps
Latency: 5 ms



Encodings

[0.79]
[0.50]
[0.002]
...

Join key: ...

$$\omega \bowtie$$

Join
Encoder

...

Event rate: ...



Source
Encoder

...

Novel Neural Message Passing

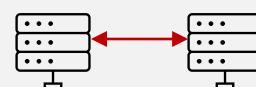
1. Message
passing from
operators to hosts



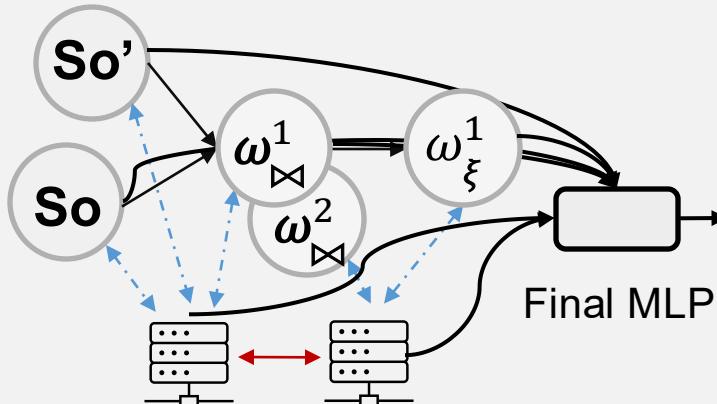
2. Message
passing from **hosts**
to **operators**



3. Message
passing **between**
(parallel)
resources



4. Message passing **through operator graph**



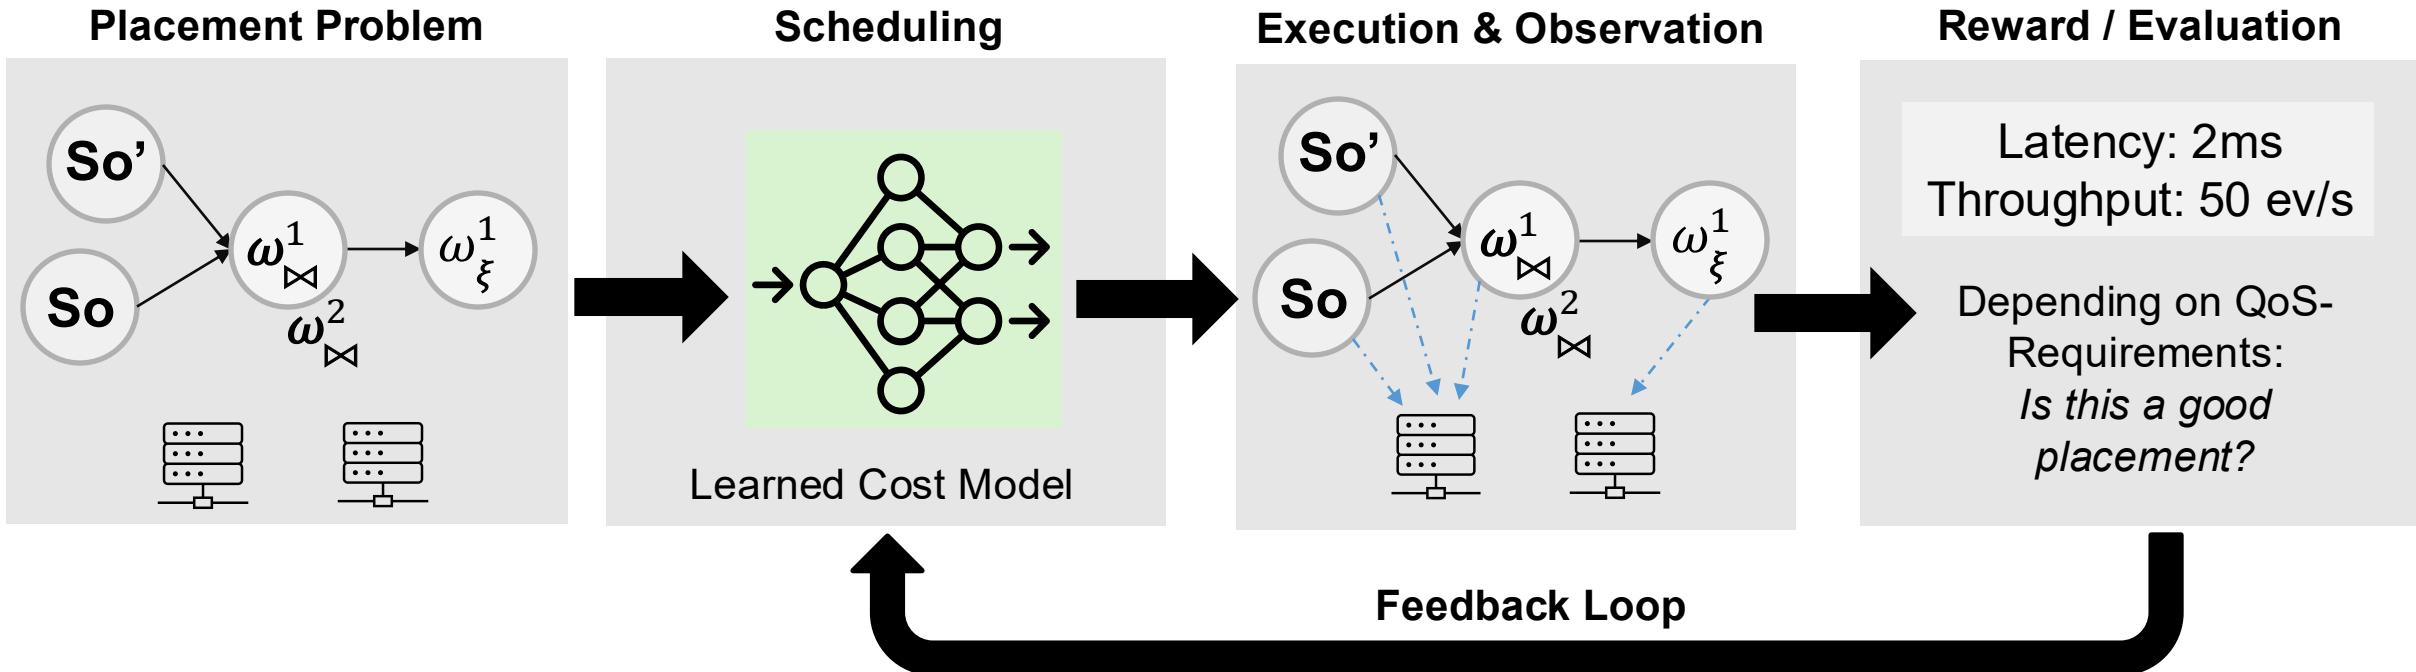
Predicted Costs
Throughput: 25 ev/s,
E2E-Latency: 212ms

Final MLP

Specialized Models in a Nutshell

Specialized LCM
(e.g. RL tuned for workload)

Intuition: Improving model over time for a given workload by monitoring the results and iteratively updating the model



Advantages

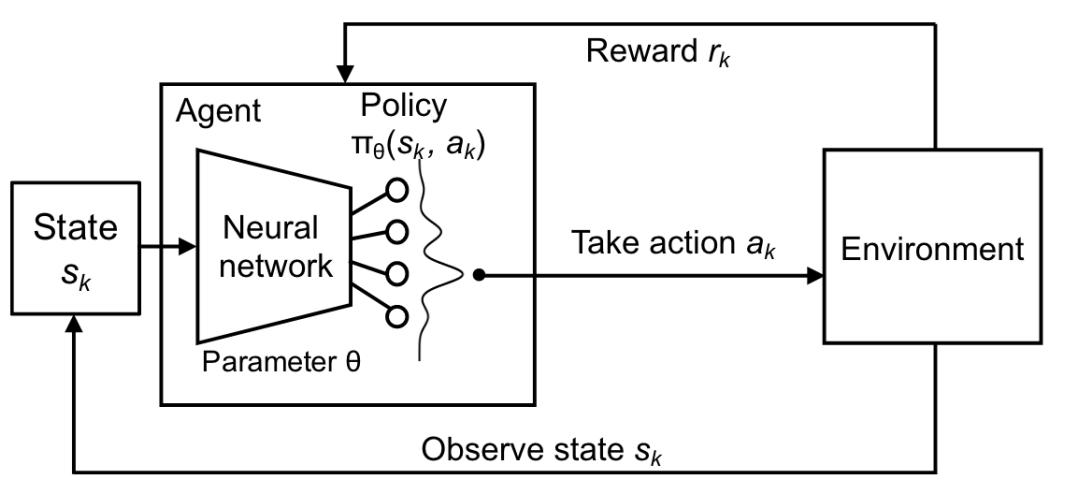
- No human interaction required as policy improves over time
- Avoids the massive collection of training data

Disadvantages

- Model gets tied towards seen workloads and does not generalize
- Retraining required if workloads change

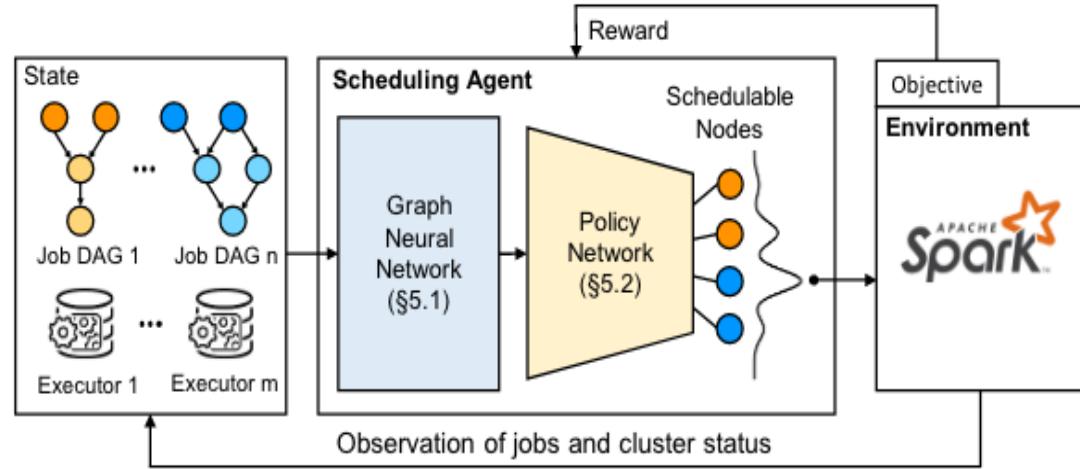
LCMs using Reinforcement Learning

Background: Reinforcement Learning



- Learning an **agent** by interacting with the environment
- Learning **policy** over time: Which actions to take given a system state?
- Assuming markov process: Actions are conditionally independent of the past

Decima: Learning Scheduling Algorithms with RL



- For a given query an agent uses a GNN and a policy network to come up with a schedule
- The schedule is executed on a spark cluster and observed
- The agent is updated by learning from a reward function of the given placement

H.Mao, M. Schwarzkopf, S.B. Venkatakrishnan, Z. Meng, M. Alizadeh
Learning Scheduling Algorithms for Data Processing Clusters SIGCOMM 2019.

Tutorial: Learned Cost Models for Query Optimization: From Batch to Streaming Systems

More approaches follow this idea:
- Moira (Foroni et al)
- Li et al

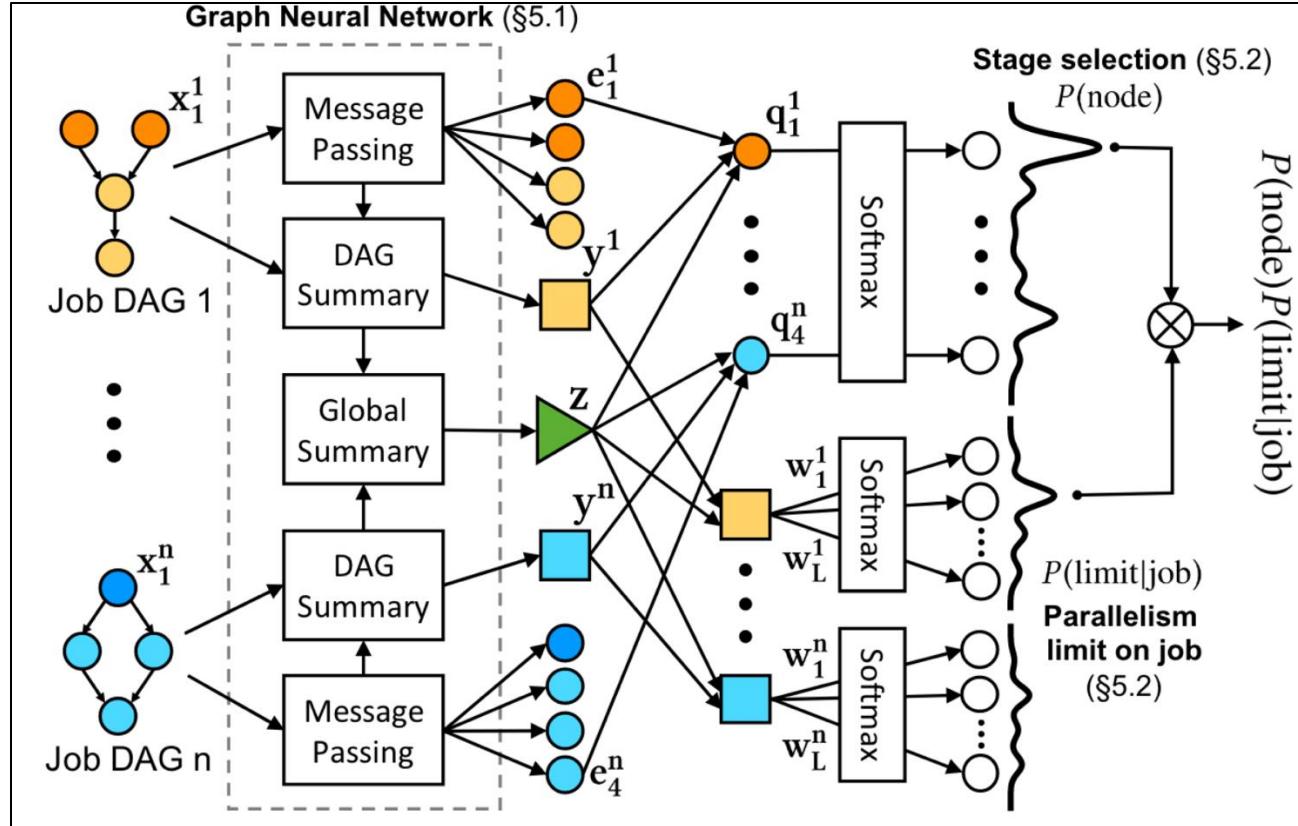
Decima: Learning Placement Costs

Features:

- number of tasks within operator
- average task duration
- number of executors working on the node
- available local executors

Embeddings:

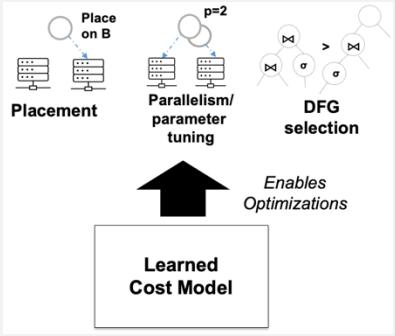
- per-node embedding (e)
- per query embeddings (y)
- global summary



Output:

- the score of the schedule
- maximal parallelism degree

Summary: LCMs for Stream Processing



Role of LCMs in Optimizing Stream Processing Systems

Model	Intended Task	Model Architecture	Input Features			
			Stream Statistics	Hardware Charact.	Hardware Monitoring	Query Plan
Moira	Cost Estimation (latency, throughput)	SVM	✓		✓	
Imai et al.	Cost Estimation (throughput)	Linear Reg.	✓		✓	
Li et al.	Cost Estimation (latency)	SVR	✓		✓	
Zero-shot	Cost Estimation (latency, throughput)	GNNs	✓			✓
ZeroTune	Operator Parallelism	GNNs	✓	✓		✓
COSTREAM	Operator Placement	GNNs	✓	✓		✓
Li et al.	Operator Placement	RL	✓			✓
Decima	Operator Placement	RL	✓			✓
Ni et al.	Operator Placement	RL	✓			✓

Taxonomy on existing work

**Generalizable LCM
(e.g., Zero-shot)**

**Specialized LCM
(e.g. RL tuned for workload)**

**Key Dimension:
Generalizability vs. Specialization**

| Open problems for batch and streaming systems

Training Data Collection

- *How can we collect the right labels efficiently?*
- *Do we need both optimal and non optimal query plans?*
- *How do we capture load fluctuations (streaming)?*



LCMs for Query Optimization

- *Which are the right models for query optimization?*
- *Which are the right metrics for LCMs beyond Q-error?*



I Open problems for batch and streaming systems

LCMs Evaluation



- *Which is a right benchmark with fixed training/validation/testing split?*
- *What are good metrics that reflect the downstream task?*

LCMs Interpretability & Explainability

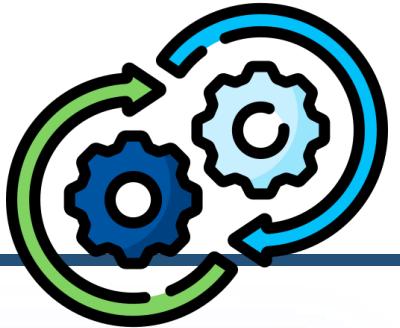


- *Shall we aim for white-box models instead of NNs?*
- *What's the trade-off between “accuracy” and interpretability?*
- *How do we explain the results stemming from a black-box LCM?*

Open problems for batch and streaming systems

LCMs for Hybrid Workloads

- *How can we build LCMs that support batch-stream workloads, commonly found in data lake settings?*
- *Do we need specialized LCMs per type, or could one be used?*



Summary

Learned Cost Models for Batch and Streaming Systems

