

Opening the Black-Box

Explaining Learned Cost Models For Databases



TECHNISCHE
UNIVERSITÄT
DARMSTADT

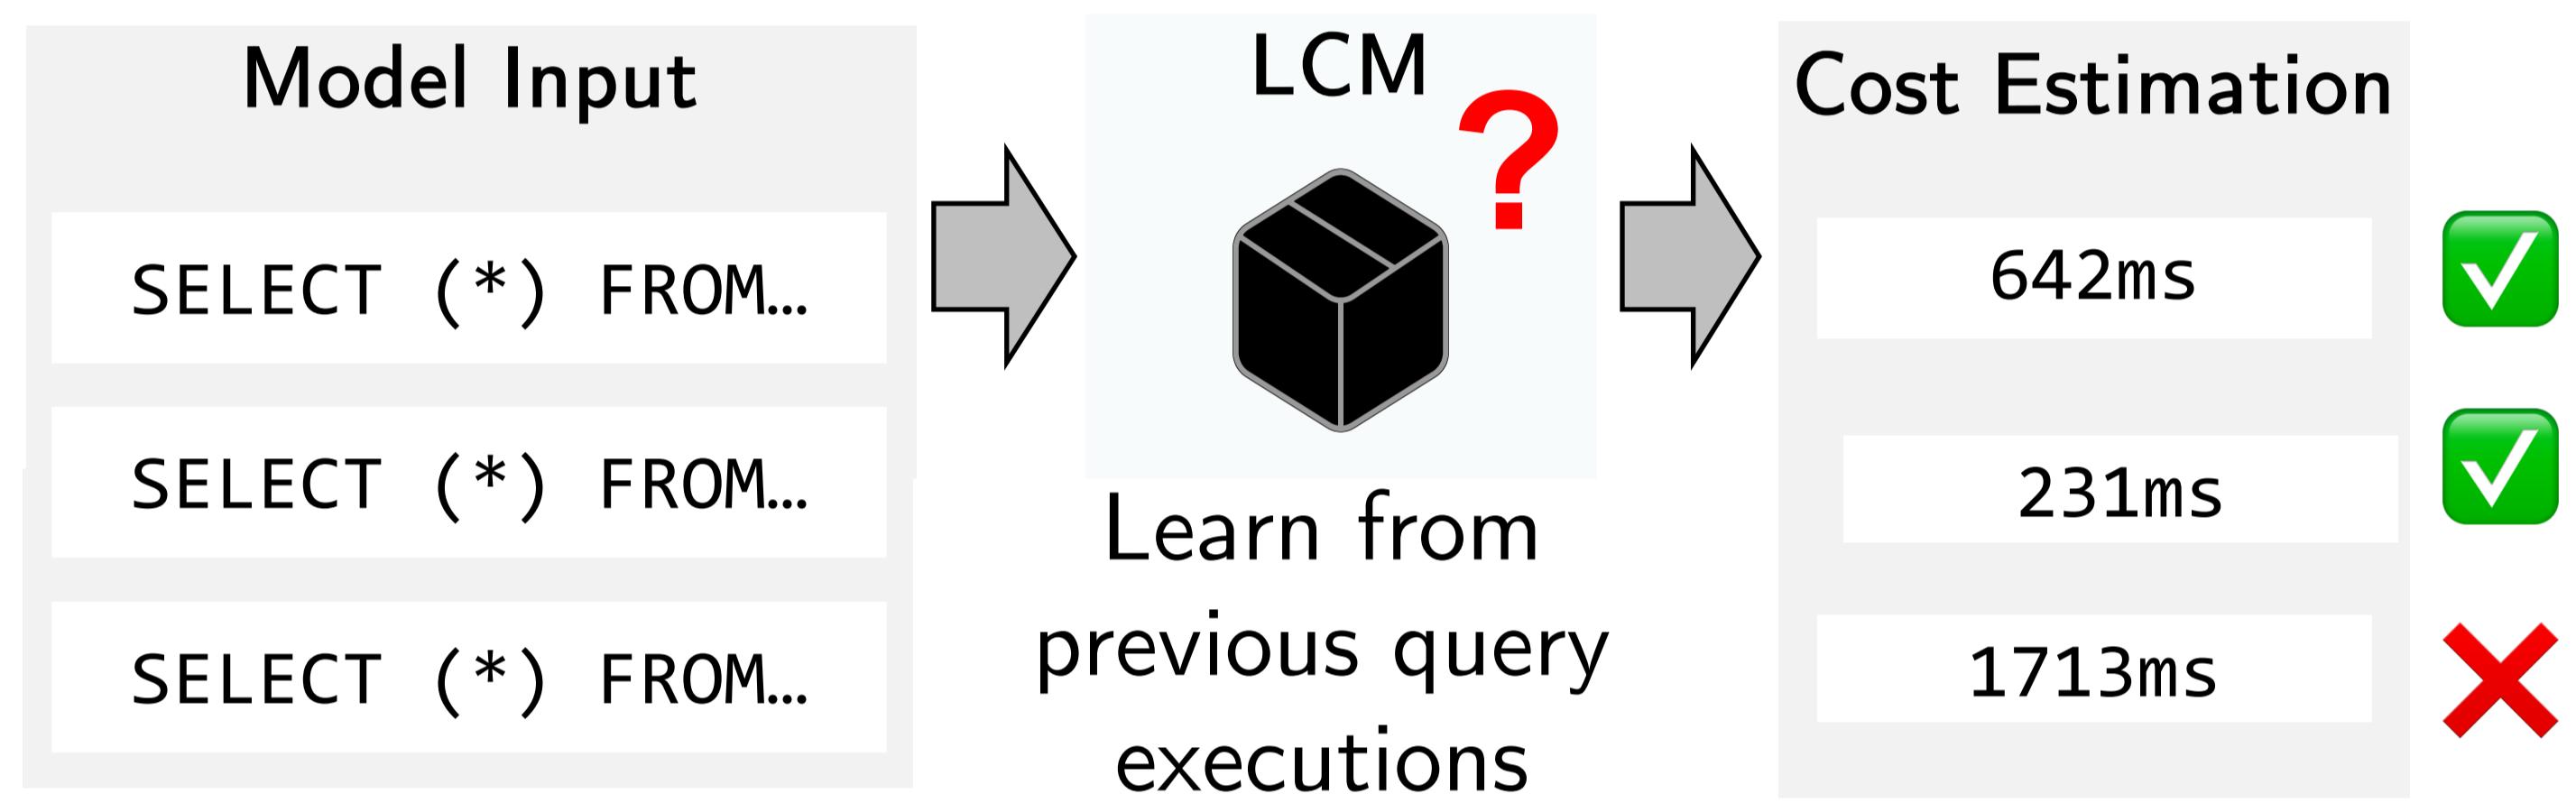
Roman Heinrich, Oleksandr Havrylov, Manisha Luthra,
Johannes Wehrstein & Carsten Binnig



What has a Learned Cost Model (LCM) really learned?

Motivation: LCMs Behave as a Black-Box

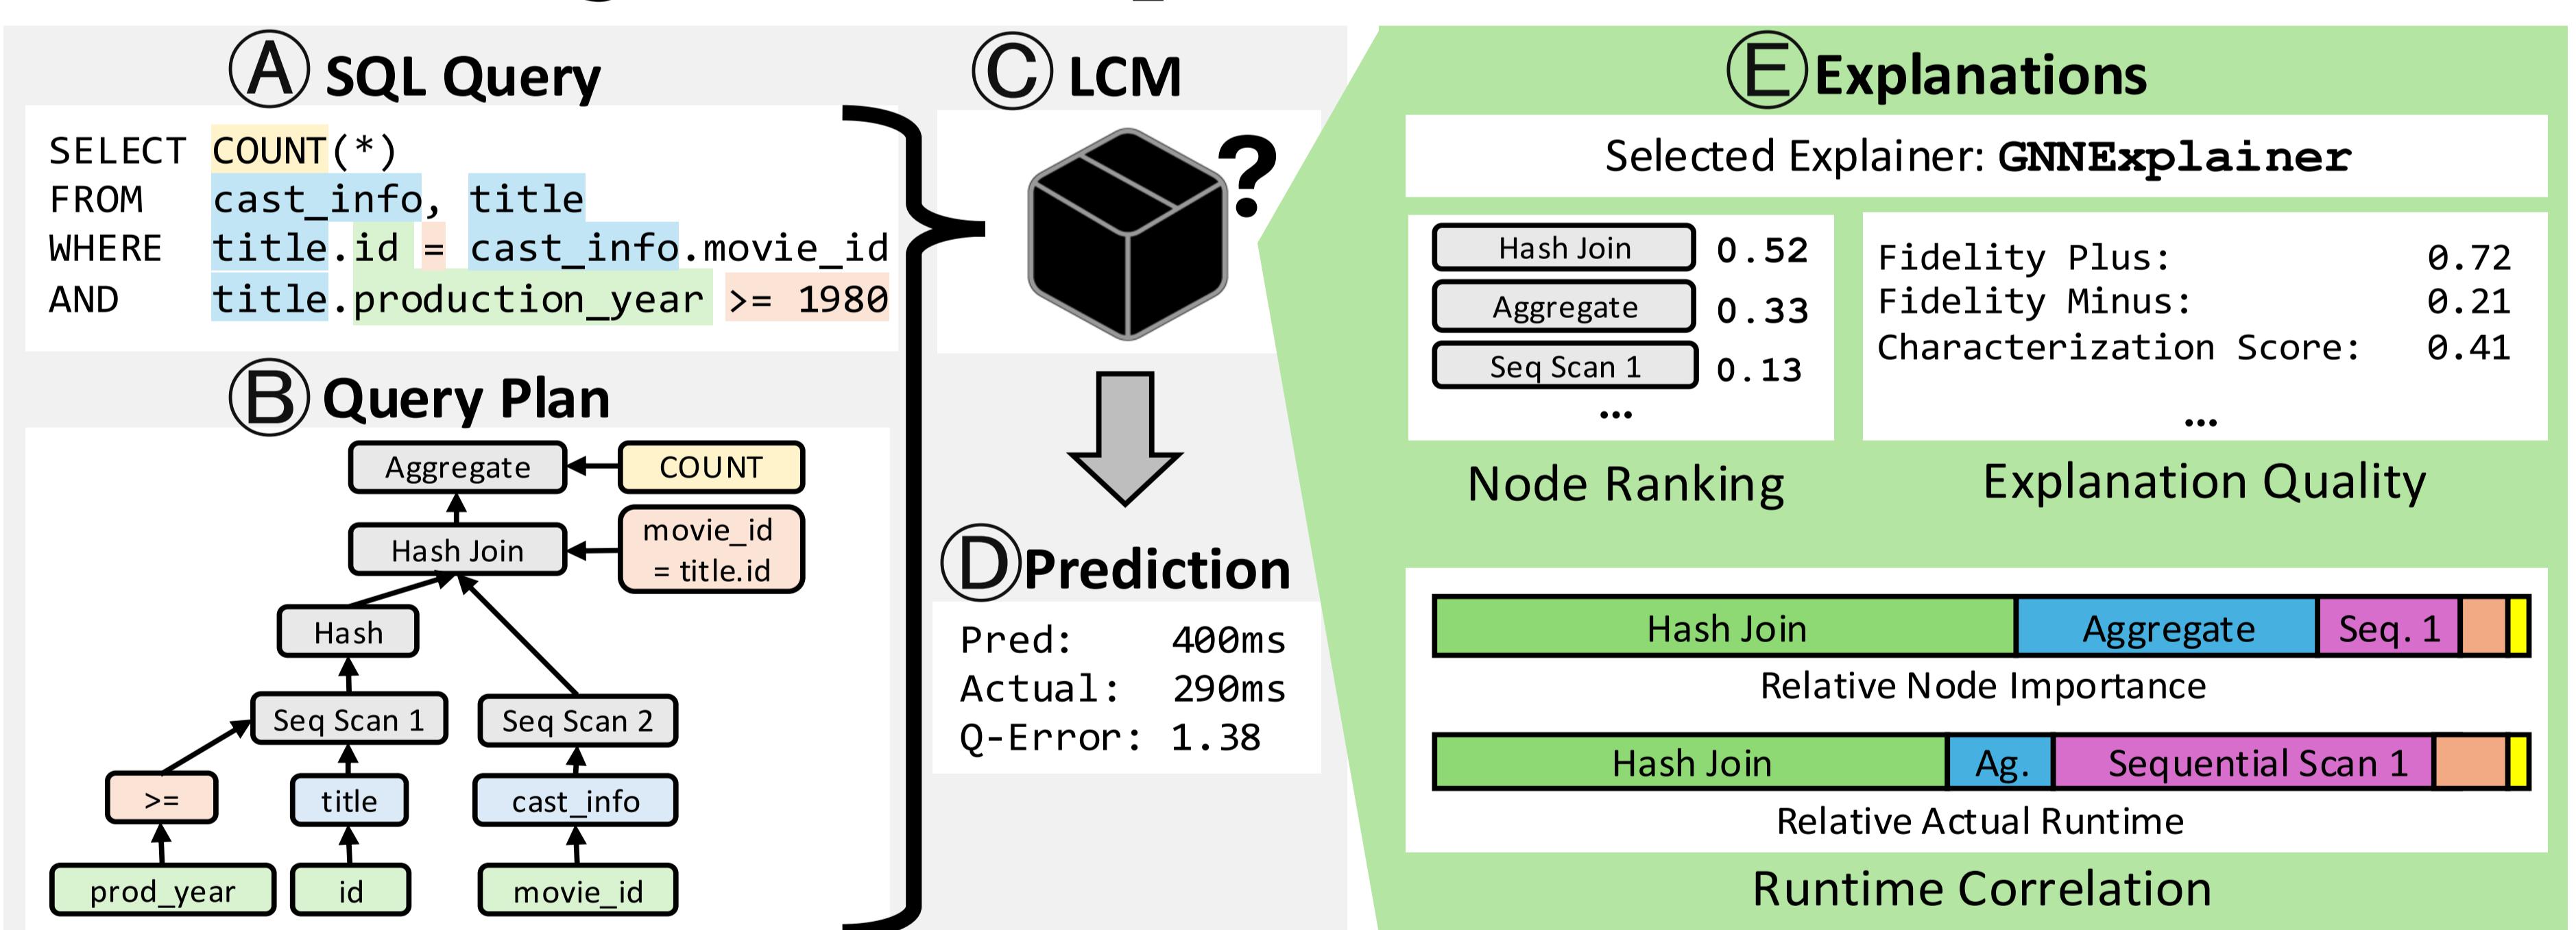
- Learned Cost Models (LCMs) predict query execution costs.
- LCMs show highly promising results, often outperforming classical models.
- However, LCMs are trained with deep learning and thus behave as **black box** in their predictions
- The black-box nature of LCMs makes it difficult to understand **why** and **how** they came up with given results



It is unclear why LCMs mispredict query execution costs – because they act as a black-box during prediction

Contribution: Making LCMs Explainable

- Highlighting what the LCM has learned by showing the **importance** of the **nodes** (=query operators) during GNN inference.
- Analyzing how well existing explainers (e.g. GNNExplainer) can help.
- Using **gradient-based explainers** and **feature-based explainers** by adapting node masking for Zero-Shot cost models.

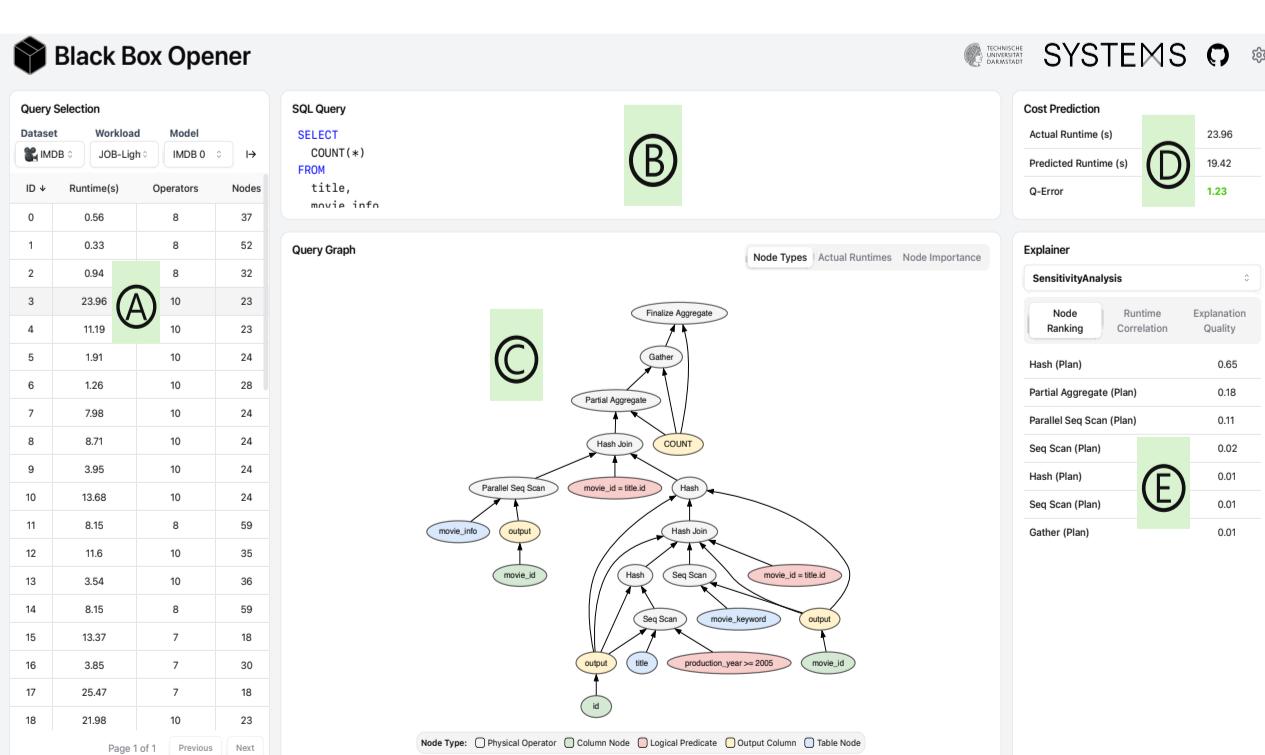


Introducing **novel metrics**:

- Node Ranking*: What node was most important?
- Runtime Correlation*: How much do the importance scores correspond to the actual runtimes?
- Explanation Quality*: How good is the explanation?

Our Idea: Show what LCMs have really learned by visualizing how important is a given node for the final prediction

Interactive Demo



Visualizing node importance scores of Learned Cost Models

Initial Results

- Often, the node importance scores match with the actual operator runtimes.



- Sometimes, LCMs tends to model puts too much **emphasis on aggregation** nodes!
- Still, LCMs often achieve **good predictions** despite **incorrect assumptions**

Outlook

- Exploring more explainer algorithms and LCM architectures
- Take feature importance and subgraph importance into account as well



Paper



Code